



UNIVERSITÀ DI PISA

Department of Philology, Literature, and Linguistics
Master's Degree Program in Digital Humanities

GRADUATION THESIS

MAI(A) the Sound Be With You

**MAIA-AV: A Ground-Up Audio-Visual Benchmark for
Multimodal AI Assessment**

Supervisor

Prof. Alessandro Lenci

Student

Michela Deodati

ID: 597983

Co-Supervisor(s)

Dr. Bernardo Magnini

PhD Student Davide Testa

Academic Year 2025/2026

Contents

Introduzione	1
1 Literature review	4
2 Data Collection	11
2.1 From MAIA to MAIA-AV: Dataset Expansion and Curation	11
2.2 Expanding the Questions and Answers Dataset	23
3 Study Design and Experimental Setup	45
3.1 Overview of the Evaluated Models	45
3.2 Preliminary Analysis as a Diagnostic Tool for Multimodal Assessment . .	50
3.3 Core Reasoning Assessment Framework	83
4 Results and Discussions	112
4.1 Preliminary Analysis Result	112
5 Conclusions	115
6 Future Perspectives	116
A Appendix	123

Introduction

Determining whether the strong performance of Visual Language Models (VLMs) reflects genuine multimodal understanding, rather than the exploitation of superficial correlations and linguistic biases, remains an open challenge. *MAIA: Multimodal AI Assessment* was introduced to address this question, as a native Italian benchmark for evaluating video-based reasoning in Vision-Language Models [22, 21]. MAIA investigates the extent to which VLMs coherently integrate visual and linguistic information. More specifically, it assesses whether model predictions are grounded in visual evidence rather than primarily driven by distributional regularities learned by the language component. This distinction is crucial because a model may answer a question correctly while failing to recognise the same information consistently when it is presented in a discriminative format. In such cases, the prediction may result from linguistic plausibility rather than genuine visual understanding. To address this issue, MAIA employs two aligned tasks: Visual Statement Verification and Open-Ended Visual Question Answering. Through these tasks, the benchmark evaluates not only model accuracy, but also the robustness and consistency of model behaviour across different evaluation formats. The benchmark further defines twelve fine-grained reasoning categories to identify which competencies are used by the models and in which areas linguistic shortcuts, unimodal collapse, or hallucinations emerge. A limitation of the original framework is that the dataset was designed and evaluated primarily on the basis of visual information. Real-world videos, however, are inherently audiovisual: relevant meaning may also be conveyed through dialogue, environmental sounds, music, and other acoustic events. Recent multimodal and omnimodal architectures can process visual, auditory, and textual inputs, but the availability of multiple modalities does not guarantee that they will be integrated effectively. Models may over-rely on a single channel or be negatively affected by redundant or weakly relevant information. Building on this limitation, *Lost in Transcription: Investigating the Role of Audio Information for Video-based Visual Statement Verification* [4] extends the original MAIA research framework to a more complex analysis of the interaction between visual, linguistic, and auditory information,

while preserving its competence-oriented perspective. The study considers two complementary representations of audio: raw signal and automatic transcriptions provided as textual input. By evaluating these representations both independently and in combination with video information, the analysis investigates whether audio contributes meaningfully to task performance and whether one representation has a stronger influence on model predictions. This leads to two research questions:

RQ1: *To what extent does raw audio affect video-based Visual Statement Verification when it is jointly processed with visual and textual information?*

RQ2: *How effectively do multimodal models integrate visual, textual, and auditory information when performing video-based reasoning and understanding tasks?*

The experiments show that raw audio has a limited and strongly context-dependent effect. It is most useful when visual information is unavailable, while its addition to an already informative visual input may instead provide little benefit or introduce noise. The results also show that transcriptions can produce greater interference than raw audio, suggesting that multimodal architectures do not necessarily achieve more effective cross-modal understanding simply because additional modalities are available. These findings must nevertheless be interpreted in relation to the nature of the dataset. The one hundred original MAIA videos were selected and annotated according to their visual content, and the questions were formulated without considering the audio track. Consequently, the benchmark does not systematically require auditory information to solve its items. My thesis project begins from this limitation. The following chapters investigate whether models can integrate multiple input modalities more coherently when the dataset is specifically designed so that auditory information is at least partially relevant. The central hypothesis is that, under these conditions, models may use complementary signals across modalities to reconstruct missing information and develop a more genuinely multimodal understanding of video content, producing answers grounded in the available evidence rather than in linguistic plausibility or chance.

1. Literature review

The rapid evolution of Visual Language Models (VLMs) has progressively expanded the ability of neural systems to jointly process perceptual and linguistic information. This line of research builds on the hypothesis that linguistic meaning cannot be constructed exclusively from distributional regularities learned from textual corpora, but must also be grounded in perceptual experience [2]. Within this context, Visual Question Answering (VQA) emerged as one of the first systematic paradigms for assessing whether a model can answer questions about the content of an image [1]. Subsequent benchmarks, such as GQA, extended this setting by introducing compositional questions involving entities, attributes, and visual relations [9]. However, high task accuracy alone does not provide sufficient evidence of multimodal comprehension. Several studies have shown that models can exploit statistical regularities in answers, questions, or dataset distributions, producing correct predictions without necessarily *understanding* the associated visual content. Evaluation has therefore progressively moved beyond task performance towards the assessment of grounding, understood as the ability to associate a prediction with the relevant perceptual evidence. One widely adopted strategy for evaluating grounding relies on minimally contrastive pairs, in which a correct instance is paired with a foil obtained by modifying a semantically relevant element. *FOIL-it!*, for example, evaluates whether a model can detect localized inconsistencies between an image and its linguistic description [19]. *Winoground*, instead, presents pairs of images and descriptions containing the same lexical elements but arranged according to different compositional relations, showing that recognizing individual entities does not necessarily entail understanding the relational structure connecting them [23]. The transition from images to videos introduces additional challenges. Video understanding requires integrating information distributed over time, recognizing actions and state changes, ordering events, and establishing temporal and causal relations. *Action Genome Question Answering (AGQA)* addresses these requirements through compositional questions concerning objects, actions, and spatio-temporal relations [8]. The *Multi-modal Video Understanding Benchmark (MVBench)*

similarly evaluates a broader range of competencies, including action perception, movement, temporal ordering, and dynamic reasoning [13]. Nevertheless, video benchmarks remain vulnerable to unimodal shortcuts and to strategies that rely predominantly on the most informative frame. *Video Language Model Assessment (VILMA)* addresses this limitation as a zero-shot benchmark based on video–text contrastive pairs [10]. For each complex capability, VILMA introduces a proficiency test designed to verify the preliminary perceptual ability required to solve the corresponding task. The results show that, once these prerequisite conditions are considered, part of the apparently correct performance can be attributed to accidental or spurious strategies. In particular, the examined VLMs do not display a systematic advantage over statistical image-based models on tasks requiring strict temporal grounding. Within this line of research, MAIA introduced a natively Italian benchmark designed to provide a fine-grained evaluation of video-linguistic reasoning [22]. The resource covers twelve reasoning categories related to linguistic phenomena, including causal, counterfactual, implicit, spatial, temporal, sentimental, planning, uncertainty, and out-of-scope reasoning. Its main contribution lies in aligning two tasks built on the same videos and questions: *Visual Statement Verification (VSV)*, where the model selects the correct statement from a minimally contrastive pair, and *Open-Ended Visual Question Answering (OEVQA)*, where the model freely generates the answer. This setting makes it possible to assess not only accuracy but also coherence between discriminative comprehension and linguistic generation. A model may correctly answer a multiple-choice question while failing to reproduce the same information in an open-ended setting, or may behave inconsistently across equivalent linguistic formulations. To capture this phenomenon, MAIA introduces the *Aggregate Accuracy* metric, which rewards only those instances in which the model correctly answers all true–false pairs associated with the same question and simultaneously generates the correct response in the corresponding open-ended task. The reduction in Aggregate Accuracy relative to conventional single-instance accuracy highlights the fragility of apparently strong performance and the need to jointly evaluate robustness, consistency, and grounding. Although VLMs have traditionally focused on text–vision interactions, real videos are intrinsically audiovisual. Their interpretation may depend on music, prosody, background noise, and other acoustic signals

that cannot be reconstructed from visual frames alone. An increasing body of research has accordingly extended multimodal language model architectures to support joint processing of text, images, video, and audio. *ImageBind*, for instance, introduced a shared representation space spanning images, text, audio, depth, thermal information, and inertial movement [7]. Other multimodal language models (MLMs) rely on modality-specific encoders combined with projection or fusion modules that align heterogeneous representations with the linguistic space. *Video-SALMONN: Speech-Enhanced Audio-Visual Large Language Models* introduces a multi-resolution Q-former¹ to align audio signals with synchronized video [20], while *VideoLLaMA2* combines spatio-temporal modeling and acoustic processing within a generative multimodal architecture [3]. More recent omnimodal architectures process multiple modalities within a unified framework. *Qwen2.5-Omni* adopts a thinker–talker architecture that accepts text, image, video, and audio inputs and produces textual or audio responses [14]. The subsequent *Qwen3.5-Omni* architecture integrates visual and audio encoders into a common backbone and introduces explicit temporal references to improve audio–video alignment [15]. These developments demonstrate that heterogeneous signals can be processed within a single architecture. The availability of a modality, however, does not imply that it is used functionally or complementarily. The distinction between **multimodal availability** and **multimodal integration** is essential here: the former indicates that a system can accept multiple input sources, the latter requires the model to select, coordinate, and combine relevant information across modalities. A system may thus be multimodal from an architectural perspective while relying predominantly on unimodal information during prediction. This issue is evident in *Audio-Visual Question Answering* benchmarks such as MUSIC-AVQA, which evaluates spatio-temporal relations in musical performance videos and extends the analysis to objects, entities, and everyday activities [12]. Subsequent studies reveal a strong visual bias, showing that many questions can be answered using the video stream alone, without exploiting the audio channel. Under these conditions, similar audio-only and video-only performance does not demon-

¹A multi-resolution Q-former, such as a Multi-Resolution Causal Q-Former, processes continuous audio, video, or image information at different temporal or spatial scales through distinct query banks, balancing fine-grained signals with broader contextual information.

strate genuine integration, but rather suggests a form of unimodal collapse in which one modality dominates the decision process. To address this problem, *Dave* was introduced as a benchmark in which each question jointly requires visual and acoustic information, making either modality insufficient in isolation [16]. The benchmark distinguishes among action recognition, sound classification, temporal ordering, and multimodal synchronization, a decomposition that helps determine whether an error arises from the failure to perceive an event or from the inability to establish temporal relations between sounds and visual actions. The results show that models may achieve strong performance on individual components while encountering greater difficulty in audio–visual synchronization, indicating that cross-modal integration remains a central limitation. The increasing attention to audio has also led to benchmarks specifically designed to assess acoustic intelligence. *MMAU-Pro* evaluates 49 competencies distributed across speech, environmental sounds, music, and their combinations [11]. These tasks include long-audio comprehension, multi-track reasoning, spatial localization, prosody interpretation, and the integration of dialogue, music, and background noise. The persistently limited performance of current models indicates that acoustic processing cannot be reduced to dialogue recognition alone, since audio also conveys spatial, temporal, emotional, and contextual information that must be evaluated independently. A complementary perspective is provided by *SONIC-01*, a benchmark for omnimodal models based on audio–visual interactions drawn from real conversational and everyday contexts [17]. *SONIC-01* combines open-ended summarization, multiple-choice questions, and temporal event localization, while also showing that replacing raw audio with transcription removes paralinguistic information related to tone and communicative intention. Its results identify temporal localization as one of the most challenging tasks, confirming that the integration of verbal content, acoustic signals, and visual sequences remains an open problem. The comparison between native audio and transcription is also central to my recent work, *Lost in Translation: Investigating the Role of Audio Information for Video-based Visual Statement Verification* [4]. The study extends the MAIA framework by introducing audio into the VSV task and comparing two representations of the acoustic modality: raw audio directly processed by *Qwen2.5-Omni-7B* and automatic transcriptions supplied as additional textual input to *Qwen2.5-VL-7B*. To characterize the

audio content, the videos are divided into four classes:

Music: Instances characterized by repetitive vocalizations, onomatopoeic patterns, strongly repetitive tokens, or explicit lexical cues associated with non-speech audio.

Dialogue: Videos whose transcriptions display sufficiently coherent linguistic content, acceptable log-probability values, low degeneracy, and limited evidence of no-speech segments.

Silence/Noise: Videos with empty or highly unreliable transcriptions, particularly when associated with low log-probability values and relatively high no-speech probabilities.

Unknown: Residual cases that appear speech-related but do not satisfy the dialogue criteria with sufficient confidence, thereby preserving potentially informative content while explicitly marking its uncertainty.

Class labels and speech transcriptions were obtained by combining the outputs of *Whisper-1* and *GPT-4o-transcribe*, together with an assessment of output plausibility. This classification distinguishes results according to the informativeness of the available audio content. The results reveal a marked dominance of the visual channel. Audio provides only a limited benefit, mainly when no visual information is available. Once rich visual input is provided, audio’s contribution becomes negligible and may even introduce noise or errors, with limited robustness observed when a single black frame is added instead. Under visually rich conditions, the difference between with-audio and without-audio settings is minimal, suggesting that audio functions more as a surrogate channel than as information fully integrated with vision. Audio transcriptions exhibit a different behavior. When added to already rich visual input, they reduce accuracy relative to the visual-only condition. This shows that transcription is not a neutral transformation: once projected into the same representational space as the questions and answer options, it can compete with them, increasing the influence of textual regularities and introducing redundant or irrelevant information. *Lost in Translation* thus shows that different representations of the same acoustic content can affect the model’s decision process differently. Raw audio can be

partially ignored when it is not useful, whereas transcription interacts directly with the linguistic component of the model. Introducing an additional modality does not therefore guarantee an improvement, and may instead compromise predictions that were already correct on the basis of reliable visual evidence. These findings must be interpreted in light of the original MAIA design. The benchmark was created to investigate visual content, and annotators were explicitly instructed not to consider audio information, so its questions and answers were not designed to exploit acoustic evidence. The limited contribution of audio may therefore reflect both the models’ modality-integration mechanisms and the limited relevance of audio within the original dataset. Recent literature further emphasizes the importance of identifying the distinct sources of multimodal errors. *MIRAGE* disentangles hallucinations caused by perceptual failures from those arising from incorrect reasoning over correctly represented information [5]. Although primarily focused on visual reasoning, its methodological principle extends naturally to audio–visual settings, where an incorrect response may derive from failed sound recognition, inaccurate temporal synchronization, selection of the wrong modality, or erroneous inference over otherwise correctly represented evidence. Complementary attribution approaches seek to identify which parts of the input contribute to the generation of an answer. *OmniTrace* proposes an attribution framework for omnimodal autoregressive generation that links generated tokens to textual segments, visual regions, and audio/video intervals, aggregating attention signals into cross-modal explanations associated with linguistically interpretable units. This type of analysis is particularly relevant for determining whether a model actually exploits audio information or whether its prediction is driven primarily by text and vision. The current state of the art thus reveals a clear tension between the rapid development of omnimodal architectures and their effective ability to integrate heterogeneous modalities. Contemporary models can receive multiple input types within the same context, yet diagnostic benchmarks provide no guarantee that these sources are used complementarily. Performance may instead depend on a dominant modality, linguistic shortcuts, dataset correlations, or additional signals that introduce noise. A central unresolved limitation is the construction of benchmarks in which audio and visual information are explicitly supervised. Assessing genuine audio–visual comprehension requires items in which audio contributes relevant

information and the correct answer cannot be determined from a single modality alone. Such a design would distinguish complementarity, dominance, substitution, and interference, while evaluating whether the model selects the appropriate evidence and combines it coherently across modalities. The present study follows this direction. Building on the competence-oriented MAIA framework and the recent findings of *Lost in Translation*, it investigates model behavior on a dataset specifically designed to require a genuine contribution from audio in order to answer the questions correctly. The central hypothesis is that, when audio and video provide partial and complementary evidence, multimodal comprehension depends on the model’s ability to reconstruct the missing information through cross-modal integration, avoiding both unimodal collapse and linguistic shortcuts based on plausibility.

2. Data Collection

2.1 From MAIA to MAIA-AV: Dataset Expansion and Curation

Video Selection

The dataset adopted in MAIA-AV builds upon the guidelines of Multimodal AI Assessment (MAIA), comprising 100 short videos, approximately 30 seconds each, selected from YouTube Italia to represent everyday scenarios of Italian culture. Content domains include sports, culinary arts, urban life, and cultural activities, with a strong focus on human presence and highly informative scenes. In the original formulation, video selection and annotation were oriented strictly toward visual comprehension, completely omitting audio content. Building upon this initial set, we collected an additional 100 videos, each uniquely identified from `video001` to `video100`. These identifiers map each video to its corresponding audio stream and text transcriptions. Each video was precisely trimmed to extract targeted segments, creating a controlled corpus of clips. Particular attention was paid to the alignment between dialogue content and the elements appearing in the scene. This targeted selection of relevant dialogue enables a precise investigation of multimodal interactions. By prioritizing scenarios where spoken dialogue is semantically grounded in the visual content, we ensure that the audio signal plays a genuinely informative role in the multimodal instance. This distinction represents the core departure from the original dataset: MAIA-AV is a benchmark engineered to combine and compare distributed evidence across visual, auditory, and textual modalities.

Implementation Details The video acquisition pipeline relies on `yt_dlp` to retrieve targeted temporal segments from selected YouTube videos. Rather than downloading complete sources for post-hoc editing, each clip is specified prior to acquisition via three parameters: source URL, start timestamp, and duration. These parameters configure a pre-

cise temporal interval directly governing the retrieval process. Timestamp conversion into seconds is handled by `parse_time`, which reduces HH:MM:SS or MM:SS strings into integer values:

```
1 def parse_time(time_str):
2     """Converts HH:MM:SS or MM:SS string to seconds."""
3     parts = [int(p) for p in time_str.split(':')]
4
5     if len(parts) == 3:
6         return parts[0] * 3600 + parts[1] * 60 + parts[2]
7     elif len(parts) == 2:
8         return parts[0] * 60 + parts[1]
9     return parts[0]
```

Adding the requested duration to the start position determines the upper boundary of the segment:

```
1 start_sec = parse_time(start_time_str)
2 end_sec = start_sec + int(durata_str)
```

The resulting bounds, `start_sec` and `end_sec`, define the exact temporal window extracted from the source. Output paths are structured prior to downloader configuration, ensuring MP4 container compliance and consistent dataset indexing under the target directory:

```
1 filename_input = input(
2     "Enter output filename (e.g., clip): "
3 ).strip()
4 if not filename_input.endswith(".mp4"):
5     filename_input += ".mp4"
6 output = f"Dataset/video{filename_input}"
```

Prefixing filenames with `video` maintains a unified identifier convention across all dataset instances. Stream selection, temporal trimming, and media normalization are encapsulated within the `yd1_opts` dictionary. Format selection prioritizes high-resolution video streams

alongside optimal audio feeds:

```
1 ydl_opts = {  
2     'format': 'bestvideo[height<=2160]+bestaudio/best',  
3 }
```

This filter selects video streams up to 2160p paired with the highest quality audio stream, falling back to pre-merged streams (/best) if separate tracks are unavailable. Temporal constraints are applied during acquisition via `download_range_func`:

```
1     'download_ranges': download_range_func(  
2         None,  
3         [(start_sec, end_sec)]  
4     ),
```

Embedding segment bounds directly into the retrieval mechanism eliminates separate media cropping steps and reduces bandwidth overhead. To maintain frame-accurate cutting across keyframe-based compressed video formats, the pipeline enforces keyframe generation at boundary points:

```
1     'force_keyframes_at_cuts': True,
```

Forcing keyframes ensures temporal alignment at cut points, preserving action continuity for downstream multimodal evaluation in MAIA-AV. Output formats and post-processing codecs are explicitly declared to standardize the corpus:

```
1     'merge_output_format': 'mp4',  
2     'postprocessor_args': {  
3         'ffmpeg': [  
4             '-c:v', 'libx264',  
5             '-preset', 'fast',  
6             '-c:a', 'aac'  
7         ]  
8     },
```

FFmpeg re-encodes visual streams via H.264 (libx264, fast preset) and acoustic streams

via AAC, eliminating encoding heterogeneity across raw YouTube sources. The unified configuration dictionary combines acquisition parameters, formatting rules, output paths, and authentication credentials:

```
1 ydl_opts = {
2     'format': 'bestvideo[height<=2160]+bestaudio/best',
3     'download_ranges': download_range_func(
4         None,
5         [(start_sec, end_sec)]
6     ),
7     'force_keyframes_at_cuts': True,
8     'merge_output_format': 'mp4',
9     'postprocessor_args': {
10         'ffmpeg': [
11             '-c:v', 'libx264',
12             '-preset', 'fast',
13             '-c:a', 'aac'
14         ]
15     },
16     'outtmpl': output,
17     'cookiefile': 'src/youtube.com_cookies.txt',
18 }
```

Passing a session cookie file (cookiefile) enables retrieval of age-restricted or session-gated content without altering core extraction logic. Execution is wrapped in exception handling blocks, logging errors and returning non-zero exit codes upon failure:

```
1 try:
2     with yt_dlp.YoutubeDL(ydl_opts) as ydl:
3         ydl.download([url])
4         print(f"\nSuccessfully downloaded clip to {output}")
5 except Exception as e:
6     print(f"\nError: {str(e)}")
7     sys.exit(1)
```

This setup allows `yt_dlp` to coordinate format resolution, stream extraction, interval clipping, and FFmpeg transcoding within a single execution pass.

Audio Conversion

In the MAIA-AV framework, the auditory element of each video is regarded as an autonomous source of information rather than a mere supplementary part of the visual narrative. Spoken language, ambient sounds, and acoustic phenomena offer significant evidence that enhances visual data and aids in scene interpretation. Isolating the audio allows this modality to be independently analyzed, ensuring its contribution is maintained throughout subsequent processing stages. For every video, the original audio track is extracted and archived as a separate file. This separation retains the identical temporal parameters and identifier associated with the corresponding audiovisual entity. Such a one-to-one mapping upholds alignment across modalities, enabling visual, acoustic, and textual representations to be traced back to the same foundational event. The resulting audio files function as direct inputs for further pipeline stages, including automatic speech transcription and specialized audio analyses. The extraction process neither generates a modified version of the original content nor alters it; instead, it separates the auditory channel, allowing independent processing while remaining completely synchronized with the originating video content.

Implementation Details The audio extraction stage converts each audiovisual dataset instance into a standalone acoustic representation while preserving multimodal alignment with the source video. Implemented via Python’s `subprocess` module, the pipeline invokes FFmpeg to automate batch processing across the complete video collection. Input and output paths are declared explicitly at execution setup:

```
1 INPUT_FOLDER = Path("data/video")
2 OUTPUT_FOLDER = Path("data/audio")
3 VIDEO_EXTENSIONS = {".mp4"}
```

Decoupling input and output directories isolates raw video resources and establishes an aligned audio mirror directory. Restricting VIDEO_EXTENSIONS to MP4 matches the output format produced by the preceding acquisition stage. The core conversion routine is encapsulated within `mp4_to_mp3`:

```
1 def mp4_to_mp3(  
2     input_path: Union[str, Path],  
3     output_path: Optional[Union[str, Path]] = None,  
4     bitrate: str = "192k",  
5 ) -> Path:  
6     """Extract the audio from an MP4 file and save it as MP3."""
```

When an explicit destination path is omitted, the function retains the base filename and replaces the container extension:

```
1     input_path = Path(input_path)  
2     output_path = (  
3         input_path.with_suffix(".mp3")  
4         if output_path is None  
5         else Path(output_path)  
6     )
```

This mapping strategy guarantees cross-modal alignment across representations (e.g., `video001.mp4` yields `video001.mp3`). Prior to invoking FFmpeg, parent directory existence is enforced:

```
1     output_path.parent.mkdir(  
2         parents=True,  
3         exist_ok=True,  
4     )
```

This step allows nested input folder hierarchies to be automatically replicated within the output directory. Extraction is executed using a single `subprocess.run` call:

```
1     subprocess.run(  
2         [  

```

```

3         "ffmpeg",
4         "-y",
5         "-i",
6         str(input_path),
7         "-vn",
8         "-c:a",
9         "libmp3lame",
10        "-b:a",
11        bitrate,
12        str(output_path),
13    ],
14    check=True,
15 )

```

The `-vn` flag discards video streams, ensuring that FFmpeg isolates the acoustic signal rather than performing generic media transcode operations. Acoustic streams are normalized using the `libmp3lame` codec at a constant target bitrate of 192 kbit/s (`-b:a 192k`). Passing `-y` overwrites pre-existing destination files without prompt interruption, facilitating seamless pipeline re-execution. Setting `check=True` ensures runtime errors raise a `CalledProcessError` rather than failing silently. Upon completion, the function returns the confirmed target path:

```

1     return output_path

```

Batch processing begins by recursively indexing and sorting compatible media files:

```

1 def find_videos(folder: Path) -> list[Path]:
2     """Return all supported video files found recursively."""
3     return sorted(
4         path
5         for path in folder.rglob("*")
6         if path.is_file()
7         and path.suffix.lower() in VIDEO_EXTENSIONS
8     )

```

Sorting file paths guarantees deterministic execution order across operating environments.

Directory validity is verified before initializing conversion loops:

```
1 if not INPUT_FOLDER.exists():
2     raise FileNotFoundError(
3         f"Input folder not found: {INPUT_FOLDER}"
4     )
5 videos = find_videos(INPUT_FOLDER)
6 if not videos:
7     print(f"No videos found in: {INPUT_FOLDER}")
8     return
```

Validating the input directory early avoids redundant FFmpeg invocations when processing empty or invalid paths. For each video, relative path resolution preserves source sub-directory structures under OUTPUT_FOLDER:

```
1 for index, video_path in enumerate(videos, start=1):
2     relative_path = video_path.relative_to(INPUT_FOLDER)
3     audio_path = (
4         OUTPUT_FOLDER
5         / relative_path.with_suffix(".mp3")
6     )
```

Replicating source directory trees ensures structural parity between visual and acoustic datasets. Each file conversion is executed inside localized exception blocks:

```
1     try:
2         mp4_to_mp3(
3             video_path,
4             audio_path,
5         )
6         print("Audio extraction completed.")
7     except subprocess.CalledProcessError as error:
```

```
8         print(f"FFmpeg error: {error}")
```

Capturing `CalledProcessError` exceptions locally prevents single-file execution failures from interrupting broader batch operations while logging specific processing errors.

Transcription Extraction

Another key representation in our dataset is the audio transcription, following the methodology established in our previous work [4]. Transcriptions were generated using the *Faster-Whisper* model (large-v3), a reimplementation of *OpenAI's Whisper* powered by the *CTranslate2* fast inference engine for Transformer architectures. The model processes raw audio content directly, with the language parameter set to Italian. Additionally, a *Voice Activity Detection* (VAD) filter is applied to isolate segments containing actual speech, thereby minimizing non-speech transcriptions. The overall transcription procedure aims to convert spoken dialogue into textual content that can be reliably recognized and processed by downstream models. This methodology represents an evolution compared to previous approaches on acoustic content analysis [4], which jointly employed *Whisper-1* and *GPT-4o-transcribe* to compare and align outputs for higher precision. In contrast, the current pipeline streamlines the transcription workflow by deploying a single, locally executed model, providing an efficient and reproducible framework.

Implementation Details The transcription stage was implemented as an independent component of the data preparation pipeline to associate each video with a textual representation of its spoken content. Powered by *Faster-Whisper* and the large-v3 model executed locally on GPU, this setup ensures full reproducibility while eliminating dependencies on external cloud services. The execution environment explicitly configures a selected CUDA device and establishes a local cache for model files:

```
1     os.environ["CUDA_VISIBLE_DEVICES"] = "0"  
2  
3     HF_CACHE = Path.cwd() / ".hf_cache"  
4     HF_CACHE.mkdir(parents=True, exist_ok=True)
```

```

5
6     os.environ["HF_HOME"] = str(HF_CACHE)
7     os.environ["HF_HUB_CACHE"] = str(HF_CACHE / "hub")

```

Utilizing a local cache confines model weights to the working directory, preventing redundant downloads across system locations. The transcription model is subsequently instantiated on CUDA using float16 precision:

```

1     MODEL_NAME = "large-v3"
2     LANGUAGE = "it"
3
4     model = WhisperModel(
5         MODEL_NAME,
6         device="cuda",
7         device_index=0,
8         compute_type="float16",
9         download_root=str(HF_CACHE / "hub"),
10    )

```

Setting language="it" explicitly conditions decoding on Italian. This parameter bypasses automatic language identification and directly aligns model execution with the linguistic profile of the dataset. Prior to processing, the pipeline recursively retrieves and filters all .mp4 files within the input directory, sorting the resulting collection to guarantee deterministic execution order:

```

1     def find_videos(folder: Path) -> list[Path]:
2         return sorted(
3             path
4             for path in folder.rglob("*")
5             if path.is_file()
6             and path.suffix.lower() in VIDEO_EXTENSIONS
7         )

```

The core transcription workflow relies on `model.transcribe()`. By directly parsing

the audio embedded within the video container, Faster-Whisper eliminates intermediate extraction steps and yields chronologically ordered transcription segments:

```
1     segments, _ = model.transcribe(  
2         str(video_path),  
3         language=LANGUAGE,  
4         beam_size=5,  
5         vad_filter=True,  
6         condition_on_previous_text=False,  
7     )
```

Three key hyperparameters govern decoding precision and stability:

- `beam_size=5`: Activates beam-search decoding, evaluating multiple competing hypotheses during generation to balance accuracy and computational overhead.
- `vad_filter=True`: Enables Voice Activity Detection to strip environmental noise, music, and unvoiced pauses, preventing non-linguistic signals from reaching the ASR engine.
- `condition_on_previous_text=False`: Disables cross-segment context conditioning, isolating individual speech segments and mitigating error propagation across the audio stream.

The resulting segment iterator is aggregated into a continuous transcript by stripping whitespace, discarding empty segments, and concatenating valid text chronologically:

```
1     return " ".join(  
2         segment.text.strip()  
3         for segment in segments  
4         if segment.text.strip()  
5     )
```

Omitting downstream linguistic post-processing or semantic filtering preserves an unadulterated representation of raw ASR performance. The corpus is processed sequentially, using each video filename as a unique primary key:


```

1     transcriptions = {}
2
3     for index, video_path in enumerate(videos, start=1):
4         try:
5             transcription = transcribe_video(video_path)
6             transcriptions[video_path.name] = transcription
7         except Exception as error:
8             transcriptions[video_path.name] = (
9                 f"ERROR: {error}"
10            )

```

Robust error handling isolates processing failures on individual files, preventing execution halts and cleanly decoupling unvoiced media from runtime exceptions. Finally, extracted transcripts are exported to a structured CSV file indexed by `video_name` and `transcription`:

```

1     writer = csv.DictWriter(
2         csv_file,
3         fieldnames=[
4             "video_name",
5             "transcription"
6         ],
7         delimiter=";",
8     )
9     writer.writeheader()
10    writer.writerows(
11        {
12            "video_name": video_name,
13            "transcription": transcription,
14        }
15        for video_name, transcription
16        in transcriptions.items()

```

2.2 Expanding the Questions and Answers Dataset

The construction of linguistic data in MAIA-AV mirrors the methodological framework established in the original MAIA benchmark. The process retains the question as the foundational element for subsequent answer collection and the creation of an evaluation dataset. In the original MAIA benchmark, each video was associated with 2,400 questions. Expert annotators, operating under specific category guidelines, were responsible for generating these questions. The questions were designed in an open-ended format to preclude simple binary responses, such as yes or no. Given that the original benchmark was intended for models based solely on visual input, annotators were explicitly instructed to disregard all acoustic data during the question formulation phase. MAIA-AV extends the original framework by modifying the informational basis for question generation. This version includes a total of 300 questions, with each video offering three inquiries that delve into spatial, temporal, and causal dimensions of the presented information. These categories are aligned with the structure of the original MAIA benchmark. The spatial category focuses on identifying locations or entities, either visible or referenced, along with detailing the spatial relationships between them. The temporal category examines event sequences, identifying when specific occurrences take place and establishing temporal relationships such as preceding, succeeding, or simultaneous actions. The causal category investigates the reasons behind or consequences of actions, events, or behaviors observed within the video context. The primary distinction between the original MAIA and its successor, MAIA-AV, is the integration of audio information, which is crucial in the latter version. Certain questions are specifically designed to address audio elements of the video, including character dialogue, vocal commands to infotainment systems, and ambient sounds. This addition incorporates linguistic content that connects information across both visual and auditory channels. The process of acquiring answers adheres to the MAIA guidelines. For the newly devised set of 300 questions, data was gathered through

a questionnaire distributed among four volunteers. These individuals are native speakers of Italian, each born in Italy, and possessing at least a high school diploma. These participants engaged with 100 videos, responding to the corresponding 300 questions, thus enabling the collection of independent responses and four distinct articulations of each answer. The structured approach remains intact: the videos operate as the multimodal information source, the questions identify the specific information to be extracted, and the human-generated answers provide the necessary linguistic reference for constructing subsequent components of the benchmark.

Caption-Foil Generation

Building upon the human-annotated responses, the construction of *MAIA-AV* proceeds with the automated generation of *caption-foil* pairs for the Visual Statement Verification (VSV) task. This procedure follows the methodology established in *MAIA*: each question-answer pair is first converted into a declarative true statement (*True Statement*, TS) and subsequently modified to create a minimally contrastive false statement (*False Statement*, FS). Consequently, the caption and foil share maximum linguistic structure, differing primarily in question-relevant information. While the original benchmark collects eight human responses per question to yield eight TS-FS pairs, *MAIA-AV* adopts a scaled-down approach: four independent responses per question yield four captions and four corresponding foils. With three questions per video (spatial, temporal, and causal), each video comprises twelve caption-foil pairs. The automated generation employs GPT-4o-2024-08-06, matching the exact model version used in the original *MAIA* pipeline. Generation proceeds in two distinct stages. In the initial stage, the model receives only the question and a single human answer. Rather than generating novel content, the prompt directs the model to synthesize the question-answer pair into a single declarative statement preserving the original semantic content. The instructions mandate minimal alterations to annotators’ phrasing while requiring the conversion of first-person expressions into an impersonal stance. Accordingly, phrases like “I think that” or “I believe that” are removed from the caption without discarding any underlying uncertainty. Restricting outputs to concise declarative sentences leverages natural annotator variance while avoiding

reliance on the model for factual core generation. This initial phase employs the prompt used for True Statement generation in *MAIA*:

Sei un assistente che crea affermazioni in italiano sui video.
Data una domanda Q e una risposta A relative a un video, devi creare un'affermazione S basata su A.
Mentre generi S, cerca di non alterare le parole che compongono A.
Se A include verbi o frasi in prima persona (ad es., 'penso', 'credo'), riformula S in modo impersonale, evitando la prima persona.
L'affermazione deve essere una frase breve e dichiarativa.¹

The instructions include few-shot examples demonstrating question–answer fusion, with specific coverage of negative responses, non-deducible information, and expressions of uncertainty. Appended to the prompt, the target question and annotator response appear in Q: and A: format, conditioning the model to output solely the statement S:. The second stage processes the generated caption to yield the corresponding foil. Unlike the initial phase, this prompt adapts to the target question category. The core methodology relies on constructing *minimal pairs*: the model modifies only the information necessary to invalidate the caption, avoiding extraneous additions and preserving sentence structure. Consequently, distinguishing between caption and foil tests a model’s ability to ground claims in audiovisual content rather than relying on surface linguistic cues. This design adheres to *MAIA*, directing GPT-4o to alter exclusively elements relevant to the specified semantic domain. Across all categories, the system prompt assigns a uniform role:

Sei un assistente progettato per creare foil a partire dalle caption.²

Subsequent instructions vary by targeted semantic dimension. For **spatial** questions, foil generation alters strictly the entity’s position or location:

¹English translation: *You are an assistant that creates statements in Italian about videos. Given a question Q and an answer A concerning a video, you must create a statement S based on A. While generating S, try not to alter the words composing A. If A includes first-person verbs or phrases (e.g., 'I think,' 'I believe'), rephrase S to be impersonal, avoiding a first-person perspective. The statement should be a concise, declarative sentence.*

²English translation: *You are an assistant designed to create foils based on captions.*

Data una caption in italiano (C) riguardante la posizione o la localizzazione di qualcuno o qualcosa, il tuo compito è creare il suo foil (F) modificando solo l'informazione spaziale.

Non aggiungere altre informazioni rispetto a quanto dichiarato in C.³

An accompanying example relocates a woman standing in a poppy field to a school setting. Preserving the subject and surrounding text isolates spatial information as the sole determinant of truth value. For **temporal** questions, instructions dictate altering the temporal relationships within the caption:

Data una caption in italiano (C) riguardante gli eventi e quando avvengono, il tuo compito è creare il suo foil (F) modificando solo l'informazione temporale.

Non aggiungere altre informazioni rispetto a quanto dichiarato in C.⁴

An illustrative example converts the chronological order of two events into an alternative sequence, preserving core semantic content while modifying event timing or order. For **causal** questions, foil construction alters the cause-and-effect relationship expressed in the caption:

Data una caption in italiano (C) riguardante le cause o gli effetti degli eventi, il tuo compito è creare il suo foil (F) modificando le cause o l'effetto dell'evento principale.

Non aggiungere altre informazioni rispetto a quanto dichiarato in C.⁵

³English translation: *Given an Italian caption (C) regarding the position or location of someone or something, your task is to create its foil (F) by changing only the spatial information. Don't add other information with respect to what is stated in C.*

⁴English translation: *Given an Italian caption (C) regarding events and when they happen, your task is to create its foil (F) by changing only the temporal information. Don't add other information with respect to what is stated in C.*

⁵English translation: *Given an Italian caption (C) regarding the causes or the effects of events, your task is to create its foil (F) by changing the causes or the effect of the main event. Don't add other information with respect to what is stated in C.*

The provided example retains the principal action, such as a person beginning to run, while altering solely the underlying cause. This focuses the contrast on the causal link rather than arbitrary scene modifications. In the final task setup, option ordering is randomized, preventing positional bias and ensuring performance metrics reflect semantic comprehension rather than structural pattern exploitation.

Implementation Details The *caption-foil* generation pipeline is implemented in `caption_foil.py`. The script relies on pandas for tabular data processing and the OpenAI Python client for model inference. Both caption and foil generation utilize `gpt-4o-2024-08-06`, with authentication managed via the `OPENAI_API_KEY` environment variable. Input and output paths, along with the expected annotation structure, are defined at initialization:

```
1 MODEL = "gpt-4o-2024-08-06"
2 RAW_ANSWERS_DIR = Path("data/vsv/raw_human_answers")
3 QUESTIONS_FILE = Path("data/vsv/question_classification.csv")
4 OUTPUT_DIR = Path("data/vsv/caption-foil")
5 CHECKPOINT_FILE = OUTPUT_DIR / "caption_foil_checkpoint.csv"
6 OUTPUT_FILE = OUTPUT_DIR / "caption_foil.csv"
7 RANDOM_SEED = 42
8 EXPECTED_MACROAREAS = {"spaziale", "temporale", "causale"}
9 EXPECTED_ANNOTATORS = 4
10 client = OpenAI(api_key=os.environ["OPENAI_API_KEY"])
```

Questions are loaded from `data/vsv/question_classification.csv`. Only the video identifier, principal dimension, and question text are required for this stage:

```
1 questions = pd.read_csv(
2     QUESTIONS_FILE,
3     sep=";",
4     encoding="utf-8-sig",
5 )
6 questions = questions[
```

```

7     ["video_name", "principal_dimension", "question_text"]
8 ].copy()

```

Prior to processing, the script verifies column availability, discards missing values, normalizes category labels, and confirms the presence of only the three expected macroareas. Video identifiers are standardized using `normalize_video_name()`, which extracts the numeric identifier and converts it to the format used across the dataset:

```

1 def normalize_video_name(value):
2     value = str(value).strip()
3     match = re.search(r"(\d+)", value)
4     if not match:
5         raise ValueError(
6             f"Impossibile ricavare il numero del video da: {value}"
7         )
8
9     return f"video{int(match.group(1)):03d}.mp4"

```

Additionally, the script ensures that each video is associated with exactly one spatial, one temporal, and one causal question. Duplicate video–macroarea pairs are rejected prior to generation. Human responses are automatically loaded from all CSV files within `data/vsv/raw_human_answers`. The file count is validated against the expected number of annotators:

```

1 files = sorted(RAW_ANSWERS_DIR.glob("*.csv"))
2 if len(files) != EXPECTED_ANNOTATORS:
3     raise ValueError(
4         f"Attesi {EXPECTED_ANNOTATORS} file di annotatori in "
5         f"{RAW_ANSWERS_DIR}, trovati {len(files)}"
6     )

```

Each file must contain the columns `titolo_video`, `macro-area`, and `risposta`. For backward compatibility, the alternative column name `risposte` is also accepted and automatically renamed. Annotator identifiers are extracted directly

from the filenames:

```
1 if "risposta" not in df.columns and "risposte" in df.columns:
2     df = df.rename(columns={"risposte": "risposta"})
3 df = df[["titolo_video", "macroarea", "risposta"]].copy()
4 df["video_name"] = df["titolo_video"].map(
5     normalize_video_name
6 )
7 df["macroarea"] = (
8     df["macroarea"]
9     .astype(str)
10    .str.lower()
11    .str.strip()
12 )
13 df["annotator"] = file.stem.removeprefix("risposte_")
```

Questions and responses are subsequently aligned via the shared `video_name` and `macroarea` fields. The merge is constrained as *many-to-one*, matching multiple annotator responses to a single question:

```
1 data = answers.merge(
2     questions,
3     on=["video_name", "macroarea"],
4     how="left",
5     validate="many_to_one",
6 )
```

Additional consistency checks validate question keys against annotator responses. Missing answers or orphan responses interrupt execution prior to invoking any API requests. Responses corresponding to the same video and macroarea are grouped into a common pool. The pool identifier is formed by concatenating the video identifier with the corresponding macroarea:

```
1 data["video_id"] = (
2     data["video_name"].str.removesuffix(".mp4")
```



```

3 )
4 data["pool_id"] = (
5     data["video_id"] + "_" + data["macroarea"]
6 )

```

Each pool must contain exactly four responses. Individual entries are assigned a sequential position and a unique identifier:

```

1 data["pool_item"] = (
2     data.groupby("pool_id").cumcount() + 1
3 )
4 data["id"] = (
5     data["pool_id"]
6     + "_"
7     + data["pool_item"].astype(str).str.zfill(2)
8 )

```

API calls are centralized within `call_openai()`, which serves both caption and foil generation functions:

```

1 response = client.chat.completions.create(
2     model=MODEL,
3     messages=messages,
4     max_tokens=max_tokens,
5 )
6 return response.choices[0].message.content.strip()

```

The function executes up to five attempts for recoverable API failures. Rate-limit errors trigger a 30-second sleep interval, whereas generic API errors pause execution for 10 seconds. Authentication errors terminate execution immediately:

```

1 except RateLimitError:
2     if attempt == retries - 1:
3         raise
4     time.sleep(30)

```

```

5 except AuthenticationError:
6     raise RuntimeError(
7         "Errore di autenticazione: controlla OPENAI_API_KEY."
8     )
9 except APIError:
10     if attempt == retries - 1:
11         raise
12     time.sleep(10)

```

For each annotated response, caption and foil generation proceed sequentially. The former accepts the question and human answer as input, whereas the latter accepts the generated caption along with its associated macroarea:

```

1 caption = generate_caption(
2     row["question_text"],
3     row["risposta"],
4 )
5 foil = generate_foil(
6     caption,
7     row["macroarea"],
8 )

```

Outputs are post-processed via `clean_generation()`, which strips optional S: or F: prefixes returned by the model:

```

1 def clean_generation(text, prefix):
2     text = text.strip()
3     if text.startswith(f"{prefix}: "):
4         text = text[3:]
5     elif text.startswith(f"{prefix}:"):
6         text = text[2:]
7
8     return text.strip()

```

To accommodate long-running executions, the script implements a checkpointing mechanism. Cached entries are reused only if their question, annotator, and human answer match the current source data and both generated fields are non-null:

```
1 checkpoint_is_valid = (  
2     saved is not None  
3     and str(saved["question"]) == str(row["question_text"])  
4     and str(saved["annotator"]) == str(row["annotator"])  
5     and str(saved["human_answer"]) == str(row["risposta"])  
6     and pd.notna(saved["caption"])  
7     and pd.notna(saved["foil"])  
8 )
```

Newly generated entries are appended incrementally to `caption_foil_checkpoint.csv`:

```
1 pd.DataFrame(generated_rows).to_csv(  
2     CHECKPOINT_FILE,  
3     index=False,  
4     encoding="utf-8-sig",  
5 )
```

Following generation, caption and foil positions are randomized independently within each question category. A fixed random seed guarantees reproducibility while maintaining a balanced distribution of options across positions:

```
1 rng = random.Random(RANDOM_SEED)  
2 for category, indices in (  
3     df.groupby("question_category").groups.items()  
4 ):  
5     indices = list(indices)  
6     n = len(indices)  
7     labels = [0] * (n // 2) + [1] * (n // 2)  
8  
9     if n % 2:  
10         labels.append(rng.randint(0, 1))
```

```
11
12     rng.shuffle(labels)
```

The resulting target variable indicates the position of the correct caption: a value of 0 places the caption in answer1, whereas a value of 1 assigns it to answer2:

```
1 df["answer1"] = df.apply(
2     lambda row:
3     row["caption"]
4     if row["target"] == 0
5     else row["foil"],
6     axis=1,
7 )
8 df["answer2"] = df.apply(
9     lambda row:
10    row["foil"]
11    if row["target"] == 0
12    else row["caption"],
13    axis=1,
14 )
```

The final dataset is saved to data/vsv/caption-foil/caption_foil.csv. Alongside the randomized options and target labels, the output retains all metadata required to trace each pair back to its originating video, question, annotator, and human response:

```
1 columns = [
2     "id",
3     "video_id",
4     "video_name",
5     "pool_id",
6     "pool_item",
7     "question_category",
8     "question",
```

```

9     "annotator",
10    "human_answer",
11    "caption",
12    "foil",
13    "answer1",
14    "answer2",
15    "target",
16 ]
17 final_df[columns].to_csv(
18     OUTPUT_FILE,
19     index=False,
20     encoding="utf-8-sig",
21 )

```

Caption–Foil Quality Validation

The automatic generation of a foil does not, by itself, guarantee that the resulting sentence constitutes a valid negative counterpart of the corresponding caption. Although generation is constrained through category-specific prompts, the model may introduce changes that are semantically unrelated to the targeted reasoning phenomenon, alter more information than necessary, or produce a sentence that remains compatible with the original caption. For this reason, the generated *caption–foil* pairs undergo an additional semi-automatic validation stage before being included in the final VSV evaluation set. The validation procedure follows the two-step strategy adopted in *MAIA*, adapted here to the three reasoning dimensions considered in *MAIA-AV*. The first stage is a *structural check*, whose purpose is not simply to determine whether the caption and foil differ, but to verify that their difference concerns the semantic dimension the corresponding question is meant to evaluate. This distinction is fundamental for preserving the diagnostic nature of the benchmark. If, for instance, a temporal foil modified the identity of an entity rather than the ordering of two events, a model could distinguish the two statements through entity recognition alone, without any temporal reasoning. Similarly, a causal pair would stop

providing a controlled test of causal understanding if the foil simultaneously changed the cause, the participants, and the location of the event. Structural verification is what keeps the contrast introduced during foil generation aligned with the competency each question is meant to probe. For **causal** instances, validation requires the foil to modify the cause, the effect, or the causal relation expressed by the caption, while preserving the remaining content as much as possible. For **temporal** instances, the contrast must concern temporal information, including event ordering, before–after relations, or the moment at which an event occurs. The **spatial** category calls for a slightly broader criterion. Unlike the original *MAIA* taxonomy, where spatial reasoning is split into Total and Partial subcategories, the spatial questions in *MAIA-AV* combine phenomena from both settings. The validation prompt therefore accepts changes involving locations, positions, directions, and relative spatial relations, including relations that hold only during a particular phase of the video. This broader formulation avoids forcing the new dataset into a distinction its annotation scheme does not retain, while keeping the core requirement that the foil differ from the caption specifically in its spatial content. Pairs that satisfy this first criterion move on to a second validation step based on *Natural Language Inference* (NLI), where the caption is treated as the premise and the foil as the hypothesis. The model assigns one of three standard NLI labels: *Entailment*, *Contradiction*, or *Neutral*. The expected relation is *Contradiction*, since the foil is designed to offer an alternative statement that should not hold under the same interpretation as the caption. This step addresses a different source of error than the structural check. A foil may correctly manipulate the intended semantic dimension while still failing to create a sufficiently distinct alternative. Conversely, a sentence may differ substantially from the caption yet fail the structural test because the modification does not isolate the phenomenon under evaluation. The NLI result is nevertheless treated as a diagnostic signal rather than an infallible automatic decision, particularly for causal and, to a lesser extent, spatial statements. Two alternative causes, for example, may be logically compatible in isolation even though only one is intended to describe the event represented in the dataset. In such cases a general-purpose NLI model may classify the relation as *Neutral* rather than *Contradiction*. A similar pattern was observed during the original *MAIA* validation, where qualitative inspection showed that many pairs labelled as

neutral were nevertheless valid caption–foil contrasts. For this reason, only pairs that pass the structural check and receive the *Contradiction* label are accepted automatically. Structurally valid pairs classified as *Neutral* are marked for manual review, while *Entailment* cases are given higher review priority, since they provide stronger evidence that the foil may remain semantically compatible with the caption. Category-sensitive structural verification and category-independent semantic inference thus provide two complementary safeguards. The former preserves the correspondence between each pair and the reasoning ability it is meant to test. The latter checks whether the resulting negative statement establishes an adequate semantic opposition with the corresponding positive statement. Importantly, neither stage is used to silently discard uncertain data. Every instance flagged as structurally invalid, entailed, neutral, or affected by execution errors is isolated in a dedicated review set, manually inspected, and, whenever the validation reveals an actual problem, corrected so that the intended semantic contrast is restored. The corrected pair is then run back through the same validation procedure. Human verification at this final stage keeps automatic classifier errors from propagating into the benchmark, while the pipeline as a whole remains predominantly automated.

Automatic Validation and Manual Revision

The automatic validation procedure was used as a quality-control step to identify caption–foil pairs that required further inspection. The structural check evaluated whether the foil modified the semantic dimension associated with the corresponding category, while the NLI stage assessed the relation between caption and foil in terms of *Entailment*, *Contradiction*, or *Neutral*. The output of these automatic checks was not treated as a definitive decision on the validity of each pair. In particular, some instances labelled as *FAIL_STRUCTURAL*, *Neutral*, or *Entailment* could still represent acceptable contrasts when considered in relation to the original question, the human answer, and the intended semantic category. This is especially relevant for cases in which the contrast is expressed through the answer-bearing element itself, such as a change of location in spatial questions or a change in the participant occupying a specific temporal position. Similarly, in causal questions, two alternative explanations may be classified as *Neutral* by an NLI

model even when only one of them is supported by the video. For this reason, the complete caption–foil file was subsequently revised manually. Each automatically flagged instance was inspected together with its question, human answer, semantic category, caption, and foil. The automatic labels were used only to prioritize potentially problematic cases, not to determine whether a pair had to be discarded or modified. When a caption–foil pair was considered semantically valid and sufficiently consistent with the intended contrast, it was retained without changes. Whenever the manual inspection revealed an actual issue, such as a foil targeting the wrong semantic dimension, an ambiguous or insufficiently discriminative contrast, a formulation that did not adequately reflect the original human answer, or a mismatch between the question and the generated statements, the pair was manually corrected instead. The final dataset thus reflects a combination of automatic quality-control signals and human verification. Modifications were introduced only when they were considered necessary after direct inspection of the corresponding instance.

Implementation Details The validation procedure is implemented as a separate post-processing stage operating on the output of the caption–foil generation pipeline. Its input is the complete dataset stored in `data/vsv/caption-foil/caption_foil.csv`. Validation is performed on the canonical `caption` and `foil` fields rather than on the randomized `answer1` and `answer2` columns used by the final VSV task, since option randomization affects only presentation order and must not interfere with the semantic relation being assessed. The relevant paths and expected semantic categories are defined explicitly:

```
1 INPUT_PATH = Path(  
2     "data/vsv/caption-foil/caption_foil.csv"  
3 )  
4 OUTPUT_DIR = Path(  
5     "data/vsv/caption-foil/validation"  
6 )  
7  
8 STRUCTURAL_OUTPUT = (  
9     OUTPUT_DIR /  
10    "caption_foil_structural_validation.csv"
```



```

11 )
12
13 FINAL_OUTPUT = (
14     OUTPUT_DIR /
15     "caption_foil_validation.csv"
16 )
17
18 VALID_CATEGORIES = {
19     "causale",
20     "temporale",
21     "spaziale",
22 }

```

Before any API request is issued, the script verifies the availability of the fields required for validation and checks that all category labels belong to the three supported macroareas:

```

1  REQUIRED_COLUMNS = {
2      "id",
3      "question_category",
4      "caption",
5      "foil",
6  }
7
8  missing_columns = (
9      REQUIRED_COLUMNS - set(df.columns)
10 )
11
12 if missing_columns:
13     raise ValueError(
14         "Mancano le seguenti colonne nel CSV: "
15         + ", ".join(sorted(missing_columns))
16     )

```

The first validation stage selects a category-specific prompt according to the value stored in `question_category`. All three prompts share the same decision format but define validity according to different semantic constraints. Causal pairs, for instance, explicitly require the modification to concern causal information:

```
1  if category == "causale":
2      return (
3          "Given an Italian caption (C) dealing "
4          "with causal relations between events "
5          "and its foil (F), your task is to assess "
6          "whether F is a valid causal foil of C.\n\n"
7
8          "To be valid, F must modify the cause, "
9          "the effect, or the causal relation "
10         "expressed in C while preserving the "
11         "remaining relevant content as much "
12         "as possible.\n"
13
14         "If F is a valid causal foil, output "
15         "'correct'. Otherwise output "
16         "'not correct'.\n\n"
17
18         f"C: {caption}\n"
19         f"F: {foil}\n"
20         "Output:"
21     )
```

The temporal prompt follows the same structure but constrains the acceptable difference to event ordering or other temporal properties:

```
1  if category == "temporale":
2      return (
3          "To be valid, F must modify the temporal "
4          "information expressed in C, such as "
```

```

5         "the ordering of events, before/after "
6         "relations, the moment at which an event "
7         "occurs, or another temporally relevant "
8         "relation.\n"
9     )

```

For spatial instances, the prompt is deliberately formulated to cover both stable locations and time-dependent spatial configurations:

```

1  if category == "spaziale":
2      return (
3          "The spatial category considered here "
4          "includes both locations and spatial "
5          "relations, including cases in which "
6          "spatial information refers to a specific "
7          "moment or phase of the video.\n\n"
8
9          "To be valid, F must modify spatial "
10         "information expressed in C, such as "
11         "the location, position, direction, or "
12         "spatial relation of a person, object, "
13         "or event.\n"
14     )

```

The model is constrained to produce only one of two labels, correct or not correct. Outputs are normalized before being stored, to remove minor formatting variations:

```

1  def normalize_structural_label(text):
2      value = clean_text(text).lower()
3
4      if value.startswith("output:"):
5          value = value[len("output:"):].strip()
6
7      value = value.rstrip(".").strip()

```

```

8
9     if value == "not correct":
10         return "not correct"
11
12     if value == "correct":
13         return "correct"
14
15     return "invalid"

```

The structural result is saved for every caption–foil pair in the `structural_eval` field. Incremental serialization is performed after each processed instance, so that completed evaluations survive an interruption:

```

1 df.at[index, "structural_eval"] = result
2
3 df.to_csv(
4     STRUCTURAL_OUTPUT,
5     index=False,
6 )

```

The second stage applies only to pairs that pass structural validation. Caption and foil are mapped respectively to the NLI premise S1 and hypothesis S2:

```

1 prompt = (
2     "Your task is to determine the natural "
3     "language inference (NLI) relationship "
4     "between S1 and S2.\n\n"
5
6     "The possible labels are:\n"
7     "- Entailment: S2 logically follows "
8     "from S1.\n"
9     "- Contradiction: S2 contradicts S1.\n"
10    "- Neutral: S1 and S2 are related, "
11    "but S2 neither follows from nor "

```

```

12     "contradicts S1.\n\n"
13
14     f"S1: {caption}\n"
15     f"S2: {foil}\n"
16     "Output:"
17 )

```

The resulting output is normalized to one of the three expected NLI labels:

```

1  if value == "contradiction":
2      return "Contradiction"
3
4  if value == "neutral":
5      return "Neutral"
6
7  if value == "entailment":
8      return "Entailment"

```

The two validation signals are then combined into an explicit status variable, with automatic acceptance limited to pairs that are structurally valid and classified as contradictions:

```

1  def get_validation_status(
2      structural_eval,
3      nli_eval,
4  ):
5      if structural_eval != "correct":
6          return "FAIL_STRUCTURAL"
7
8      if nli_eval == "Contradiction":
9          return "PASS"
10
11     if nli_eval == "Neutral":
12         return "REVIEW"
13

```

```

14     if nli_eval == "Entailment":
15         return "REVIEW_HIGH_PRIORITY"
16
17     return "ERROR"

```

Pairs that fail the structural check do not proceed automatically to semantic acceptance and are marked as FAIL_STRUCTURAL. Structurally valid pairs classified as Neutral are assigned REVIEW, while Entailment produces REVIEW_HIGH_PRIORITY. This hierarchy keeps uncertain model judgements from being treated as definitive dataset errors. At the end of the procedure, all non-passing cases are extracted into a dedicated manual-review file:

```

1  review_mask = df[
2      "validation_status"
3  ].isin(
4      [
5          "REVIEW",
6          "REVIEW_HIGH_PRIORITY",
7          "FAIL_STRUCTURAL",
8          "ERROR",
9      ]
10 )
11
12 review_df = df.loc[
13     review_mask
14 ].copy()
15
16 review_df.to_csv(
17     OUTPUT_DIR /
18     "caption_foil_manual_review.csv",
19     index=False,
20 )

```

This file serves as the interface between automatic and human validation. Each flagged instance is manually inspected by comparing its caption, foil, question, and semantic category. Cases where the automatic validator is overly conservative but the contrast is in fact appropriate are retained unchanged. When inspection instead confirms that the foil modifies the wrong semantic property, remains compatible with the caption, introduces unrelated information, or otherwise fails to provide the intended minimal contrast, the foil is corrected by hand. Corrected instances are then re-evaluated through the same structural and NLI checks before being incorporated into the final dataset. The complete validation output retains the original metadata alongside `structural_eval`, `nli_eval`, and `validation_status`. This preserves full traceability from each final VSV item back to its generated caption, foil, question, annotator response, automatic validation decisions, and any manual correction.

3. Study Design and Experimental Setup

3.1 Overview of the Evaluated Models

Qwen Family

The Qwen models considered in this work share an omnimodal architecture designed to process heterogeneous information within a unified computational framework. Both Qwen2.5-Omni and Qwen3-Omni accept textual, visual, and acoustic inputs, with modality-specific encoders transforming raw signals into representations compatible with a shared language-model backbone. Multimodal integration takes place after perceptual encoding, and particular attention is paid throughout the family to the temporal coordination of audio and video [24, 25].

Qwen2.5-Omni-3B

Qwen2.5-Omni introduces the *Thinker-Talker* architecture, separating multimodal understanding from speech generation. The *Thinker* processes the input modalities and produces textual representations and responses, while the *Talker* takes the Thinker’s internal representations and generates speech from them. Perceptual understanding and spoken generation thus remain functionally distinct components of an otherwise end-to-end pipeline [24]. Text is processed through the Qwen tokenizer. Audio is resampled to 16 kHz and converted into a 128-channel mel-spectrogram, with the encoder producing representations at a temporal resolution of roughly 40 ms. Visual information is handled by a Vision Transformer derived from Qwen2.5-VL and shared across image and video inputs; frames are sampled dynamically to preserve relevant visual content while matching the temporal granularity of the acoustic stream [24]. The modality-specific encoders are limited to perceptual extraction: their output is passed to the *Thinker*, where multimodal information is processed within the shared Transformer. Both audio and visual encoders additionally support block-wise processing, allowing long multimodal sequences to be handled incre-

mentally [24]. Multimodal alignment relies on *Time-aligned Multimodal Rotary Position Embedding* (TMRoPE), which represents positional information along temporal, vertical, and horizontal dimensions. Text is assigned identical positional values across all three; images retain a constant temporal coordinate while preserving spatial position along height and width; audio receives temporal positions tracking its progression through the acoustic stream [24]. Video combines both axes: each frame keeps its two-dimensional spatial layout while carrying a temporal coordinate tied to its position in the clip, on a scale coordinated with the acoustic representation. Audio and video can thus be placed within a common temporal reference system, which is what allows the *Thinker* to relate acoustic and visual evidence occurring at corresponding moments. Encoded representations from both streams are arranged within a single multimodal sequence and processed jointly by the language model’s attention mechanism, so integration occurs before response generation rather than through a late combination of independent modality-specific predictions [24]. The *Talker* is a dual-track autoregressive Transformer decoder, conditioned on the high-dimensional representations it receives from the *Thinker*; the resulting discrete speech units are converted into a waveform through a sliding-window diffusion Transformer [24]. The technical report describes this organization at the level of the Qwen2.5-Omni family as a whole, without a fusion mechanism specific to the 3B variant [24].

Qwen3-Omni-30B-A3B-Instruct

Qwen3-Omni keeps the Thinker-Talker organization of its predecessor while redesigning several components. It handles text, images, audio, and video, and supports both textual and spoken generation; in the 30B-A3B configuration, both the *Thinker* and the *Talker* adopt a Mixture-of-Experts architecture, increasing overall capacity while limiting the parameters activated at each inference step [25]. The *Thinker* still carries out multimodal understanding and text generation by combining representations from dedicated perceptual encoders. The *Talker*, however, is more loosely coupled to it than in Qwen2.5-Omni: rather than depending directly on the Thinker’s high-level textual hidden states, it conditions speech generation on textual context together with high-dimensional visual and acoustic representations [25]. Audio is processed by the *Audio Transformer* (AuT). Raw

audio is resampled to 16 kHz and represented as a 128-channel mel-spectrogram; convolutional downsampling brings the resulting representations to roughly 12.5 Hz, or about 80 ms per acoustic frame. AuT uses attention windows of different sizes to support both global audio understanding and block-wise processing [25]. Visual input is processed by an encoder derived from Qwen3-VL and initialized from SigLIP2-So400m, shared across images and video; as in Qwen2.5-Omni, frames are sampled dynamically to preserve visual content while matching the temporal granularity of the audio stream [25]. TMRoPE is retained for multimodal alignment, again along temporal, vertical, and horizontal dimensions: text keeps identical coordinates across all three, images hold a constant temporal coordinate while preserving spatial position, and audio now receives absolute temporal identifiers at roughly 80 ms resolution [25]. Video frames are assigned temporal identifiers taken from their actual timestamps, on the same scale used for audio, so acoustic and visual representations share an absolute temporal coordinate system while spatial information is still carried separately through height and width [25]. This replaces the fixed temporal chunking of Qwen2.5-Omni with alignment based on absolute temporal identifiers, producing a finer correspondence between acoustic and visual events and removing the dependence on a predefined chunk size [25]. The aligned representations are then processed by the Mixture-of-Experts *Thinker*, so multimodal integration occurs within the shared Transformer computation after perceptual encoding and temporal alignment; the pre-training strategy reinforces this from the outset, drawing on unimodal and cross-modal data spanning image-text, audio-text, video-text, video-audio, and video-audio-text combinations [25]. Speech generation is likewise redesigned. The *Talker* predicts a multi-codebook acoustic representation based on residual vector quantization: the first codebook is generated autoregressively, while a Multi-Token Prediction module estimates the remaining residual codebooks. The final waveform is produced by *Code2Wav*, a causal convolutional renderer built for incremental synthesis [25].

Gemma Family

The Gemma models considered in this work implement multimodality through two distinct strategies. Gemma 3n E4B relies on dedicated perceptual encoders that convert

acoustic and visual signals into representations processed by the language-model backbone. Gemma 4 12B instead adopts a unified, encoder-free architecture, mapping raw perceptual information directly into the decoder’s embedding space through lightweight projection modules [18, 6].

Gemma 3n E4B IT

Gemma 3n accepts text, images, audio, and video and generates textual responses; separate perceptual encoders convert acoustic and visual signals into token representations for the central language model [18]. The E4B configuration is built on *MatFormer* (Matryoshka Transformer), which nests a smaller E2B computational path within the larger E4B structure; both are trained jointly, allowing the model to operate under different computational budgets without requiring separate weights [18]. Gemma 3n also introduces *Per-Layer Embeddings*, storing a substantial portion of the embedding parameters outside accelerator memory. As a result, E4B keeps only around four billion core parameters resident on the accelerator despite a larger overall parameter count, a gain in computational efficiency rather than a change to how modalities are integrated [18]. Audio is processed by an encoder based on the Universal Speech Model, producing roughly one acoustic token every 160 ms (about six tokens per second) that feed directly into the language model [18]. Visual input is handled by MobileNet-V5-300M, supporting resolutions of 256×256 , 512×512 , and 768×768 pixels across both image and video content [18]. A *Multi-Scale Fusion VLM adapter* processes the resulting visual representations before they reach the language model; its role is limited to combining features at different scales and does not amount to a general audiovisual fusion mechanism [18]. Multimodal integration remains centered on the language-model backbone: audio and visual encoders produce modality-specific representations that become available to the Transformer during generation, and the available documentation does not describe a shared temporal alignment mechanism between acoustic and visual signals comparable to TMRoPE in the Qwen models [18]. Gemma 3n further employs *KV Cache Sharing*, letting keys and values from intermediate attention layers be reused by subsequent ones, which reduces redundant computation during the prefill stage, a benefit that is particularly relevant for the long sequences produced

by audio and video input [18].

Gemma 4 12B IT

Gemma 4 is a natively multimodal family built on a decoder-only Transformer, and the 12B configuration departs from the rest of the family architecturally. Rather than relying on perceptual encoders, it is trained from scratch under a unified encoder-free paradigm, with visual and acoustic input mapped directly into the language model’s embedding space through lightweight projection modules [6]. This sets it apart from the other Gemma 4 configurations: the smaller E2B and E4B models retain distinct vision and audio encoders, and even the larger variants keep dedicated visual processing. Gemma 4 12B alone drops both the vision and the audio encoder, replacing them with a direct projection of raw perceptual information [6]. For visual input, the model operates on RGB patches of $48 \times 48 \times 3$ values. A large linear projection of roughly 35 million parameters takes the place of the Vision Transformer used elsewhere in the family, mapping the resulting representations directly into the language model’s embedding space [6]. Spatial information is preserved by adding two-dimensional coordinate-based positional embeddings to the projected patches, followed by a final LayerNorm before the visual representations enter the backbone, enough to keep spatial structure available without a separate visual encoder [6]. The audio pathway follows the same principle. Raw audio, sampled at 16 kHz, is divided into 40 ms segments of 640 values each; the USM-based Conformer encoder used by the other Gemma 4 variants is removed entirely, and each 640-dimensional segment is projected directly into the language model’s embedding space [6]. No additional positional encoding is introduced for these acoustic representations, since their sequential order already preserves the temporal structure of the audio stream. Visual and acoustic information thus enter the same embedding space through modality-specific linear projections rather than through independent high-capacity encoders [6]. Multimodal integration accordingly takes place directly inside the decoder-only Transformer: the projection modules establish representational compatibility between raw perceptual input and the language model’s embedding space, and the Transformer itself carries out the subsequent interaction between perceptual and linguistic information, narrowing the separation

between perceptual encoding and language processing found in more conventional multi-modal systems [6]. The 12B model is dense and combines local and global self-attention, with five local attention blocks per global block. Local layers use RoPE, while global layers use p -RoPE with $p = 0.25$; the architecture also reuses keys as values in global attention layers, reducing the memory footprint of the global KV

3.2 Preliminary Analysis as a Diagnostic Tool for Multi-modal Assessment

The preliminary analysis aims to determine which information relevant to the subsequent resolution of *MAIA-AV* questions can be autonomously recovered by each model from the visual content of the video alone. The analysis is accordingly conducted in a *question-independent* setting: the model receives solely the video input, with no access to the questions, corresponding answers, audio streams, or transcriptions. The output is a structured and temporally grounded representation of the visual content, which is only later compared with the informational requirements of individual questions. This separation distinguishes the model’s capacity to extract visual information from its ability to retrieve relevant evidence when guided by a linguistic query. While the preliminary analysis does not directly measure performance on the VSV task, its purpose is to characterize what the model can represent prior to the introduction of a specific question-answering objective. Isolating these two abilities addresses a key limitation of evaluations based solely on final accuracy. A correct prediction does not guarantee that the model has correctly identified relevant entities, reconstructed events, or established the relations necessary to support its output. Conversely, an incorrect answer fails to reveal which stage of the process caused the failure. The preliminary analysis thus introduces an intermediate diagnostic layer between the raw video content and the final task outcome. Methodologically, this approach relates to *VILMA* [10], specifically regarding its distinction between proficiency tests and main tests. In *VILMA*, proficiency tests evaluate fundamental capabilities deemed necessary for resolving complex temporal phenomena via a zero-shot caption–foil paradigm. While our method does not replicate this specific evaluation protocol, it adopts

the same underlying principle: complex performance must be interpreted relative to the prerequisite capabilities that underpin it. In *MAIA-AV*, these prerequisites are examined through a structured representation of the visual content. For each video, four multimodal and omnimodal models, subsequently evaluated on the VSV task, independently analyze temporally localized video segments. Outputs are initially kept separate to enable cross-representation comparison without prematurely merging predictions. Inter-model agreement is interpreted as an indicator of **prediction stability** rather than absolute ground truth, since multiple models may converge on an identical yet unsupported interpretation. The representation is constructed progressively across three levels: **entities**, **events**, and **relations between events**. This hierarchy separates directly observable perceptual information from information that requires temporal integration or higher-level inference. The **entity level** describes elements identified within the visual content and their observable properties, detailing the objects, individuals, and other scene components detected by the model together with their corresponding temporal intervals of presence. The **event level** captures actions, interactions, and state changes tied to specific temporal intervals. Because overlapping video segments may yield duplicate descriptions of the same event, extracted outputs are consolidated to eliminate redundancy while preserving genuinely recurrent events. Temporal anchoring maintains a direct link between each event and the video segment providing its supporting evidence. The **relation level** organizes extracted events according to temporal and inferential dependencies. Temporal relations encompass configurations such as *before*, *after*, *during*, and *overlaps*, supporting event reordering and the tracking of changes across the video timeline. Causal relations are treated independently of temporal succession: simple temporal precedence does not by itself establish a cause-and-effect relationship. The analysis accordingly differentiates causal information directly supported by visual evidence from relations that require further inference over events, states, or intentions, keeping observable content separate from information introduced by the model’s own inferential process. The distinction between representational levels matters because questions assigned to the same nominal category can demand substantially different types and volumes of evidence. Identifying an object’s location within a single frame, for instance, requires fundamentally different operations than determin-

ing its position following a sequence of movements. Similarly, recognizing an isolated event differs from reconstructing a sequence of multiple events or inferring causal links between them. To capture these distinctions systematically, questions are analyzed along three operational dimensions: **spatial**, **temporal**, and **causal**. A unified three-level scale is applied across each dimension. Level 0 denotes cases where the required information can be directly retrieved. Level 1 involves integrating an explicit relation or a limited number of evidence units. Level 2 demands broader reconstruction across multiple events, temporal phases, state changes, or implicit relations. These criteria are summarized in Table 3.1. This classification reflects the informational requirements of the questions rather than their linguistic complexity, accounting for properties such as the temporal distribution of evidence, the number of relevant evidence units, the presence of inter-event relations, and whether the required information is explicit or implicit. Once constructed independently, the visual representation can be compared against the question requirements to determine whether the entities, events, and relations necessary to answer a question are already captured within the vision-only representation. The evaluation therefore does not depend directly on the model’s capacity to generate the correct answer. Instead, it assesses whether the supporting information is present within its visual representation. This comparison is particularly relevant for *MAIA-AV*, since it establishes a visual baseline prior to introducing additional modalities. Prior findings [4] showed that, in the original *MAIA* setting, incorporating acoustic information yielded no systematic performance gains, while including transcriptions led to a slight yet statistically robust decline in accuracy, consistent with a dataset originally designed around visual content rather than information explicitly conveyed via audio. The preliminary analysis thus establishes *what can be seen* prior to evaluating the contribution of *what can be heard*. It clarifies whether the visual channel inherently contains the required evidence and highlights categories where visual information alone proves insufficient, so that the incremental contributions of audio and transcriptions can be assessed against an explicit visual baseline rather than solely through shifts in downstream task accuracy.




Figure 3.1: Conceptual schema of the preliminary analysis.

Overall, the preliminary analysis transforms *MAIA-AV* from a purely result-driven evaluation into a diagnostic, competence-oriented framework. Final performance can thus be analyzed across specific sub-capabilities, including entity recognition, event identification, temporal integration, spatial reconstruction, and causal inference, disentangling evidence availability failures from errors in downstream selection or integration during the VSV task.

Category	Level 0	Level 1	Level 2
Spatial (S)	Direct retrieval of an entity’s location or spatial property from a single frame.	Explicit spatial relation or event-conditioned localization involving at most two evidence units.	Spatial reconstruction across multiple phases, including movement, relocation, or state change.
Temporal (T)	Recognition of a single event or state without temporal comparison.	Temporal localization or pairwise ordering between two explicit events.	Multi-event temporal reconstruction, including duration, repetition, onset, or state evolution.
Causal (C)	Retrieval of an explicitly stated cause without inference.	Explicit cause-and-effect relation involving at most two evidence units.	Implicit or multi-event causal reconstruction requiring contextual or intentional inference.

Table 3.1: Complexity levels for spatial, temporal, and causal questions.

Video Preprocessing: Temporal Segmentation and Frame Sampling

Prior to the entity extraction phase, each video is preprocessed by dividing it into four-second temporal windows with a stride of three seconds. Consecutive spans share one second of visual content. This overlap reduces the risk that an action or state change occurring near the boundary between two segments is only partially observed, while maintaining shared context across consecutive observations without altering the chronological order of the video. For each segment, the `start_time` and `end_time` are preserved, keeping each segment associated with its original temporal position in the video. This information is maintained throughout all subsequent stages of the pipeline, allowing observations to be reconnected with their exact position in the video. The windowing process serves exclusively to provide temporal indexing and structure visual evidence, leaving content and semantic interpretation to subsequent phases. Redundant information is deliberately retained at this stage to prioritize continuity of evidence, with redundancy resolution deferred to the output consolidation step. The preprocessing phase ultimately yields an ordered sequence of temporally localized visual units that serve as the shared foundation for all subsequent stages of analysis.

Implementation Details The preprocessing stage is implemented in `video_preprocessing.py` to convert each source video into a set of temporally indexed visual representations. The script first retrieves core metadata via `ffprobe`, including duration, frame rate, spatial resolution, codec, and audio stream presence. Although audio metadata is recorded, the acoustic signal is excluded from visual representation construction. Frames are extracted using OpenCV by seeking directly to target temporal positions, and each is saved as a JPEG image with its original timestamp embedded in both the filename and the JSON manifest, preserving temporal provenance independently of downstream processing. A core output of the script is the dense visual stream, sampled by default at 4 FPS:

```
1 parser.add_argument("--dense-fps", type=float, default=4.0)
```

Extracted frames are saved in the `dense_frames` directory for each video. Simultaneously,

the script generates 32 uniform global frames and performs shot-boundary detection via PySceneDetect. These boundaries define variable-length segments with a target duration of 4 seconds, ranging from 2 to 6 seconds, and a default overlap of 0.5 seconds, extracting eight frames per segment. The preprocessing output is serialized into `metadata.json`, `frames.json`, `shots.json`, and `segments.json`. For the subsequent preliminary analysis, semantic extraction scripts operate directly on the dense stream rather than shot-based segments. Utility functions locate each video's `dense_frames` directory, parse timestamps from filenames, and sort frames chronologically, so that the model receives temporally ordered visual observations instead of isolated or semantically segmented images. The inference temporal windows are configured as follows:

```
1 parser.add_argument("--window-seconds", type=float, default=4.0)
2 parser.add_argument("--stride-seconds", type=float, default=3.0)
```

A 4-second window with a 3-second stride yields a 1-second overlap between consecutive observations. The pipeline selects dense frames within these nominal boundaries, and if frame counts exceed the allowed maximum, frames are sampled uniformly across the window. At 4 FPS, a 4-second interval contains 16 frames, remaining well below the default 32-frame ceiling. Duplicate consecutive windows are automatically omitted. Each temporal unit is encapsulated in a `TemporalWindow`, mapping a unique segment ID to its start time, end time, and frame paths. Recorded boundaries reflect the timestamps of the first and last selected frames rather than nominal window limits, with millisecond precision maintained to ensure exact alignment with the source video. The transition to semantic extraction relies strictly on *ordered dense frames*. For each window, the model receives selected RGB images alongside a prompt specifying the video ID, temporal interval, and frame-to-timestamp map, with no questions, answers, transcriptions, or acoustic signals provided. Temporal segmentation thus serves purely as an indexing mechanism for visual evidence, deferring semantic interpretation to the subsequent extraction stage.

Semantic Representation Extraction

Building upon the defined temporal windows, the pipeline proceeds with semantic representation extraction. This stage generates a structured semantic description for each window, processing every window independently. The extraction procedure is repeated across all four evaluated models. Although the models employ distinct architectures and inference modalities, each receives identical inputs and is constrained to produce a standardized representation adhering to a common schema, as shown in Table 3.2. Consequently, performance differences can be attributed directly to model capabilities rather than annotation formatting. For each window, the representation spans seven semantic categories:

entities : Operating at the most fundamental descriptive level, each entity receives a local identifier, a concise label, visually observable attributes, and its role within the scene when discernible. The representation tracks the temporal occurrence of the entity and the specific frames providing visual evidence. These local identifiers enable consistent references across other descriptive categories within the same segment.

actions : Actions represent visually observable occurrences, specifying the performing agent and target object whenever feasible.

events : Events denote broader semantic units linked to one or more participants across a given temporal span. This distinction separates the detection of an isolated action from the broader event context in which it occurs.

spatial_relations : Spatial configurations between entities or scene elements are defined using a subject-relation-object structure. Permissible relations include configurations such as *in front of*, *behind*, *above*, *below*, *inside*, *outside*, *left of*, *right of*, *near*, and *on*. These spatial descriptors are temporally localized, meaning a spatial configuration observed in one segment is not assumed to hold throughout the video.

state_changes : State changes capture observable transformations affecting an entity. By contrasting a prior state with a subsequent state, this structure represents temporal variations within a segment. This level is essential when answering a question requires tracking changes in an entity’s properties, position, or condition over time.

temporal_relations : The model describes temporal links such as *before*, *after*, *overlaps*, *during*, *starts*, *finishes*, and *simultaneous*. These relations, however, remain strictly localized to the current window: global video temporal structure is not reconstructed at this stage, nor are event correspondences across different segments systematically resolved.

causal_hypotheses : Causal candidate relations are explicitly separated from directly observed facts. The framework requires causal hypotheses to be flagged as inferred evidence, preventing temporal precedence from being automatically interpreted as a cause-and-effect relationship, and keeping a clear boundary between observed evidence and model inference.

The distinction between observed and inferred evidence is preserved through the *evidence_type* field, accompanied by a confidence score between 0 and 1. Each element also includes an *evidence_frames* list identifying supporting frames. During output normalization, frame references are retained only if they fall within the processed window, and generated temporal intervals are similarly bounded by segment limits, keeping semantic representations strictly aligned with their supporting visual evidence. Categories may remain empty if a model detects insufficient evidence within a segment. At the segment level, the pipeline preserves *segment_id*, *start_time*, *end_time*, and *input_frames*. At this initial stage, a unified global video description is not yet constructed. Because consecutive windows overlap, identical entities, actions, or events may be extracted repeatedly across adjacent segments. This redundancy is intentionally preserved during extraction to maintain a direct mapping between each prediction and its supporting local evidence, leaving entity reconciliation and global temporal aggregation to subsequent pipeline stages.

Field	Information represented
entities	Identified entities, observable attributes, roles, and temporal occurrence.
actions	Local actions, their actors, and involved objects.
events	Temporally extended local events and their participants.
spatial_relations	Spatial configurations between entities or scene elements.
state_changes	Observable transitions between an entity’s previous and subsequent states.
temporal_relations	Temporal relations between events identified within the same local context.
causal_hypotheses	Candidate causal relations explicitly marked as inferred evidence.

Table 3.2: Semantic information extracted from each temporal window.

Implementation Details Entity and semantic analysis is implemented through a unified extraction pipeline shared across all evaluated models. Model-specific wrapper scripts load the corresponding architectures and format selected input frames, while temporal framing, prompt generation, schema definitions, validation logic, and serialization are centralized in `utils.py`, so that every model is evaluated against an identical semantic protocol regardless of underlying architectural differences. The common semantic space is defined across seven categories:

```

1 SEMANTIC_CATEGORIES = (
2     "entities",
3     "actions",
4     "events",
5     "spatial_relations",
6     "state_changes",
7     "temporal_relations",
8     "causal_hypotheses",
9 )

```

The categories illustrated in Table 3.2 specify the structured information required from each temporal window. For each temporal window, the pipeline constructs a standardized

extraction prompt via `build_semantic_prompt`, which supplies the video ID, window boundaries, and an explicit frame-to-timestamp map alongside the category definitions:

```
1 frame_map = "\n".join(
2     f'- frame {index}: "{path.name}", '
3     f'timestamp={get_timestamp(path):.3f}s'
4     for index, path in enumerate(
5         window.frame_paths,
6         start=1,
7     )
8 )
```

Rather than eliciting free-form descriptions, the prompt¹ enforces strict visual grounding: every extracted element must reference supporting frames, temporal boundaries, a confidence score, and an `evidence_type`. Directly visible details are classified as `observed`, whereas inferential claims are labeled as `inferred`, and the prompt explicitly instructs the model not to treat simple temporal succession as causality:

```
1     """Non inventare dettagli non visibili.
2     Usa evidence_type="observed"
3     per fatti direttamente visibili.
4     Usa evidence_type="inferred"
5     soltanto per inferenze inevitabili
6     o ipotesi causali.
7     Per ogni elemento indica confidence
8     tra 0 e 1 e i nomi esatti dei frame
9     che lo supportano.
10    Non confondere una successione
11    temporale con una relazione causale.
```

¹English translation: "Do not invent details that are not visible. Use `evidence_type="observed"` for directly visible facts. Use `evidence_type="inferred"` only for unavoidable inferences or causal hypotheses. For each element, indicate confidence between 0 and 1 and the exact frame names that support it. Do not confuse temporal succession with a causal relationship. If a category contains no elements, return an empty list. Return exclusively a valid JSON object."

```
12     Se una categoria non contiene elementi,  
13     restituisci una lista vuota.  
14     Restituisci esclusivamente  
15     un oggetto JSON valido."""
```

Models are constrained to output structured JSON objects containing standardized meta-data, including `start_time`, `end_time`, `evidence_frames`, `evidence_type`, and `confidence`.

Entity identifiers generated within a window can be cross-referenced across actions, events, and relations in the same segment. The main execution loop is encapsulated in `process_video_with_inf`. After dense frames are partitioned into temporal windows, each window is processed sequentially:

```
1  for window_index, window in enumerate(  
2      windows,  
3      start=1,  
4  ):  
5      prompt = build_semantic_prompt(  
6          video_directory.name,  
7          window,  
8      )  
9      response = inferencer(  
10         window.frame_paths,  
11         prompt,  
12     )  
13  
14     parsed = parse_model_json(response)  
15  
16     segment = normalize_semantic_output(  
17         parsed,  
18         window,  
19     )  
20  
21     segments.append(segment)
```

The inferencer component is the sole model-dependent step in this loop, accepting frame paths and prompts and returning textual output. All downstream parsing and validation operations remain identical across architectures. To manage output formatting inconsistencies from generative models, the parsing stage strips Markdown wrappers and extracts the primary JSON object. If standard decoding fails, the pipeline attempts recovery using `json_repair` before marking the response invalid:

```
1  try:
2      parsed = json.loads(candidate)
3  except json.JSONDecodeError as first_error:
4      try:
5          from json_repair import repair_json
6          parsed = json.loads(
7              repair_json(candidate)
8          )
9      except Exception as repair_error:
10         raise SemanticOutputError(
11             f"JSON non valido: {first_error}"
12         ) from repair_error
```

Structural constraints are subsequently enforced by `normalize_semantic_output`, which clips temporal boundaries to window limits, restricts confidence scores to the interval $[0, 1]$, and filters out invalid frame references. Evidence types are validated against observed, inferred, or uncertain, with causal hypotheses automatically reclassified as inferred if initially marked as observed:

```
1  confidence = _safe_float(
2      item.get("confidence"),
3      0.0,
4  )
5  item["confidence"] = round(
6      max(0.0, min(1.0, confidence)),
```



```

7     4,
8 )
9 evidence_type = str(
10     item.get(
11         "evidence_type",
12         "observed",
13     )
14 ).lower()
15 if evidence_type not in {
16     "observed",
17     "inferred",
18     "uncertain",
19 }:
20     evidence_type = "uncertain"
21 if (
22     category == "causal_hypotheses"
23     and evidence_type == "observed"
24 ):
25     evidence_type = "inferred"

```

Frame-level grounding validation retains only filenames present in the input temporal window:

```

1 valid_evidence = set(
2     window.evidence_frames
3 )
4 item["evidence_frames"] = [
5     str(name)
6     for name in evidence_frames
7     if str(name) in valid_evidence
8 ]

```

Normalized segments are enriched with temporal metadata from the preprocessing stage:

```

1 normalized = {
2     "segment_id": window.segment_id,
3     "start_time": window.start_time,
4     "end_time": window.end_time,
5     "input_frames": window.evidence_frames,
6 }

```

Outputs are saved for each video in a `_semantic.json` file alongside model metadata and window configuration settings. Window failures are logged in an `errors` field containing segment identifiers, temporal bounds, frame lists, and error descriptions, preserving optional raw responses for diagnostic review. This phase ultimately transforms temporally indexed visual evidence into structured local representations. Cross-window deduplication is deliberately omitted at this step to preserve local grounding, leaving element consolidation and full video reconstruction to subsequent analysis stages.

Event Analysis

Spatial Relations

Spatial information captures entity configurations within the scene and tracks positional changes across localized visual observations. Unlike temporal relations, spatial features are extracted directly during the semantic representation phase, while visual windows are still accessible. Spatial relations encode explicit configurations between entities, or between an entity and a scene element. These include positional, proximity, containment, and orientation relations such as *left of*, *right of*, *above*, *below*, *inside*, *outside*, *in front of*, *behind*, *near*, and *on*. Because these descriptors are tied to localized segments, a single entity pair may exhibit distinct spatial arrangements across different temporal intervals. Displacements and state transitions also drive spatial variation. A positional change is represented by combining the spatial relations observed at different timestamps with the corresponding action and any recorded `state_changes`, distinguishing static scene layouts from dynamic rearrangements. Event-dependent localization further evaluates spatial configurations relative to an event interval, before, during, or after, isolating how a prior

action alters an entity’s final position. Methodologically, co-occurrence does not imply a spatial relation. The presence of two entities within the same temporal segment provides insufficient grounds to assume proximity, orientation, containment, or contact: a relation is retained only when explicitly captured in the semantic analysis. Spatial evidence currently remains bound to these local semantic descriptions. The event extraction stage draws on `spatial_relations` for event reconstruction but does not construct a global spatial graph.

Temporal Relations

Temporal information is formalized during event extraction, where local semantic descriptions are consolidated into a chronological timeline covering the full video. Each consolidated event retains a `start_time` and an `end_time` defining its temporal boundary and, implicitly, its duration. Relations between events are modeled explicitly through pairs of `event_id` references. Supported relation types are *before*, *after*, *overlaps*, *during*, and *simultaneous*, each linking two existing events in the consolidated representation together with an associated confidence score. *Before* and *after* express temporal precedence and sequence, *during* denotes temporal containment, *overlaps* indicates partial interval sharing, and *simultaneous* marks concurrent occurrences. This structure captures chronological succession, repetition, containment, and overlap across the video timeline while keeping distinct occurrences separate. Generation order within the JSON output does not itself constitute temporal evidence: dependencies derive strictly from segment timestamps and relational metadata rather than from the model’s output sequence. The temporal schema adopts the following structure:

```
1 "temporal_relations": [  
2   {  
3     "first_event": "E0001",  
4     "relation": "before|after|overlaps|during|simultaneous",  
5     "second_event": "E0002",  
6     "confidence": 0.0  
7   }
```

This consolidated event structure also serves as input for the subsequent causal analysis, though chronological organization is kept distinct from causal inference: temporal precedence does not by itself establish a cause-and-effect relationship.

Implementation Details The event analysis is implemented as a second-stage transformation over the structured semantic representations produced by the previous phase. No visual input is processed again at this stage. For each model, the corresponding semantic output is loaded and used as the only source of information for event reconstruction. The event analyzer can reorganize and consolidate previously extracted information, but it cannot introduce evidence obtained from a new inspection of the video, preserving the separation between perceptual extraction and subsequent structural reasoning. The input therefore consists of the temporally indexed semantic segments generated for the same model. Each segment retains its identifier and temporal boundaries together with the semantic categories extracted from the corresponding visual window. The information made available to the event analysis includes entities, actions, local events, spatial relations, state changes, temporal relations, and causal hypotheses:

```
1 keys = (  
2   "segment_id",  
3   "start_time",  
4   "end_time",  
5   "entities",  
6   "actions",  
7   "events",  
8   "spatial_relations",  
9   "state_changes",  
10  "temporal_relations",  
11  "causal_hypotheses",  
12 )
```

These fields are considered jointly because an event may not be represented exclusively

in the local events field. An action or an observable state transition, for instance, may provide relevant information for reconstructing an event even when the local semantic description does not encode it as an independent event. The event analysis therefore operates over the complete local semantic representation rather than simply concatenating the previously extracted event lists. Before inference, the semantic segments are arranged according to their temporal position and serialized into a structured textual input. The model is explicitly instructed to treat this representation as a closed source of evidence, with the common prompt constraining the reconstruction process through a set of rules designed to prevent additional perceptual or commonsense information from being introduced:

```
1
2 * usa esclusivamente le informazioni presenti
3   nell'analisi semantica
4 * non inventare azioni, oggetti, intenzioni,
5   emozioni o cause mancanti
6 * unisci soltanto i duplicati causati dalla
7   sovrapposizione delle finestre
8 * mantieni separati eventi realmente distinti
9   o ripetuti
10 * ricava i tempi soltanto dai segmenti forniti
```

A central function of this phase is the consolidation of local descriptions produced by overlapping temporal windows. Since the semantic extraction uses 4-second windows with a 3-second stride, adjacent segments share part of the same visual evidence, and the same physical event can consequently appear in multiple semantic outputs. During event reconstruction, descriptions referring to the same occurrence are merged into a single event, an operation intended to remove redundancy introduced by the preprocessing strategy rather than repetitions that genuinely occur in the video. The prompt explicitly enforces this distinction. Events with similar descriptions are not automatically treated as duplicates: two occurrences are merged only when their temporal context and supporting segments indicate that they represent the same event observed through overlapping windows. Repeated actions occurring at different moments, by contrast, remain represented as distinct events,

so that repetition is preserved as a meaningful temporal property of the video. The consolidated representation follows a common JSON schema. Each reconstructed event contains a unique identifier, a concise description, temporal boundaries, the participants involved, the semantic segments supporting the event, the type of evidence, and a confidence value:

```
1 {  
2   "event_id": "E0001",  
3   "description": "string",  
4   "start_time": 0.0,  
5   "end_time": 0.0,  
6   "participants": ["entity or participant"],  
7   "evidence_segments": ["segment_0001"],  
8   "evidence_type":  
9     "observed|inferred|uncertain",  
10  "confidence": 0.0  
11 }
```

The `event_id` field provides a stable reference for subsequent relational analyses. Events are assigned progressive identifiers such as E0001, E0002, and E0003, associated with the consolidated chronological representation rather than inherited from the local semantic outputs, so that events can be referred to consistently once multiple local descriptions have been merged. Temporal localization is inherited from the semantic analysis, since the event extractor does not estimate event boundaries by accessing the original frames. Instead, `start_time` and `end_time` must remain compatible with the temporal intervals of the semantic segments supporting the event. When the same occurrence is represented across multiple overlapping windows, its consolidated interval can reflect the temporal evidence distributed across those segments. The `evidence_segments` field preserves provenance after consolidation. An event extracted from a single semantic window may contain only one segment identifier, whereas an event observed across overlapping windows may be associated with several:

```
1 "evidence_segments": [  
2   "segment_0004",
```

```
3 "segment_0005"  
4 ]
```

Consolidation removes duplicated event descriptions without discarding the evidence from which they originated, keeping the resulting event representation traceable to the local observations produced in the previous phase, even though it is more compact than the semantic representation itself. Participants are derived from entities or participants already represented in the semantic analysis and are not generated as independent new entities at this stage, so that the event representation preserves semantic continuity between entity extraction and event reconstruction. The epistemic status of each event is represented through `evidence_type`. The analysis distinguishes three values:

```
1 "observed"  
2 "inferred"  
3 "uncertain"
```

Observed is used for events supported by information previously represented as directly visible. Inferred is reserved for information that already requires an inferential interpretation of the available semantic evidence. Uncertain represents cases in which the available evidence does not support a stronger classification. The event extraction stage is not intended to transform directly observed information into unsupported inference. Confidence is represented numerically in the $[0, 1]$ interval and subsequently normalized to ensure a consistent representation. Normalization also verifies the structural validity of the generated output before serialization. In addition to individual events, the same inference produces an initial global temporal organization. Temporal relations connect events through their newly assigned identifiers rather than through their textual descriptions:

```
1 "temporal_relations": [  
2 {  
3 "first_event": "E0001",  
4 "relation":  
5 "before|after|overlaps|during|simultaneous",  
6 "second_event": "E0002",
```

```

7  "confidence": 0.0
8  }
9  ]

```

The allowed relations distinguish precedence, succession, containment, partial temporal overlap, and simultaneity, providing an explicit representation of the temporal structure that emerges once local events have been consolidated. Temporal relations cannot be derived solely from the order in which events are written in the generated JSON. The prompt explicitly requires temporal organization to depend on the intervals and relations available in the semantic representation: the order in which the text is generated is not treated as temporal evidence, and temporal succession is likewise kept distinct from causal interpretation. The model output is parsed as structured JSON rather than retained as free-form text. Responses are cleaned of possible Markdown delimiters, and the JSON object is extracted before decoding. The same robust parsing strategy used in the semantic stage allows malformed but recoverable outputs to be repaired when possible:

```

1  text = (
2  text.replace("`json", "")
3      .replace("`", "")
4  .strip()
5  )
6
7  start = text.find("{")
8  end = text.rfind("}")
9
10 if start < 0 or end < start:
11     raise ValueError(
12         "La risposta non contiene un oggetto JSON"
13     )
14
15 payload = json.loads(
16     text[start:end + 1]

```


This phase primarily changes the granularity of the representation. The semantic analysis describes evidence locally within individual overlapping windows, while the event analysis reorganizes those observations into a coherent event-based representation of the complete video, removing redundancies caused by temporal overlap, preserving genuinely repeated occurrences, assigning stable identifiers, maintaining temporal and evidential provenance, and establishing the first explicit relations among events, all without any new inspection of the visual input.

Causal Inference

Causal inference operates exclusively on the consolidated event representation, independently of perceptual extraction. This phase receives no visual frames or raw multimodal input, relying solely on previously reconstructed events and their established temporal relations. Isolating causal reasoning in this way prevents the recovery of omitted visual details or the insertion of unobserved events. Before processing, the event representation is pruned to its essential fields: `event_id`, `description`, `temporal interval`, `participants`, `supporting semantic segments`, `evidence_type`, and `confidence`. The analysis considers four types of causal dependency:

- `causes`: one event directly produces the outcome represented by another.
- `enables`: one event establishes the conditions required for another to occur.
- `motivates`: the available evidence supports an intentional or motivational link between events.
- `prevents`: one event inhibits or counteracts another.

Each entry links a `cause_event` to an `effect_event`, optionally supplemented by `supporting_events` where additional events are needed to substantiate the inference. No new events are introduced during this phase. The resulting structure follows this schema:

```

1  "causal_relations": [
2    {
3      "cause_event": "E0001",
4      "relation": "causes|enables|motivates|prevents",
5      "effect_event": "E0002",
6      "evidence_type": "direct|inferred|uncertain",
7      "supporting_events": ["E0003"],
8      "explanation": "string",
9      "confidence": 0.0
10   }
11 ]

```

The epistemic status of each relation is indicated by `evidence_type`. Dependencies grounded in explicit physical interactions or state transitions are labeled `direct`. Connections that require contextual interpretation are labeled `inferred`, while partially supported links receive `uncertain`. Motivational dependencies require explicit support within the event representation rather than speculative plausibility. Temporal succession does not imply causation: a chronological link such as *before* or *after* remains insufficient on its own to establish a causal relation. Dependencies are introduced only when the semantic content explicitly supports causation, enabling, motivation, or prevention. Keeping these stages separate ensures that temporal ordering is never mistaken for causal evidence, preserving the distinction between extracted facts and inferential reasoning.

Implementation Details The causal inference stage is a text-only processing step applied to the consolidated event representation produced in the preceding phase. For each video, the pipeline loads the corresponding `_events.json` file generated by the model. No raw video frames, audio signals, transcriptions, or downstream text prompts are supplied, so the causal analyzer operates exclusively on a closed representation of previously extracted visual information. Prior to inference, the event file is pruned by the `compact` function, which retains only the fields required for causal reasoning:

```

1 def compact(payload):
2     return {
3         "id_video": payload.get("id_video"),
4         "events": [
5             {
6                 "event_id": event.get("event_id"),
7                 "description": event.get("description"),
8                 "start_time": event.get("start_time"),
9                 "end_time": event.get("end_time"),
10                "participants": event.get("participants", []),
11                "evidence_segments": event.get("evidence_segments", [])
12            },
13            {
14                "evidence_type": event.get("evidence_type"),
15                "confidence": event.get("confidence"),
16            }
17            for event in payload.get("events", [])
18        ],
19        "temporal_relations": payload.get("temporal_relations", []),
20    }

```

This filtering isolates causal reasoning from the original perceptual features, requiring the model to derive dependencies strictly from the consolidated events and their temporal layout. The allowable relation types and evidence categories are bounded explicitly in code:

```

1 ALLOWED_RELATIONS = {
2     "causes",
3     "enables",
4     "motivates",
5     "prevents",
6 }
7 ALLOWED_EVIDENCE = {
8     "direct",

```

```

9     "inferred",
10    "uncertain",
11  }

```

These definitions separate direct cause-and-effect links (causes) from enabling conditions (enables), intentional drivers (motivates), and inhibitory actions (prevents). Inference is guided by a standardized prompt that instructs the model to rely solely on the provided events and temporal relations, explicitly forbidding it from treating simple chronological succession as causal evidence:

```

1
2  * usa esclusivamente gli eventi e le relazioni
3  temporali presenti nell'input
4  * non usare domande, caption, foil o altre
5  informazioni esterne
6  * la semplice successione temporale NON implica causalita'
7  * crea una relazione soltanto quando il contenuto
8  degli eventi supporta una dipendenza causa-effetto,
9  una condizione abilitante, una motivazione oppure
10 una prevenzione
11 * cause_event ed effect_event devono essere ID
12 di eventi presenti nell'input
13 * non inventare eventi, oggetti, intenzioni o cause
14 non ricavabili dalla rappresentazione fornita
15 * se non esiste evidenza sufficiente, restituisci
16 causal_relations come lista vuota

```

Within this schema, relationships are classified by epistemic confidence: direct for explicit physical or state-change interactions, inferred for contextual or intentional deductions, and uncertain for weakly supported cues. Output formatting enforces the same structured JSON schema used elsewhere in the pipeline:

```

1  {
2  "causal_relations": [

```

```

3  {
4      "cause_event": "E0001",
5      "relation": "causes|enables|motivates|prevents",
6      "effect_event": "E0002",
7      "evidence_type": "direct|inferred|uncertain",
8      "supporting_events": ["E0003"],
9      "explanation": "string",
10     "confidence": 0.0
11 }
12 ]
13 }

```

The fields `cause_event` and `effect_event` reference existing event IDs, while `supporting_events` optionally logs contextually relevant secondary events. Output parsing strips Markdown delimiters and decodes the outer JSON object. If formatting errors occur, `json_repair` attempts structural recovery, and if parsing fails entirely, a secondary generation call requests strictly formatted JSON without altering the input prompt:

```

1  text = text.replace("json", "").replace("\n", "").strip()
2  start = text.find("{")
3  end = text.rfind("}")
4  candidate = text[start : end + 1]
5  try:
6      return json.loads(candidate)
7  except json.JSONDecodeError:
8      from json_repair import repair_json
9      return json.loads(repair_json(candidate))

```

A normalization step then applies a sequence of structural checks before the relations are accepted. It first collects the set of valid event identifiers parsed from the input:

```

1  valid_ids = {
2      event.get("event_id")
3      for event in payload.get("events", [])

```

```

4     if event.get("event_id")
5 }

```

Candidate relations that reference an unknown event, or that link an event to itself, are discarded:

```

1 if (
2     cause not in valid_ids
3     or effect not in valid_ids
4     or cause == effect
5 ):
6     continue

```

Relations using an unrecognized type are dropped, while evidence labels outside the allowed set default to inferred:

```

1 if rel not in ALLOWED_RELATIONS:
2     continue
3 if evidence not in ALLOWED_EVIDENCE:
4     evidence = "inferred"

```

Relations sharing the same cause, relation, and effect are then deduplicated, keeping only the first occurrence of each triplet:

```

1 key = (cause, rel, effect)
2 if key in seen:
3     continue
4 seen.add(key)

```

Supporting-event references are filtered in the same way, excluding invalid identifiers as well as the cause and effect events themselves:

```

1 supporting = [
2     event_id
3     for event_id in relation.get("supporting_events", [])
4     if event_id in valid_ids and event_id not in {cause, effect}

```

```
5 ]
```

Finally, confidence scores are clipped to the $[0, 1]$ interval:

```
1 try:
2     confidence = float(relation.get("confidence", 0.0))
3 except (TypeError, ValueError):
4     confidence = 0.0
5 confidence = max(0.0, min(1.0, confidence))
```

Final outputs are serialized per video into `_causal.json` files containing metadata, model parameters, and the resulting relations. Unresolvable errors yield empty relation arrays rather than speculative fallbacks.

NLI-Based Visual Proficiency Assessment

The final stage of the preliminary analysis integrates the information produced at the semantic, event, temporal, spatial, and causal levels into a consolidated representation of the visual evidence recovered by each model. The procedure operates exclusively on the outputs generated during the preceding phases, preserving the separation between visual analysis and question-dependent evaluation without performing any further inspection of the video. The integrated representation combines complementary information produced at different levels of the pipeline. Entity descriptions, spatial relations, and state changes originate from the semantic analysis. Consolidated events provide a global event-level representation together with their participants, temporal boundaries, supporting segments, evidence type, and confidence. Temporal relations connect these events through their identifiers, while causal relations encode the dependencies established during causal inference. These components are treated not as independent observations, but as progressively structured views of the same visual evidence. Finalization does not amount to a simple concatenation of all intermediate outputs. Redundancies introduced by overlapping temporal windows have already been resolved during event consolidation, where multiple descriptions of the same occurrence are merged while genuinely repeated events remain distinct, and duplicated or structurally invalid causal relations have already been

removed during causal normalization. The epistemic status assigned during the previous stages is preserved as well: information classified as observed, inferred, or uncertain remains distinguishable in the consolidated representation, together with the corresponding confidence values. Uncertain information is not automatically promoted to positive evidence, nor are conflicting descriptions resolved through an additional interpretation introduced at this stage. The final representation thus captures the entities, events, spatial configurations, temporal organization, and causal dependencies recovered by the model, without expanding the evidence beyond what the visual analysis actually produced. Crucially, **finalization does not constitute an answer to any VSV question**. Its purpose is to fix a representation of what the model has recovered from the visual modality before the linguistic content of the evaluation item is introduced. Questions, captions, and foils therefore play no role in constructing the representation, and once it has been finalized, its content remains fixed throughout the evaluation of the corresponding question: the question may provide context for interpreting the linguistic alternatives, but it cannot trigger the extraction of new entities, events, relations, or visual evidence. The proficiency assessment itself is formulated as a Natural Language Inference task. The finalized visual representation acts as the **premise**, while the caption and the corresponding foil are evaluated independently as alternative hypotheses. GPT-4o is employed exclusively as an NLI judge, not to answer the original question, generate a new description of the video, or select directly between the two alternatives. Its role is restricted to determining the logical relation between the evidence contained in the finalized representation and each hypothesis. For both caption and foil, the judge assigns one of three labels: entailment, contradiction, or neutral. Entailment indicates that the hypothesis is supported by the available visual representation. Contradiction indicates that it is incompatible with the represented evidence. Neutral is used when the available information is insufficient to establish either relation, so that missing evidence is not interpreted as contradiction and absence of information is not treated as negative evidence. The caption and foil are judged separately. A caption–foil pair receives PASS only when the finalized representation entails the caption and contradicts the foil. Every other combination produces NOT PASS. This rule adopts a deliberately conservative criterion. Entailment of the caption alone is

insufficient, since the same representation must also provide enough evidence to reject the minimally contrastive foil. Contradiction of the foil, in turn, cannot compensate for a caption classified as `neutral`. A model is therefore considered proficient on a given pair only when its visual representation supports both sides of the required distinction. The procedure is repeated independently for the four caption–foil pairs associated with each question. Question-level proficiency requires consistency across the complete set: a question receives `PASS` only when all four pairs satisfy the pair-level criterion, and is classified as `NOT PASS` if at least one pair fails. This strict aggregation reduces the influence of individual linguistic formulations and requires the recovered visual evidence to remain sufficient across multiple, minimally different realizations of the same underlying item. For each pair, the procedure records the classification assigned to the caption, the classification assigned to the foil, and the resulting binary decision, which are then aggregated into the final proficiency label associated with the question. The resulting measure is accordingly distinct from conventional question-answering accuracy: rather than evaluating whether the model can generate or select the correct answer after being exposed to the question, it evaluates whether the information independently recovered from the visual channel is sufficiently complete and discriminative to support the correct statement while rejecting its minimally contrastive alternative, connecting question-independent visual analysis to the competence-oriented evaluation objective of the preliminary analysis.

Implementation Details The final evaluation is implemented in `final_evaluation.py`, which operates on the finalized representations produced for the four evaluated models. These representations are stored as JSON files in `data/preliminar_analysis/final_results` and constitute the fixed visual premises used throughout the proficiency assessment. The script expects one finalization file for each model and verifies their presence before starting the evaluation:

```
1 EXPECTED_MODELS = (  
2     "gemma",  
3     "gemma-4-12B",  
4     "qwen",
```

```

5  "qwen-30B",
6  )
7
8  FINAL_DIR = Path(
9  "data/preliminar_analysis/final_results"
10 )
11
12 final_files = [
13 FINAL_DIR / f"{model}.json"
14 for model in EXPECTED_MODELS
15 ]

```

Finalized representations are loaded independently for each model and indexed by normalized video identifiers. The loading procedure accepts different top-level JSON organizations, including lists of video representations, dictionaries containing videos or results, individual video objects, and dictionaries directly indexed by video ID, allowing the subsequent NLI procedure to retrieve the complete representation associated with a video without modifying its content. Question metadata and caption–foil pairs are loaded separately from `question_classification.csv` and `caption_foil.csv`. Video identifiers and question categories are normalized before the two sources are joined on the video and principal question dimension. The question is therefore introduced only after the finalized visual representation has been loaded, and is retained as contextual information for the NLI judge rather than incorporated into the premise itself. The NLI evaluation uses `gpt-4o-2024-08-06` with deterministic generation:

```

1  MODEL = "gpt-4o-2024-08-06"
2  risposta = client.chat.completions.parse(
3  model=MODEL,
4  temperature=0,
5  messages=[
6  {
7  "role": "system",

```

```

8  "content":
9  f"{PROMPT}\n\n{REGOLE[categoria]}",
10 },
11 {
12  "role": "user",
13  "content": contenuto,
14 },
15 ],
16 response_format=RispostaNLI,
17 )

```

Each request contains three distinct fields: the question, the finalized representation used as the premise, and one hypothesis. The caption and foil are never evaluated jointly. Instead, two independent calls are performed using the same premise and question context:

```

1  nli_caption = classifica(
2  premessa,
3  row["domanda"],
4  row["categoria"],
5  row["caption"],
6  )
7
8  nli_foil = classifica(
9  premessa,
10 row["domanda"],
11 row["categoria"],
12 row["foil"],
13 )

```

The system prompt explicitly restricts the judge to the information contained in the premise. External knowledge, completion of missing information, and plausibility-based reasoning are excluded, and incomplete, weak, or ambiguous evidence must be classified as *neutro*. Additional constraints are introduced according to the question category: spatial relations

cannot be inferred from simple co-occurrence, temporal relations cannot be inferred from the textual order of events, and causal relations cannot be established from temporal succession alone. The model response is constrained through a Pydantic schema rather than parsed from unrestricted natural-language output:

```
1 class RispostaNLI(BaseModel):
2     etichetta: Literal[
3         "implicazione",
4         "contraddizione",
5         "neutro",
6     ]
```

This design restricts each inference to one of the three admissible NLI labels, preventing the judge from generating explanations or free-form answers. A missing or structurally invalid classification raises an error rather than being converted into an evaluation decision. Pair-level proficiency is subsequently computed through a deterministic rule implemented outside GPT-4o:

```
1 def valuta_coppia(caption, foil):
2     return (
3         "PASS"
4         if caption == "implicazione"
5         and foil == "contraddizione"
6         else "NOT PASS"
7     )
```

GPT-4o's role is thus limited to assigning the two NLI labels: it does not directly decide whether a proficiency test is passed. The binary outcome is derived programmatically, receiving PASS only when the finalized representation entails the caption and contradicts the corresponding foil. All other label combinations result in NOT PASS. Results are written incrementally to a model-specific checkpoint in the `nli` subdirectory. Before evaluating a pair, the script checks whether its identifier has already been processed, allowing interrupted executions to resume without repeating completed API calls. Each stored row con-

tains the model, pair identifier, video, category, question, caption, foil, the two NLI classifications, and the resulting pair-level outcome. Once all model-specific evaluations have been completed, pair-level records are concatenated into `nli_risultati.csv`. Question-level proficiency is then computed by grouping results according to model, video, category, and question:

```
1 df = risultati.groupby(  
2 [  
3 "modello",  
4 "video_id",  
5 "categoria",  
6 "domanda",  
7 ]  
8 ).agg(  
9 totale_coppie=("esito", "size"),  
10 coppie_pass=("pass", "sum"),  
11 ).reset_index()  
12  
13 df["proficiency"] = [  
14 "PASS" if passate == totale else "NOT PASS"  
15 for passate, totale in zip(  
16 df["coppie_pass"],  
17 df["totale_coppie"],  
18 )  
19 ]
```

The aggregation adopts the same conservative principle used at pair level: proficiency is assigned only when every caption–foil pair associated with the question passes. Since the experimental data contain four pairs for each question, this corresponds to the required 4/4 criterion, and a single NOT PASS is sufficient to classify the question as NOT PASS. The script additionally stores the number of evaluated pairs, the number of successful pairs, and their ratio in `question_proficiency.csv`. The implementation keeps the

responsibilities of the two components distinct: the finalized JSON representation remains unchanged and functions solely as the visual premise, GPT-4o performs constrained NLI classification over each caption and foil, and the final proficiency decision is produced by deterministic aggregation rules rather than by an additional generative judgment.

3.3 Core Reasoning Assessment Framework

Experimental Logic and Hypotheses

The experimental design of MAIA-AV extends beyond a simple comparative evaluation of model accuracy. Its primary objective is to analyze the role played by individual modalities in the decision-making process and verify whether correct answers are effectively derived from input evidence. Experiments are structured around a **Modality-Ablation Design**. By holding the test item constant while systematically varying the input modalities, it is possible to observe how predictions shift when a specific information source is added or removed. This experimental control enables **paired comparisons**, allowing behavioral variance to be attributed more precisely to the selected modality configuration. Furthermore, this setup establishes a clear distinction between **Modality Sufficiency** and **Modality Necessity**. A modality is defined as sufficient when it independently contains all the information required to solve the task. However, sufficiency does not imply that the modality is uniquely useful, nor does it preclude the utility of additional sources. Conversely, supplementing a sufficient modality with another does not guarantee performance gains; the added input may be redundant, introduce irrelevant noise, or, in the worst-case scenario, interfere with a previously correct prediction. Our experimental setup thus shifts the focus from assessing overall model capability toward evaluating how effectively a system integrates available modalities and whether distinct modality preferences emerge. From this perspective arise the concepts of **Modality Dominance**, **Modality Gain**, **Repair**, **Damage**, and **Cross-Modal Interaction**, which are formally defined as evaluation metrics in subsequent sections. This methodological framework is rooted in the principle of **controlled information removal** established in *Lost in Transcription* [4]. In both

paradigms, the underlying task remains constant while the input information accessible to the model is systematically varied. This setup allows us to directly evaluate whether an added modality compensates for the absence of another or provides an advantage when partial evidence is already available. In *Lost in Transcription*, analysis was primarily oriented toward assessing acoustic contributions within a benchmark built around visual evidence. The main focus was verifying whether the acoustic/textual signal could compensate for absent visual evidence and how this contribution shifted as visual evidence is added. With MAIA-AV, we extend this paradigm through a more articulated manipulation of input modalities. The framework comprises the following unimodal and baseline conditions: complete absence of input, isolation of the acoustic signal, isolation of visual information, and isolation of audio transcriptions. These initial conditions establish how much relevant information can be retrieved from a single modality alone. We then introduce two multimodal configurations: video with native audio (the full audiovisual clip) and video with audio transcriptions (muted video paired with text transcriptions). In these multimodal settings, the focus extends beyond whether audio merely assists a primarily visual decision. Instead, the goal is to comprehend how the availability and representation of information across different modalities reshape the model’s overall reasoning. This setup transforms the core question from a simplistic inquiry—*does audio help?*—into a broader investigation of how distinct information sources jointly contribute to multimodal reasoning.

Visual Statement Verification

All experimental conditions adopt an identical task formulation: **Visual Statement Verification** (VSV). Each decision is evaluated on a pair consisting of a caption (a correct description) and a foil (a minimally contrastive false description). These two alternatives are presented to the model as a binary choice, and the system must return exclusively 0 or 1 depending on its selection. This constrained formulation minimizes confounding factors associated with open-ended text generation, enabling a direct comparison across models and input configurations. Consequently, the exact same experimental unit can be observed across different multimodal configurations without introducing variations in prompt struc-

ture or evaluation criteria. This task structure allows us to distinguish two analytical levels. The first corresponds to single caption-foil choices, designated as **Item-Level Decisions**. The second is **Aggregate Accuracy**, which evaluates sets of items linked to the same underlying conceptual category to assess whether the model maintains consistent behavior across different formulations. These metrics are formally defined in subsequent sections.

Domanda	
Dove cammina la bambina alla fine della scena?	
Caption	Foil
Alla fine della scena, la bambina cammina su un tappeto .	Alla fine della scena, la bambina cammina su un sentiero .
Question	
Where does the little girl walk at the end of the scene?	
Caption	Foil
At the end of the scene, the little girl walks on a carpet .	At the end of the scene, the little girl walks on a path .

Table 3.3: Example of a minimally contrastive caption–foil pair, reported in Italian with its English translation.

Experimental Setup

The protocol is operationalized through a standardized pipeline and a set of adapters tailored to specific families of models. This pipeline centralizes various tasks, including the preparation of items, the definition of experimental conditions, the construction of prompts, the management of inputs, and the storage of outputs. Inputs consist of caption-foil pairs, original videos, versions without audio tracks, separate audio files, and precomputed transcripts. The pre-established order of alternatives in the dataset is preserved, and the target retains essential information for verifying decision accuracy. For each item, the

raw response generated by the model is also preserved. From this response, the binary label required for the task is subsequently extracted. This mechanism enables the differentiation between valid responses and unparsable outputs, facilitating the documentation of potential execution errors. A significant distinction concerns modes not supported by the architecture. A lack of support for a specific type of input is not treated as an incorrect response. Instead, such a condition is marked as unavailable, and the model is only excluded from comparisons requiring that specific configuration. Therefore, the protocol should be interpreted through a model-by-mode matrix, where the availability of conditions may vary across different architectures.

Cross-Architecture Input Standardization The experimental protocol remains uniform across model families, with input delivery adapted to specific multimodal interfaces, including varying processors and formats for Qwen and Gemma. These adaptations ensure identical information payload across different native architectures without altering the experimental design.

Video Processing The preparation of visual content adheres to a standardized sampling policy. In the experiments, video content is encapsulated using 32 frames evenly distributed across its duration. This strategy ensures that the amount of visual information provided to various architectures remains comparable, while also accommodating the specific preprocessing requirements of the respective processors. In visual-only conditions, a version of the video devoid of audio components is used. In audiovisual configurations, visual information is coupled with the acoustic tracks specified by the protocol. Thus, sampling represents a shared experimental choice. Internal differences among the visual encoders of the models fall outside this segment of the protocol and should be addressed separately in the architectural descriptions.

Audio and Transcription Processing The acoustic component is prepared in a format compatible with the frontend of models capable of processing native audio. When needed, the track is extracted from the video and normalized according to the requirements of the specific adapter. Implementations requiring a 16 kHz signal apply resampling before in-

ference. The transcription follows a different logic. It is not produced during the main experiment and does not represent an intermediate output of the model under evaluation. Instead, it serves as a precomputed input, kept consistent across the conditions using it and directly integrated into the textual context. This distinction is methodologically significant. The transcription primarily preserves recognized linguistic content, while the audio signal retains additional acoustic properties. Comparing these two representations allows for the study not just of the presence of information, but also of the form in which it becomes accessible to the model.

Inference Protocol Inference parameters are configured to minimize stochastic variability. Because the task demands a binary classification choice rather than open-ended text generation, deterministic decoding is enforced using greedy search (e.g., setting `temperature = 0` or `do_sample=False`, depending on the model interface). Maximum output length is strictly constrained. Restricting generation to a minimal token budget mitigates non-compliant formatting and ensures consistency across models and experimental conditions. Inference is executed using `bfloat16` precision where supported. Models are deployed across execution backends, including `vLLM` and Hugging Face Transformers. These infrastructure settings preserve the underlying experimental logic, while hardware-dependent parameters, such as GPU allocation and framework-specific workarounds, are detailed in the technical appendix for reproducibility.

Zero-input Condition

The no-input experimental condition completely removes perceptual information, providing the model exclusively with the textual caption–foil pair required to complete the task. No video frames, audio signals, or transcriptions of any kind are supplied. This configuration does not assess audiovisual comprehension; rather, its function is strictly diagnostic. It allows us to establish the extent to which decisions can be derived from the linguistic plausibility of the options, structural regularities within the dataset, or systematic parametric biases of the model. Performance significantly above chance level in this condition serves as a strong indicator of potential linguistic shortcuts, operating independently of

genuine content understanding.

Implementation Details The *No Input* experiment was implemented through a model-independent evaluation pipeline shared across all architectures considered in the study. The purpose of this condition is to remove every source of perceptual information while preserving the binary decision structure of the Visual Statement Verification task. For each item, the model is presented exclusively with the two minimally contrastive textual alternatives. The original question, video frames, audio track, and corresponding transcription are entirely omitted. The experimental logic is centralized in a common evaluation module, independent from model-specific inference adapters. Each adapter is responsible only for formatting textual requests according to the native interface of its target architecture. This separation ensures that identical items, instructions, target labels, output parsing procedures, and evaluation criteria are applied consistently across all models. The system instruction is shared across the entire experimental pipeline:

```
Sei un assistente binario per un task di Visual Statement Verification.  
Rispondi esclusivamente con 0 oppure 1.2
```

The prompt used for the *No Input* condition is likewise identical for all models:

```
Basandoti esclusivamente sulle due alternative testuali, scegli.3  
  
0: {answer1}  
1: {answer2}  
  
Rispondi esclusivamente con 0 oppure 1.4
```

Here, `answer1` and `answer2` correspond to the two members of the caption–foil pair. Their numerical ordering is independent of their semantic status; dataset target labels indicate whether the correct caption occupies position 0 or 1. The model is thus asked to perform a binary selection between competing linguistic realizations rather than gener-

²English translation: *You are a binary assistant for a Visual Statement Verification task. Respond exclusively with either 0 or 1.*

³English translation: *Based exclusively on the two textual alternatives, choose.*

⁴English translation: *Respond exclusively with either 0 or 1.*

ating open-ended text. At the implementation level, prompt construction is performed centrally before passing the request to any model adapter:

```
1 def prompt(row, mode, transcript=""):
2     return PROMPTS[mode].format(
3         answer1=row["answer1"],
4         answer2=row["answer2"],
5         transcript=transcript,
6     )
```

In *No Input* mode, no transcription is retrieved and no multimodal content is attached. Consequently, only the formatted textual prompt reaches the model. Although conversation formatting relies on each architecture's native chat template, this implementation detail does not alter the information available to the model. All models are constrained to produce binary outputs. To handle cases where generative models wrap predictions in extraneous text, responses are normalized via a regex parsing function that extracts the first isolated digit (0 or 1):

```
1 def prediction(text):
2     m = re.search(r"(?<!\d)[01](?!\\d)", str(text))
3     return int(m.group()) if m else None
```

If no valid binary label can be recovered, the prediction is flagged as invalid rather than being forcibly mapped to a candidate answer, preventing malformed outputs from skewing accuracy. The evaluation loop remains uniform across models. For every item, the shared pipeline builds the prompt, invokes the appropriate adapter, parses the response, and checks correctness against the ground-truth target:

```
1 raw = infer(
2     mode,
3     row,
4     prompt(row, mode, transcript),
5     paths(row),
6 )
```

```

7 pred = prediction(row)
8 result = (
9     "correct" if pred == int(row["target"])
10    else "wrong" if pred is not None
11    else "invalid"
12 )

```

Output records are serialized uniformly across models. Each entry stores item metadata, model ID, experimental condition, raw model output, parsed binary label, and correctness status:

```

1 FIELDS = [
2     "id", "video_id", "video_name", "pool_id",
3     "pool_item", "question_category", "question",
4     "caption", "foil", "answer1", "answer2",
5     "target", "model", "mode", "raw_model_output",
6     "predicted_label", "is_correct", "result", "error",
7 ]

```

Results are written to model-specific output files for the *No Input* condition. Existing item identifiers are cached prior to execution, allowing interrupted runs to resume seamlessly without recomputing completed instances. Outputs are flushed immediately after writing to prevent data loss. Inference is executed deterministically across all adapters by disabling stochastic sampling. Disabling sampling is particularly crucial in the *No Input* condition to eliminate output variance unrelated to the underlying linguistic representations. Overall, model-specific code is strictly limited to weight loading, chat template formatting, and generation execution. Prompt construction, decision formatting, output parsing, correctness evaluation, error handling, and serialization remain completely standardized. This ensures that performance variations in the *No Input* condition reflect differences in model capability rather than inconsistencies in the experimental protocol.

Audio-only Condition

In the *Audio Only* condition, the model exclusively receives the raw acoustic signal, with no visual frames or textual transcriptions provided. This setting allows us to identify which caption-foil pairs can be resolved based on spoken content, environmental sounds, background noise, music, or prosodic speech characteristics. Furthermore, comparing performance in this condition against other experimental configurations reveals whether acoustic information retains its independent utility when combined with visual input, or if its contribution is diminished or overshadowed by another dominant modality.

Implementation Details The *Audio-Only* condition was implemented as a controlled unimodal setting designed to isolate the contribution of the acoustic channel. For each VSV item, the model receives the raw audio associated with the corresponding video together with the two minimally contrastive textual alternatives. No visual frames, video representations, or audio-derived transcriptions are provided. The original question is retained solely as dataset metadata and is omitted from the inference prompt. Consequently, the model must discriminate between caption and foil by relying on acoustic evidence and the linguistic content of the two candidate statements. As in the *No Input* condition, experimental logic is centralized in a common evaluation module shared across all four models. Model-specific adapters are responsible only for converting audio and textual content into the representation expected by each architecture. This separation preserves the same prompt, candidate ordering, target labels, prediction parsing, and evaluation procedure across models, while allowing each architecture to process the acoustic signal through its native multimodal interface. The system instruction remains unchanged across experimental conditions, but for the *Audio-Only* condition, the following prompt is used:

```
Ascolta esclusivamente l'audio fornito.  
Scegli quale delle due descrizioni è corretta sulla base delle  
informazioni ricavabili esclusivamente dall'audio.  
  
0: {answer1}  
1: {answer2}
```

Rispondi esclusivamente con 0 oppure 1.

The two alternatives correspond to the caption–foil pair associated with the VSV item. Their numerical positioning does not encode semantic status, as the correct caption may occur in either position. The target field stored in the dataset determines whether the true statement corresponds to label 0 or 1. The model is thus required to perform a binary selection rather than produce an open-ended description. At the pipeline level, the acoustic resource associated with each item is resolved from its video identifier alongside other media paths:

```
1 def paths(row):
2     stem = Path(row["video_name"]).stem
3     return {
4         "audio": AUDIO_DIR / f"{stem}.mp3",
5         "mute": MUTE_DIR / f"{stem}.mp4",
6         "video": VIDEO_DIR / f"{stem}.mp4",
7     }
```

Only the acoustic component is selected by the model adapter when `only_audio` mode is active. The exact conversion of this resource into model input is architecture-dependent. Models accepting audio files directly receive the acoustic resource through their native processor, whereas architectures requiring raw waveforms decode the signal into the format expected by their inference backend. These operations strictly concern input formatting and do not introduce textual transcriptions or visual information.

Prompt construction remains entirely model-independent:

```
1 def prompt(row, mode, transcript=""):
2     return PROMPTS[mode].format(
3         answer1=row["answer1"],
4         answer2=row["answer2"],
5         transcript=transcript,
6     )
```

In the *Audio-Only* condition, the `transcript` argument remains empty. Thus, the

textual portion of the request contains only task instructions and the two competing statements. The acoustic stream constitutes the sole perceptual source available to the model.

The common evaluation loop passes the selected condition, item, formatted prompt, and resolved media paths to the corresponding model-specific inference adapter:

```
1 raw = infer(  
2     mode,  
3     row,  
4     prompt(row, mode, transcript),  
5     paths(row),  
6 )  
7  
8 pred = prediction(raw)
```

Regardless of internal generation interfaces, all models are required to return a binary decision. The generated text is normalized through the same shared parsing function used across remaining experimental conditions:

```
1 def prediction(text):  
2     m = re.search(r"(?!\\d)[01](?!\\d)", str(text))  
3     return int(m.group()) if m else None
```

The first isolated occurrence of either 0 or 1 is interpreted as the prediction. If neither label can be recovered, the output is classified as invalid instead of being forcibly mapped to candidate options. The parsed prediction is subsequently checked against the stored target:

```
1 result = (  
2     "correct" if pred == int(row["target"])  
3     else "wrong" if pred is not None  
4     else "invalid"  
5 )
```

Inference is executed deterministically across model adapters by disabling stochastic sampling. This ensures that variation between predictions is attributable to the information

available in the experimental condition rather than sampling variability.

The output structure is likewise shared with other experimental settings. For each item, the pipeline logs the model identifier, experimental condition, raw model generation, parsed binary prediction, target label, correctness status, and potential execution errors. Results are written to a dedicated `only_audio` directory, with a separate output file for each model. Completed item identifiers are detected prior to execution, allowing interrupted runs to resume seamlessly without recomputing previously processed instances.

This condition therefore differs from *No Input* solely in the availability of the acoustic signal. This controlled modification allows us to quantify whether speech, environmental sounds, music, or prosodic cues contain sufficient information to discriminate true captions from minimally contrastive foils, independently of visual evidence and audio-derived textual representations.

Video-only Condition

In the *Video Only* condition, the model receives exclusively the visual stream, with the audio track removed. This condition serves as the primary baseline to evaluate which decisions can be sustained purely by visual evidence. Furthermore, the visual condition plays a central role in paired comparisons: subsequently introducing audio or transcriptions enables us to measure the marginal contribution of these additional information sources relative to a baseline where visual context is already available.

Implementation Details. The *Video-Only* condition was implemented as a controlled visual setting designed to isolate the contribution of the visual modality. For each VSV item, the model receives the visual content of the corresponding video together with the two minimally contrastive textual alternatives. The acoustic track and the audio-derived transcription are excluded from the input. The original question is retained as dataset metadata but is not included in the inference prompt. The decision must therefore be based on the visual evidence available in the video and on the linguistic contrast between caption and foil. As in the other experimental conditions, the general evaluation procedure is centralized in a model-independent module. The same prompt structure, candidate

ordering, target labels, prediction parser, and output format are used for all four models. Model-specific adapters are responsible only for converting the visual input into the representation required by the corresponding architecture. The system instruction remains unchanged across experimental conditions: For the *Video-Only* condition, the model receives the following condition-specific prompt:

```
Osserva il video fornito.  
Scegli quale delle due descrizioni è corretta sulla base  
esclusivamente delle informazioni visive disponibili.  
  
0: {answer1}  
1: {answer2}  
  
Rispondi esclusivamente con 0 oppure 1.5
```

The two alternatives correspond to the caption–foil pair associated with the VSV item. Their numerical position does not encode whether an alternative is true or false. The correct caption can occupy either position and the target label stored in the dataset specifies the expected binary prediction. This prevents the numerical labels from carrying semantic information and preserves the same decision structure adopted in the other experimental conditions.

To guarantee that no acoustic information can reach the model, the *Video-Only* condition uses a dedicated muted version of each video. The common pipeline resolves the available media resources from the identifier of the corresponding item:

```
1 def paths(row):  
2     stem = Path(row["video_name"]).stem  
3     return {  
4         "audio": AUDIO_DIR / f"{stem}.mp3",  
5         "mute": MUTE_DIR / f"{stem}.mp4",  
6         "video": VIDEO_DIR / f"{stem}.mp4",  
7     }
```

⁵English translation: *Observe the provided video. Choose which of the two descriptions is correct based exclusively on the available visual information. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

For this condition, only `paths["mute"]` is passed to the visual processing stage. The use of a physically muted video rather than simply instructing the model to ignore audio guarantees that the acoustic channel is absent from the actual multimodal input.

The visual representation is standardized to a maximum of 32 frames sampled across the complete temporal extent of the video:

```
1 N_FRAMES = 32
```

Frames are distributed uniformly over the sequence rather than being extracted from a single local interval. Conceptually, their positions correspond to equally spaced indices between the beginning and the end of the video:

```
1 frame_indices = np.linspace(  
2     0,  
3     total_frames - 1,  
4     min(N_FRAMES, total_frames),  
5     dtype=int,  
6 )
```

When a video contains fewer than 32 frames, all available frames are retained. This sampling strategy provides a fixed visual budget while preserving information from the complete temporal development of the clip. The extracted frames are converted to RGB before being supplied to the multimodal processor. Depending on the native interface of the architecture, the resulting sequence is subsequently represented either as a video input or as an ordered sequence of image inputs. This implementation difference is confined to the model-specific adapter and does not alter the visual evidence available in the experiment.

Prompt construction remains independent from the model architecture:

```
1 def prompt(row, mode, transcript=""):  
2     return PROMPTS[mode].format(  
3         answer1=row["answer1"],  
4         answer2=row["answer2"],
```

```

5     transcript=transcript,
6 )

```

In the *Video-Only* condition, the transcript argument remains empty. Therefore, the textual component of the request contains only the task instruction and the two alternatives. No linguistic information derived from the audio track is introduced.

At inference time, the common evaluation loop combines the prompt with the media paths and delegates only the architecture-specific processing to the corresponding adapter:

```

1 raw = infer(
2     mode,
3     row,
4     prompt(row, mode, transcript),
5     paths(row),
6 )
7 pred = prediction(raw)

```

All models are required to return a binary output. The generated response is normalized through the same parser used in the remaining experimental conditions:

```

1 def prediction(text):
2     m = re.search(r"(?!\\d)[01](?!\\d)", str(text))
3     return int(m.group()) if m else None

```

The first isolated occurrence of either 0 or 1 is interpreted as the prediction. Outputs from which no valid label can be extracted are classified as invalid rather than being automatically assigned to one of the two alternatives. The resulting label is then compared with the target stored for the corresponding caption–foil pair:

```

1 result = (
2     "correct" if pred == int(row["target"])
3     else "wrong" if pred is not None
4     else "invalid"
5 )

```

Inference is performed deterministically, with stochastic sampling disabled across the model adapters. For each item, the pipeline stores the raw generation, the parsed binary prediction, the target label, the correctness status, the model identifier, and the experimental condition. Results are written to a dedicated `only_video` directory using a separate output file for each model. The *Video-Only* condition therefore preserves the same binary VSV decision protocol adopted throughout the experimental framework while restricting perceptual evidence to the visual channel. Its comparison with *No Input* makes it possible to quantify the contribution introduced by visual grounding, while comparison with the audio-visual conditions provides the reference point needed to determine whether additional modalities improve, preserve, or interfere with visually grounded predictions.

Transcription-only Condition

In the *Transcription Only* condition, the model receives exclusively the textual transcription derived from the audio track. This configuration isolates spoken linguistic content from non-verbal acoustic features, such as prosody, speaker characteristics, background noise, and music. Consequently, this setting evaluates performance driven strictly by explicit verbal information.

Implementation Details. The *Only-Transcription* condition was implemented as a controlled textual setting designed to isolate the contribution of the linguistic information derived from the audio channel. For each VSV item, the model receives the audio transcription together with the two minimally contrastive textual alternatives. No visual frames, video representation, or raw acoustic signal are provided. The original question is retained as dataset metadata but is not included in the inference prompt. The model must therefore discriminate between caption and foil using only the textual content extracted from the audio and the linguistic information contained in the two alternatives.

This condition differs from the *No Input* setting in a single controlled aspect. While the latter exposes the model only to the caption–foil pair, *Only-Transcription* additionally provides the textual representation of the audio track. It therefore makes it possible to determine whether speech-derived linguistic information contributes to the decision even

when direct perceptual access to the original video and acoustic signal is removed.

As in the other experimental conditions, the evaluation logic is centralized in a common module shared across the evaluated models. The system instruction, candidate ordering, target representation, output parser, and result serialization remain unchanged. Model-specific adapters are only responsible for applying the appropriate chat template and performing generation.

The common system instruction is:

```
Sei un assistente binario per un task di Visual Statement Verification.  
Rispondi esclusivamente con 0 oppure 1.
```

6

The condition-specific prompt is defined as follows:

```
Leggi esclusivamente la trascrizione dell'audio fornita.  
Scegli quale delle due descrizioni è corretta sulla base  
esclusivamente delle informazioni testuali disponibili nella  
trascrizione.  
  
Trascrizione dell'audio: {transcript}  
  
0: {answer1}  
1: {answer2}  
  
Rispondi esclusivamente con 0 oppure 1.7
```

The transcription is not generated during VSV inference. Instead, the previously generated transcriptions are stored in a dedicated CSV file and loaded before the evaluation begins. The common implementation maps each transcription to the corresponding video name:

⁶English translation: “You are a binary assistant for a Visual Statement Verification task. Respond exclusively with either 0 or 1.”

⁷English translation: *Read exclusively the provided audio transcription. Choose which of the two descriptions is correct based exclusively on the textual information available in the transcription. Audio transcription: {transcript}. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

```

1 TRANSCRIPTIONS = Path(
2     "data/input/transcription/transcriptions.csv"
3 )
4
5 def transcriptions():
6     with TRANSCRIPTIONS.open(
7         encoding="utf-8-sig",
8         newline=""
9     ) as f:
10         return {
11             r["video_name"]: r["transcription"]
12             for r in csv.DictReader(f, delimiter=";")
13         }

```

This mapping allows the same textual representation to be retrieved consistently for every VSV item associated with the same video. During the common evaluation loop, transcription retrieval is explicitly restricted to the two conditions in which this information is required:

```

1 transcript = (
2     transcripts.get(row["video_name"], "")
3     if mode in (
4         "only_transcription",
5         "transcript_video",
6     )
7     else ""
8 )

```

For *Only-Transcription*, the retrieved text is inserted directly into the condition-specific prompt. Prompt construction remains centralized and independent of the model architecture:

```

1 def prompt(row, mode, transcript=""):

```

```

2     return PROMPTS[mode].format(
3         answer1=row["answer1"],
4         answer2=row["answer2"],
5         transcript=transcript,
6     )

```

The two candidate statements correspond to the minimally contrastive caption–foil pair. Their numerical labels do not encode their semantic status. The true caption may occupy either position and the dataset target specifies whether the expected answer is 0 or 1. The transcription is therefore auxiliary evidence used to discriminate between the same alternatives employed in every other experimental condition.

A relevant implementation property of this setting is the complete absence of multi-modal objects from the inference request. At adapter level, `only_transcription` does not trigger the loading of either visual or acoustic resources. The general logic can be represented as follows:

```

1 content = []
2
3 if mode == "only_transcription":
4     pass
5
6 content.append({
7     "type": "text",
8     "text": prompt,
9 })

```

Consequently, the transcription is not processed through an audio encoder. It is tokenized as ordinary textual input together with the task instruction and the two candidate statements. This distinction is central to the experimental design. The condition evaluates access to the semantic content extracted from audio rather than access to the original acoustic modality itself. Prosodic information, speaker characteristics, environmental sounds, and other non-textual acoustic properties are therefore unavailable unless they

were explicitly represented in the generated transcription.

The same inference interface used by the other conditions is then called with the formatted prompt:

```
1 raw = infer(  
2     mode,  
3     row,  
4     prompt(row, mode, transcript),  
5     paths(row),  
6 )  
7  
8 pred = prediction(raw)
```

Although the common function also provides the media paths to the adapter, none of these resources is included in the model input for *Only-Transcription*. This preserves a uniform inference interface across conditions without altering the information available in the present setting.

All models are required to return a binary response. The generated output is processed through the shared parser:

```
1 def prediction(text):  
2     m = re.search(  
3         r"(?<!\d)[01](?!\\d)",  
4         str(text),  
5     )  
6     return int(m.group()) if m else None
```

The parser extracts the first isolated occurrence of 0 or 1. If no valid binary label is found, the prediction is marked as invalid rather than being assigned automatically to one of the alternatives. Valid predictions are compared with the target associated with the caption–foil pair:

```
1 result = (  
2     "correct"
```

```

3     if pred == int(row["target"])
4     else "wrong"
5     if pred is not None
6     else "invalid"
7 )

```

Generation is performed deterministically, with stochastic sampling disabled, and the maximum generated sequence is restricted to the few tokens required for the binary response. This minimizes generation variability and keeps the model output aligned with the discriminative nature of the VSV task. For each item, the common evaluation pipeline stores the raw model output, the parsed prediction, the expected target, the correctness status, the model identifier, the experimental condition, and possible execution errors. Results are written to a dedicated `only_transcription` directory, with a separate CSV file for each evaluated model. Previously completed item identifiers are detected before inference, which allows interrupted executions to resume without recomputing existing predictions. The *Only-Transcription* condition therefore provides a direct diagnostic of the contribution of textualized audio information. Its comparison with *No Input* measures the additional contribution introduced by the transcription itself. Its comparison with *Audio-Only* distinguishes the semantic information preserved by textualization from the information available in the original acoustic signal. Finally, comparison with conditions including visual evidence makes it possible to examine whether the same linguistic information remains useful when it is combined with the primary visual channel.

The Audio-Visual (AV) Condition

The Vision-Language (VL) Configuration

Metrics Selection and Aggregation Strategy

Linguistic and Structural Bias Assessment under Zero-Input

The *No Input* condition requires distinct analytical treatment compared to other performance metrics. Its primary purpose is to measure the extent to which the task can be

solved without perceptual stimuli, making it unsuitable for direct inclusion in standard performance comparisons. Instead, dedicated metrics are employed to quantify **linguistic and structural shortcuts**. The *No Input* condition is designed to quantify the extent to which a model can solve the VSV task without access to perceptual evidence. Since neither visual nor acoustic information is available, performance in this setting cannot be interpreted as evidence of multimodal understanding. Instead, the analysis focuses on whether the model can exploit linguistic regularities, structural properties of the alternatives, or systematic biases in the dataset.

Let N denote the number of evaluated items. For each item i , let $y_i \in \{0, 1\}$ be the target label and $\hat{y}_i \in \{0, 1\}$ the prediction produced by the model.

Raw Accuracy The first measure is *Raw Accuracy*, defined as the proportion of items for which the predicted label corresponds to the target:

$$\text{Acc}_{\text{NoInput}} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = y_i), \quad (3.1)$$

where $\mathbb{I}(\cdot)$ is the indicator function, equal to 1 when the condition is satisfied and 0 otherwise. Because each item requires a binary choice between two alternatives, a model selecting one of the two options uniformly at random has expected accuracy

$$\mathbb{E}[\text{Acc}_{\text{random}}] = \frac{1}{2} = 0.5. \quad (3.2)$$

An accuracy close to 0.5 is therefore compatible with behavior close to random guessing. Conversely, a substantially higher value indicates that the model can identify the target more frequently than expected by chance despite the absence of perceptual evidence. This suggests that some information contained in the linguistic form of the alternatives may be predictive of the correct answer. Raw Accuracy must not be interpreted as a measure of language comprehension. The model is not required to understand the underlying video event and does not have access to the evidence necessary to perform grounded reasoning. The metric only measures its ability to recover the target from the two textual alternatives available in the *No Input* condition.

Language-Only Advantage To express the deviation from the random baseline more directly, we compute the *Language-Only Advantage* (LOA). It is obtained by subtracting the expected binary chance level from the observed *No Input* accuracy:

$$\text{LOA} = \text{Acc}_{\text{NoInput}} - 0.5. \quad (3.3)$$

The transformation can be derived directly from the decomposition

$$\text{Acc}_{\text{NoInput}} = 0.5 + \text{LOA}. \quad (3.4)$$

Hence, $\text{LOA} = 0$ indicates performance exactly at the random baseline. Positive values indicate performance above chance:

$$\text{LOA} > 0 \implies \text{Acc}_{\text{NoInput}} > 0.5, \quad (3.5)$$

while negative values indicate performance below chance:

$$\text{LOA} < 0 \implies \text{Acc}_{\text{NoInput}} < 0.5. \quad (3.6)$$

The main advantage of this transformation is interpretability. Rather than reporting the complete accuracy value, LOA directly expresses the amount of performance exceeding the expected random level. It can therefore be used as a compact indicator of the potential contribution of linguistic or structural shortcuts. A positive LOA does not identify the nature of the shortcut. It only establishes that the textual alternatives contain sufficient regularities for the model to perform above chance. Additional diagnostics are therefore required to determine which properties may contribute to this effect.

Matthews Correlation Coefficient Accuracy alone can conceal degenerate decision strategies. For example, a model may develop a systematic preference for one output label and still obtain a non-negligible accuracy, particularly when the observed predictions are not evenly distributed across the two positions. For this reason, the analysis also includes the *Matthews Correlation Coefficient* (MCC). For the binary task, predictions can be summarized through the four entries of the confusion matrix: true positives (TP), true negatives (TN), false positives (FP), and false negatives

(FN). MCC is defined as

$$\text{MCC} = \frac{\text{TP} \cdot \text{TN} - \text{FP} \cdot \text{FN}}{\sqrt{(\text{TP} + \text{FP})(\text{TP} + \text{FN})(\text{TN} + \text{FP})(\text{TN} + \text{FN})}}. \quad (3.7)$$

The numerator compares concordant and discordant predictions. Correct positive and negative decisions contribute through $\text{TP} \cdot \text{TN}$, while the two types of error contribute through $\text{FP} \cdot \text{FN}$. The denominator normalizes this difference according to the marginal distributions of targets and predictions. MCC ranges between -1 and 1 :

$$-1 \leq \text{MCC} \leq 1. \quad (3.8)$$

A value close to 1 indicates strong agreement between predictions and targets. A value close to 0 indicates little or no association. A negative value indicates a systematic inverse association. MCC complements Raw Accuracy because it evaluates the overall correspondence between predictions and targets while taking all four components of the confusion matrix into account. It is therefore useful for distinguishing genuine target-sensitive performance from apparently successful behavior produced by strongly unbalanced prediction strategies.

Position Bias The *Position Bias* analysis directly measures whether the model systematically prefers one of the two candidate positions independently of their content. Let

$$p_0 = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = 0) \quad (3.9)$$

denote the empirical probability of selecting the first alternative, and

$$p_1 = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = 1) = 1 - p_0 \quad (3.10)$$

the probability of selecting the second alternative. Under perfectly balanced selection behavior,

$$p_0 = p_1 = 0.5. \quad (3.11)$$

We therefore define the magnitude of the positional deviation as

$$\text{PB} = |p_0 - 0.5|. \quad (3.12)$$

Because $p_1 = 1 - p_0$, the same quantity can equivalently be written as

$$\text{PB} = |p_1 - 0.5|. \quad (3.13)$$

The preferred position is defined as

$$\text{PreferredPosition} = \underset{j \in \{0,1\}}{\operatorname{argmax}} p_j. \quad (3.14)$$

A Position Bias equal to 0 indicates perfectly balanced selection frequencies. The maximum possible value is 0.5, obtained when the model always selects the same position:

$$0 \leq \text{PB} \leq 0.5. \quad (3.15)$$

This diagnostic is required because above-chance Raw Accuracy does not by itself demonstrate the exploitation of meaningful linguistic information. Part of the observed result may instead originate from a systematic preference for the first or second alternative. Position Bias therefore provides a direct measure of this potential confound.

Lexical Attraction The *Lexical Attraction* analysis evaluates whether the model preferentially selects the alternative that is lexically more similar to the original question. For each item i , let q_i denote the question and let a_{i0} and a_{i1} denote the two textual alternatives. Each text is represented as a TF-IDF vector. The cosine similarity between the question and alternative j is defined as

$$s_{ij} = \cos(\mathbf{v}_{q_i}, \mathbf{v}_{a_{ij}}) = \frac{\mathbf{v}_{q_i}^\top \mathbf{v}_{a_{ij}}}{\|\mathbf{v}_{q_i}\|_2 \|\mathbf{v}_{a_{ij}}\|_2}, \quad j \in \{0, 1\}, \quad (3.16)$$

where $\mathbf{v}_{q_i}$ and $\mathbf{v}_{a_{ij}}$ are the corresponding TF-IDF representations. The lexically closest alternative is then identified as

$$\ell_i = \underset{j \in \{0,1\}}{\operatorname{argmax}} s_{ij}. \quad (3.17)$$

The *Lexical Attraction Rate* measures how frequently the model selects this alternative:

$$\text{LAR} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = \ell_i). \quad (3.18)$$

A high value indicates that the model tends to prefer the candidate sharing greater lexical similarity with the question. However, this behavioral tendency does not establish that lexical overlap represents a dataset bias. To evaluate the structural relationship between lexical similarity and the target, a second quantity is computed:

$$\text{LTA} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(y_i = \ell_i), \quad (3.19)$$

where LTA denotes the *Lexical Target Alignment*. It measures how frequently the correct alternative is also the one with the highest lexical similarity to the question. The two measures therefore capture different phenomena:

$$\text{LAR} \rightarrow \text{model behavior}, \quad (3.20)$$

$$\text{LTA} \rightarrow \text{dataset structure}. \quad (3.21)$$

This distinction is essential. A high Lexical Attraction Rate means that the model follows lexical similarity, but this strategy is only useful if lexical similarity is correlated with the target. Conversely, a high Lexical Target Alignment indicates that such a shortcut is available in the dataset, even if a specific model does not exploit it systematically. The interaction between the two quantities can therefore be interpreted as follows:

$$\text{LAR} \gg 0.5 \quad \text{and} \quad \text{LTA} \gg 0.5 \quad (3.22)$$

provides evidence for both the availability of a lexical shortcut and the tendency of the model to exploit it.

Length Bias A parallel analysis is performed for the length of the two alternatives. Let $L(a_{ij})$ denote the length assigned to alternative j for item i , according to the length representation adopted in the implementation. The longer candidate is identified as

$$h_i = \operatorname{argmax}_{j \in \{0,1\}} L(a_{ij}). \quad (3.23)$$

The model's preference for longer alternatives is measured through the *Length Attraction Rate*:

$$\text{LenAR} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = h_i). \quad (3.24)$$

This quantity describes a behavioral property of the model. Values substantially above 0.5 indicate a tendency to select the longer alternative. As for lexical similarity, a separate measure is required to determine whether this heuristic could actually be rewarded by the dataset. We therefore compute the *Length Target Alignment*:

$$\text{LenTA} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(y_i = h_i). \quad (3.25)$$

This metric measures how frequently the correct alternative is also the longer one. The distinction again separates model behavior from dataset structure:

$$\text{LenAR} \rightarrow \text{preference of the model for longer alternatives}, \quad (3.26)$$

$$\text{LenTA} \rightarrow \text{association between alternative length and the target}. \quad (3.27)$$

A model may prefer longer alternatives even when length contains no useful information about the correct label. Conversely, the dataset may contain a correlation between length and correctness that is not exploited by a particular model. Considering both quantities prevents these two cases from being conflated.

No-Input Category Breakdown The main *No Input* diagnostics are additionally computed separately for each reasoning category. Let c denote a category and let

$$I_c = \{i \mid \text{category}(i) = c\} \quad (3.28)$$

be the subset of items belonging to that category, with $N_c = |I_c|$. Category-specific Raw Accuracy is therefore

$$\text{Acc}_c = \frac{1}{N_c} \sum_{i \in I_c} \mathbb{I}(\hat{y}_i = y_i), \quad (3.29)$$

and the corresponding Language-Only Advantage becomes

$$\text{LOA}_c = \text{Acc}_c - 0.5. \quad (3.30)$$

MCC and Position Bias are computed analogously by restricting the confusion matrix and the prediction frequencies to the subset I_c . In particular,

$$p_{0,c} = \frac{1}{N_c} \sum_{i \in I_c} \mathbb{I}(\hat{y}_i = 0), \quad (3.31)$$

with category-specific Position Bias

$$PB_c = |p_{0,c} - 0.5|. \quad (3.32)$$

This decomposition makes it possible to determine whether *No Input* solvability varies systematically across reasoning categories. A higher *No Input* accuracy for a given category must not be interpreted as evidence that the model is more competent at the corresponding reasoning process. In this experimental condition, the perceptual evidence required to perform spatial, temporal, causal, or other forms of grounded reasoning is unavailable. Category-level differences instead indicate that some groups of items may expose stronger linguistic or structural regularities than others. This distinction is particularly important when the *No Input* results are later compared with performance under perceptual conditions. For example, differences observed between spatial, temporal, and causal questions should not automatically be attributed to multimodal reasoning ability. Part of the variation may already be present before any perceptual modality is introduced. The category breakdown therefore provides a necessary baseline for separating genuine grounding effects from differences induced by the linguistic structure of the task.

Setup	Input	Description
No Input	None	The model receives no multimodal information. This condition measures the extent to which the answer can be inferred from linguistic priors or dataset biases.
Audio Only	Audio	The model receives only the raw audio track, while the visual content is removed. This setup evaluates the information that can be recovered exclusively from the acoustic modality.
Video Only	Video	The model receives only the visual content of the video, without audio. This condition measures how much task-relevant information can be extracted from the visual modality alone.
Audio + Video	Audio + Video	The complete audiovisual signal is provided to the model. This represents the fully multimodal condition and allows the contribution of joint audio-visual processing to be assessed.
Transcription + Video	Transcription + Video	The visual content is combined with a textual transcription of the audio track. This setup tests whether linguistic information extracted from audio is more effectively exploited when provided explicitly as text.

Table 3.4: Overview of the experimental setups used to evaluate the contribution of different input modalities.

4. Results and Discussions

4.1 Preliminar Analysis Result

Expected Behaviour under the Preliminar Condition

Prior to examining the experimental outcomes, qualitative expectations are delineated regarding model performance under the frame-only condition, providing a methodological framework for interpreting distinctions across question categories without retrospectively deriving explanations from observed results. For questions related to spatial dimensions, it is anticipated that there will be a relatively high presence of visual evidence, due to the inherent visibility of entity locations, movements, object configurations, and spatial relationships within the visual stream. The complexity may increase, however, when it becomes necessary to reconstruct this information over multiple temporal intervals, such as when tracking an entity through a series of movements or elucidating a spatial configuration resulting from several sequential actions. Temporal questions are expected to display a more varied profile. Actions and event sequences that are visually distinctive can often be pinpointed through timestamp data and organized by the temporal structure produced during Event Analysis. However, not all temporal information may be fully recoverable from visual input alone. Some questions might involve events where interpretation is reliant on a linguistic context or on interrelations between moments that remain ambiguous when only visual information is available. The adequacy of the resulting representation thus hinges on both the quality of chronological reconstruction and the inherent characteristics of the events involved. Causal questions are anticipated to be the most challenging within the visual-only scenario. Although the visual stream can provide insights into event succession, interactions, and observable state alterations, the underlying reasons for these events may rely on spoken dialogue or remain indistinct from mere images. Motivations, intentions, and explanatory relationships are less likely to be directly supported by visual evidence on its own. This supposition also emphasizes the methodological distinction be-

tween temporal succession and causal interdependence, as the occurrence of one event preceding another should not be interpreted as evidence of causation alone. These expectations are crafted independently of empirical performance and serve exclusively as a qualitative reference for further investigation. There is no presumption concerning the degree of category-level differences or which category might eventually attain the highest proficiency scores. Subsequent to this, proficiency results are contrasted with downstream Visual Semantic Verification (VSV) performance in order to differentiate between evidence availability and task execution. Visual proficiency determines if the independently constructed representation contains sufficient information to distinguish between corresponding captions and their counterparts (foils). In contrast, VSV performance evaluates whether the model delivers the correct response when both the question and inputs specific to the experimental condition are directly accessible. Together, these two metrics provide a diagnostic perspective that surpasses task accuracy alone. A Pass proficiency result paired with a correct VSV response indicates a scenario where relevant visual evidence was both represented and successfully applied in the downstream task. Conversely, a Pass outcome linked to an incorrect response suggests that despite having adequate evidence in the preliminary representation, there was failure in the downstream task. Such a scenario points to potential issues with evidence selection, integration, or reasoning at a later stage. When both proficiency and VSV performance are negative, it suggests an inadequacy at the perceptual or representational level, as neither the preliminary representation provides sufficient evidence for proficiency nor does the downstream task succeed. However, this comparison alone does not pinpoint the exact source of failure. The most diagnostically informative configuration emerges when proficiency indicates Not Pass but the downstream VSV response is accurate. In such cases, correct answer selection occurs despite lacking a visual representation that meets predefined proficiency criteria. This could be attributed to utilizing partial cues, alternative evidence, parametric priors, shortcut strategies, or mere guessing; however, the comparison alone does not identify the specific mechanism responsible for deriving the answer. A joint analysis of visual proficiency and downstream performance thus enables differentiation of whether the requisite evidence for a question was already present before question exposure or whether the model ultimately executed

the task successfully. This offers a richer diagnostic interpretation of VSV performance compared to assessing accuracy in isolation.

5. Conclusions

Conclusioni...

6. Future Perspectives

List of Abbreviations

VLM	Visual Language Model
VQA	Visual Question Answering
AGQA	Action Genome Question Answering
MVBench	Multi-modal Video Understanding Benchmark
VILMA	Video Language Model Assessment
MAIA	Multimodal AI Assessment
VSV	Visual Statement Verification
OEVQA	Open-Ended Visual Question Answering
MLM	Multimodal Language Model
AVQA	Audio-Visual Question Answering
Dave	Diagnostic benchmark for Audio Visual Evaluation
MMAU-Pro	Massive Multi-Task Audio Understanding and Reasoning benchmark, Pro version
SONIC-01	SOcial Natural Interaction Corpus (Omnimodal, v1)
GPT	Generative Pre-trained Transformer
MIRAGE	Assessing Hallucination in Multimodal Reasoning Chains of Multimodal Large Language Models

Bibliography

- [1] Stanislaw Antol et al. “VQA: Visual Question Answering”. In: *Proceedings of the IEEE International Conference on Computer Vision*. 2015, pp. 2425–2433.
- [2] Marco Baroni. “Grounding Distributional Semantics in the Visual World”. In: *Language and Linguistics Compass* 10 (2015). DOI: 10.1111/lnc3.12170.
- [3] Zesen Cheng et al. “VideoLLaMA 2: Advancing Spatial-Temporal Modeling and Audio Understanding in Video-LLMs”. In: *arXiv preprint arXiv:2406.07476* (2024). arXiv: 2406.07476. URL: <https://arxiv.org/abs/2406.07476>.
- [4] Michela Deodati et al. “Lost in Transcription: Investigating the Role of Audio Information for Video-Based Visual Statement Verification”. In: *Proceedings of the Twelfth Italian Conference on Computational Linguistics (CLiC-it 2026)*. Forthcoming. 2026.
- [5] Bowen Dong et al. “MIRAGE: Assessing Hallucination in Multimodal Reasoning Chains of MLLMs”. In: *Advances in Neural Information Processing Systems*. 2025. URL: <https://bit.ly/25mirage>.
- [6] Gemma Team. *Gemma 4 Technical Report*. 2026. DOI: 10.48550/arXiv.2607.02770. arXiv: 2607.02770 [cs.CL]. URL: <https://arxiv.org/abs/2607.02770>.
- [7] Rohit Girdhar et al. “ImageBind: One Embedding Space to Bind Them All”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2023, pp. 15180–15190.
- [8] Madeleine Grunde-McLaughlin, Ranjay Krishna, and Maneesh Agrawala. “AGQA: A Benchmark for Compositional Spatio-Temporal Reasoning”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2021, pp. 11287–11297.

- [9] Drew A. Hudson and Christopher D. Manning. “GQA: A New Dataset for Real-World Visual Reasoning and Compositional Question Answering”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2019, pp. 6700–6709.
- [10] Ilker Kesen et al. “VILMA: A Zero-Shot Benchmark for Linguistic and Temporal Grounding in Video-Language Models”. In: *International Conference on Learning Representations*. 2024. arXiv: 2311.07022. URL: <https://arxiv.org/abs/2311.07022>.
- [11] Sonal Kumar et al. “MMAU-Pro: A Challenging and Comprehensive Benchmark for Holistic Evaluation of Audio General Intelligence”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 40. 2026, pp. 22688–22697. URL: <https://sonalkum.github.io/mmdu-pro>.
- [12] Guangyao Li et al. “Learning to Answer Questions in Dynamic Audio-Visual Scenarios”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2022, pp. 19108–19118.
- [13] K. Li et al. “MVBench: A Comprehensive Multi-Modal Video Understanding Benchmark”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2024. DOI: 10.48550/arXiv.2311.17005. arXiv: 2311.17005.
- [14] Qwen Team. “Qwen2.5-Omni Technical Report”. In: *arXiv preprint arXiv:2503.20215* (2025). arXiv: 2503.20215. URL: <https://arxiv.org/abs/2503.20215>.
- [15] Qwen Team. “Qwen3.5-Omni Technical Report”. In: *arXiv preprint arXiv:2604.15804* (2026). arXiv: 2604.15804. URL: <https://arxiv.org/abs/2604.15804>.
- [16] Gorjan Radevski et al. “Dave: Diagnostic benchmark for audio visual evaluation”. In: *Advances in Neural Information Processing Systems* 38 (2026).

- [17] Ahmed Y. Radwan et al. “SONIC-O1: A Real-World Benchmark for Evaluating Multimodal Large Language Models on Audio-Video Understanding”. In: *arXiv preprint arXiv:2601.21666* (2026). arXiv: 2601.21666. URL: <https://arxiv.org/abs/2601.21666>.
- [18] Omar Sanseviero and Ian Ballantyne. *Introducing Gemma 3n: The Developer Guide*. Google Developers Blog, June 2025. URL: <https://developers.googleblog.com/en/introducing-gemma-3n-developer-guide/>.
- [19] Ravi Shekhar et al. “FOIL it! Find One Mismatch between Image and Language Caption”. In: *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2017, pp. 255–265. DOI: 10.18653/v1/P17-1024. URL: <https://aclanthology.org/P17-1024/>.
- [20] Guangzhi Sun et al. “Video-SALMONN: Speech-Enhanced Audio-Visual Large Language Models”. In: *arXiv preprint arXiv:2406.15704* (2024). arXiv: 2406.15704. URL: <https://arxiv.org/abs/2406.15704>.
- [21] Davide Testa et al. “All-in-One: Understanding and Generation in Multimodal Reasoning with the MAIA Benchmark”. In: *Findings of the Association for Computational Linguistics: EMNLP 2025*. 2025, pp. 20030–20050. DOI: 10.18653/v1/2025.findings-emnlp.1091. URL: <https://aclanthology.org/2025.findings-emnlp.1091/>.
- [22] Davide Testa et al. “MAIA: A Benchmark for Multimodal AI Assessment”. In: *Proceedings of the Eleventh Italian Conference on Computational Linguistics (CLiC-it 2025)*. 2025, pp. 1121–1134. URL: <https://aclanthology.org/2025.clicit-1.106/>.
- [23] Tristan Thrush et al. “Winoground: Probing Vision and Language Models for Visio-Linguistic Compositionality”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2022, pp. 5238–5248.

- [24] Jin Xu et al. “Qwen2.5-Omni Technical Report”. In: *arXiv preprint arXiv:2503.20215* (2025). DOI: 10.48550/arXiv.2503.20215. URL: <https://arxiv.org/abs/2503.20215>.
- [25] Jin Xu et al. “Qwen3-Omni Technical Report”. In: *arXiv preprint arXiv:2509.17765* (2025). DOI: 10.48550/arXiv.2509.17765. URL: <https://arxiv.org/abs/2509.17765>.

Appendix

A. Appendix