



UNIVERSITÀ DI PISA

Department of Philology, Literature, and Linguistics
Master's Degree Program in Digital Humanities

GRADUATION THESIS

MAI(A) the Sound Be With You

MAIA-AV: Extending Multimodal AI Assessment to the
Audio-Visual Domain

Supervisor

Prof. Alessandro Lenci

Student

Michela Deodati

ID: 597983

Co-Supervisor(s)

Dr. Bernardo Magnini (Company

Supervisor)

PhD Student Davide Testa

Academic Year 2025/2026

Contents

Introduzione	1
1 Literature review	5
2 Data Collection	12
2.1 From MAIA to MAIA-AV: Dataset Expansion and Curation	12
2.2 Expanding the Questions and Answers Dataset	23
3 Study Design and Experimental Setup	44
3.1 Overview of the Evaluated Models	44
3.2 Preliminary Analysis as a Diagnostic Tool for Multimodal Assessment . .	49
3.3 Core Reasoning Assessment Framework	82
3.4 Attention-Based Modality Analysis	112
4 Results and Discussions	116
4.1 Preliminary Analysis Results	116
4.2 Experimental Results and Discussion	125
5 Conclusions	150
6 Future Perspectives	153

Introduction

The present work was carried out during a research internship at Fondazione Bruno Kessler (FBK)[6], in Trento. Active since 2007 and rooted in the earlier Istituto Trentino di Cultura, founded in 1962 by Bruno Kessler, FBK is a non-profit research institute working across artificial intelligence, biomedical research, digital industry, space exploration, and quantum technologies, and is widely regarded as one of the leading research institutions in Italy. The internship was hosted by the Natural Language Processing (NLP) research unit, within FBK’s Augmented Intelligence Center.

Evaluating whether the remarkable success of Visual Language Models (VLMs) truly signifies a deep multimodal comprehension, or merely a reliance on surface-level correlations and linguistic biases, continues to be a significant area of inquiry. *MAIA: Multimodal AI Assessment* was introduced to address this question, as a native Italian benchmark for evaluating video-based reasoning in Vision-Language Models [24, 23]. MAIA investigates the extent to which VLMs coherently integrate visual and linguistic information. More specifically, it assesses whether model predictions are grounded in visual evidence rather than primarily driven by distributional regularities learned by the language component. This distinction is crucial because a model may answer a question correctly while failing to recognise the same information consistently when it is presented in a discriminative format. In such cases, the prediction may result from linguistic plausibility rather than genuine visual understanding. To address this issue, MAIA employs two aligned tasks: Visual Statement Verification and Open-Ended Visual Question Answering. Through these tasks, the benchmark evaluates not only model accuracy, but also the robustness and consistency of model behaviour across different evaluation formats. The benchmark further defines twelve fine-grained reasoning categories to identify which competencies are used by the models and in which areas linguistic shortcuts, unimodal collapse, or hallucinations emerge. A limitation of the original framework is that the dataset was designed and evaluated primarily on the basis of visual information. Real-world videos, however, are inherently audiovisual: relevant meaning may also be conveyed through dialogue, envi-

ronmental sounds, music, and other acoustic events. Recent multimodal and omnimodal architectures can process visual, auditory, and textual inputs, but the availability of multiple modalities does not guarantee that they will be integrated effectively. Models may over-rely on a single channel or be negatively affected by redundant or weakly relevant information. Building on this limitation, *Lost in Transcription: Investigating the Role of Audio Information for Video-based Visual Statement Verification* [4] extends the original MAIA research framework to a more complex analysis of the interaction between visual, linguistic, and auditory information, while preserving its competence-oriented perspective. The study considers two complementary representations of audio: raw signal and automatic transcriptions provided as textual input. By evaluating these representations both independently and in combination with video information, the analysis investigates whether audio contributes meaningfully to task performance and whether one representation has a stronger influence on model predictions. This leads to two research questions:

RQ1: *To what extent does raw audio affect video-based Visual Statement Verification when it is jointly processed with visual and textual information?*

RQ2: *How effectively do multimodal models integrate visual, textual, and auditory information when performing video-based reasoning and understanding tasks?*

The experiments show that raw audio has a limited and strongly context-dependent effect. It is most useful when visual information is unavailable, while its addition to an already informative visual input may instead provide little benefit or introduce noise. The results also show that transcriptions can produce greater interference than raw audio, suggesting that multimodal architectures do not necessarily achieve more effective cross-modal understanding simply because additional modalities are available. These findings must nevertheless be interpreted in relation to the nature of the dataset. The one hundred original MAIA videos were selected and annotated according to their visual content, and the questions were formulated without considering the audio track. Consequently, the benchmark does not systematically require auditory information to solve its items. My thesis project begins from this limitation. The following chapters investigate whether models can integrate multiple input modalities more coherently when the dataset is specifically designed so that auditory information is at least partially relevant. The central hypothesis

is that, under these conditions, models may use complementary signals across modalities to reconstruct missing information and develop a more genuinely multimodal understanding of video content, producing answers grounded in the available evidence rather than in linguistic plausibility or chance.

1. Literature review

The rapid evolution of Visual Language Models (VLMs) has progressively expanded the ability of neural systems to jointly process perceptual and linguistic information. This line of research builds on the hypothesis that linguistic meaning cannot be constructed exclusively from distributional regularities learned from textual corpora, but must also be grounded in perceptual experience [2]. Within this context, Visual Question Answering (VQA) emerged as one of the first systematic paradigms for assessing whether a model can answer questions about the content of an image [1]. Subsequent benchmarks, such as GQA, extended this setting by introducing compositional questions involving entities, attributes, and visual relations [10]. However, high task accuracy alone does not provide sufficient evidence of multimodal comprehension. Several studies have shown that models can exploit statistical regularities in answers, questions, or dataset distributions, producing correct predictions without necessarily *understanding* the associated visual content. Evaluation has therefore progressively moved beyond task performance towards the assessment of grounding, understood as the ability to associate a prediction with the relevant perceptual evidence. One popular approach to measure grounding is to create a small set of contrastive pairs, where a ground truth entry is paired with a Foil derived by a manipulation of a semantically relevant feature. For instance, *FOIL-it!* assess whether a model is sensitive to subtle errors of correspondence between image regions and lexical entities [21], and *Winoground* uses pairs of images and text descriptions that re-order their lexical elements in order to emphasize a different relational structure between them [25]. The transition from images to videos introduces additional challenges. Video understanding requires integrating information distributed over time, recognizing actions and state changes, ordering events, and establishing temporal and causal relations. *Action Genome Question Answering (AGQA)* addresses these requirements through compositional questions concerning objects, actions, and spatio-temporal relations [9]. The *Multi-modal Video Understanding Benchmark (MVBench)* similarly evaluates a broader range of competencies, including action perception, movement, temporal ordering, and dynamic reasoning [14].

Nevertheless, video benchmarks remain vulnerable to unimodal shortcuts and to strategies that rely predominantly on the most informative frame. *Video Language Model Assessment (VILMA)* addresses this limitation as a zero-shot benchmark based on video–text contrastive pairs [11]. For each complex capability, VILMA introduces a proficiency test designed to verify the preliminary perceptual ability required to solve the corresponding task. The findings indicate that, once these prerequisite conditions are considered, part of the apparently correct performance can be attributed to accidental or spurious strategies. In particular, the examined VLMs do not display a systematic advantage over statistical image-based models on tasks requiring strict temporal grounding. Within this line of research, MAIA introduced a natively Italian benchmark designed to provide a fine-grained evaluation of video-linguistic reasoning [24]. The resource covers twelve reasoning categories related to linguistic phenomena, including causal, counterfactual, implicit, spatial, temporal, sentimental, planning, uncertainty, and out-of-scope reasoning. Its main contribution lies in aligning two tasks built on the same videos and questions: *Visual Statement Verification (VSV)*, where the model selects the correct statement from a minimally contrastive pair, and *Open-Ended Visual Question Answering (OEVQA)*, where the model freely generates the answer. This setting makes it possible to assess not only accuracy but also coherence between discriminative comprehension and linguistic generation. A model may correctly answer a multiple-choice question while failing to reproduce the same information in an open-ended setting, or may behave inconsistently across equivalent linguistic formulations. To capture this phenomenon, MAIA introduces the *Aggregate Accuracy* metric, which rewards only those instances in which the model correctly answers all true–false pairs associated with the same question and simultaneously generates the correct response in the corresponding open-ended task. The reduction in Aggregate Accuracy relative to conventional single-instance accuracy highlights the fragility of apparently strong performance and the need to jointly evaluate robustness, consistency, and grounding. Although VLMs have traditionally focused on text–vision interactions, *real videos are intrinsically audiovisual*. Their interpretation may depend on music, prosody, background noise, and other acoustic signals that cannot be reconstructed from visual frames alone. An increasing body of research has accordingly extended multimodal language

model architectures to support joint processing of text, images, video, and audio. *Image-Bind*, for instance, introduced a shared representation space spanning images, text, audio, depth, thermal information, and inertial movement [8]. Other multimodal language models (MLMs) rely on modality-specific encoders combined with projection or fusion modules that align heterogeneous representations with the linguistic space. *Video-SALMONN: Speech-Enhanced Audio-Visual Large Language Models* introduces a multi-resolution Q-former¹ to align audio signals with synchronized video [22], while *VideoLLaMA2* combines spatio-temporal modeling and acoustic processing within a generative multimodal architecture [3]. More recent omnimodal architectures process multiple modalities within a unified framework. *Qwen2.5-Omni* adopts a thinker–talker architecture that accepts text, image, video, and audio inputs and produces textual or audio responses [16]. The subsequent *Qwen3.5-Omni* architecture integrates visual and audio encoders into a common backbone and introduces explicit temporal references to improve audio–video alignment [17]. All these developments show that heterogeneous signals may be treated in a unified framework. The existence of a modality, however, does not mean that it is functionally or compensatorily used. The distinction between **multimodal availability** and **multimodal integration** matters here: availability means that a system can handle a number of different input sources, while integration means that the model has to select, coordinate, and merge information across those modalities, even when it is only partially related. Hence, a system may be multimodal according to its architectural design and yet be heavily dominated by unimodal inputs in the prediction phase. This issue is evident in *Audio-Visual Question Answering* benchmarks such as MUSIC-AVQA, which evaluates spatio-temporal relations in musical performance videos and extends the analysis to objects, entities, and everyday activities [13]. Subsequent studies reveal a strong visual bias, showing that many questions can be answered using the video stream alone, without exploiting the audio channel. Under these conditions, similar audio-only and video-only performance does not demonstrate genuine integration, but rather suggests a form of unimodal collapse in which one

¹A multi-resolution Q-former, such as a Multi-Resolution Causal Q-Former, processes continuous audio, video, or image information at different temporal or spatial scales through distinct query banks, balancing fine-grained signals with broader contextual information.

modality dominates the decision process. To address this problem, *Dave* was introduced as a benchmark in which each question jointly requires visual and acoustic information, making either modality insufficient in isolation [18]. The benchmark distinguishes among action recognition, sound classification, temporal ordering, and multimodal synchronization, a decomposition that helps determine whether an error arises from the failure to perceive an event or from the inability to establish temporal relations between sounds and visual actions. Models tested on this benchmark achieve strong performance on individual components but encounter greater difficulty in audio–visual synchronization, indicating that cross-modal integration remains a central limitation. The increasing attention to audio has also led to benchmarks specifically designed to assess acoustic intelligence. *MMAU-Pro* evaluates 49 competencies distributed across speech, environmental sounds, music, and their combinations [12]. These tasks include long-audio comprehension, multi-track reasoning, spatial localization, prosody interpretation, and the integration of dialogue, music, and background noise. The persistently limited performance of current models indicates that acoustic processing cannot be reduced to dialogue recognition alone, since audio also conveys spatial, temporal, emotional, and contextual information that must be evaluated independently. A complementary perspective is provided by *SONIC-01*, a benchmark for omnimodal models based on audio–visual interactions drawn from real conversational and everyday contexts [19]. SONIC-01 combines open-ended summarization, multiple-choice questions, and temporal event localization, while also showing that replacing raw audio with transcription removes paralinguistic information related to tone and communicative intention. Its results identify temporal localization as one of the most challenging tasks, confirming that the integration of verbal content, acoustic signals, and visual sequences remains an open problem. The comparison between native audio and transcription is also central to my recent work, *Lost in Translation: Investigating the Role of Audio Information for Video-based Visual Statement Verification* [4]. The study extends the MAIA framework by introducing audio into the VSV task and comparing two representations of the acoustic modality: raw audio directly processed by *Qwen2.5-Omni-7B* and automatic transcriptions supplied as additional textual input to *Qwen2.5-VL-7B*. To characterize the audio content, the videos are divided into four classes:

Music: Instances characterized by repetitive vocalizations, onomatopoeic patterns, strongly repetitive tokens, or explicit lexical cues associated with non-speech audio.

Dialogue: Videos whose transcriptions display sufficiently coherent linguistic content, acceptable log-probability values, low degeneracy, and limited evidence of no-speech segments.

Silence/Noise: Videos with empty or highly unreliable transcriptions, particularly when associated with low log-probability values and relatively high no-speech probabilities.

Unknown: Residual cases that appear speech-related but do not satisfy the dialogue criteria with sufficient confidence, thereby preserving potentially informative content while explicitly marking its uncertainty.

Class labels and speech transcriptions were obtained by combining the outputs of *Whisper-1* and *GPT-4o-transcribe*, together with an assessment of output plausibility. This classification distinguishes results according to the informativeness of the available audio content. The findings reveal a marked dominance of the visual channel. Audio provides only a limited benefit, mainly when no visual information is available. Once rich visual input is provided, audio’s contribution becomes negligible and may even introduce noise or errors, with limited robustness observed when a single black frame is added instead. Under visually rich conditions, the difference between with-audio and without-audio settings is minimal, suggesting that audio functions more as a surrogate channel than as information fully integrated with vision. Audio transcriptions exhibit a different behavior. When added to already rich visual input, they reduce accuracy relative to the visual-only condition. This indicates that transcription is not a neutral transformation: once projected into the same representational space as the questions and answer options, it can compete with them, increasing the influence of textual regularities and introducing redundant or irrelevant information. *Lost in Translation* thus illustrates that different representations of the same acoustic content can affect the model’s decision process differently. Raw audio can be partially ignored when it is not useful, whereas transcription interacts directly

with the linguistic component of the model. Introducing an additional modality does not therefore guarantee an improvement, and may instead compromise predictions that were already correct on the basis of reliable visual evidence. These findings must be interpreted in light of the original MAIA design. The benchmark was created to investigate visual content, and annotators were explicitly instructed not to consider audio information, so its questions and answers were not designed to exploit acoustic evidence. The limited contribution of audio may therefore reflect both the models’ modality-integration mechanisms and the limited relevance of audio within the original dataset. Recent literature further emphasizes the importance of identifying the distinct sources of multimodal errors. *MIRAGE* disentangles hallucinations caused by perceptual failures from those arising from incorrect reasoning over correctly represented information [5]. Although primarily focused on visual reasoning, its methodological principle extends naturally to audio–visual settings, where an incorrect response may derive from failed sound recognition, inaccurate temporal synchronization, selection of the wrong modality, or erroneous inference over otherwise correctly represented evidence. Complementary attribution methods attempt to understand what aspects of the input are responsible for the generation of the answer. *OmniTrace* addresses this through an attribution system that applies to omnimodal autoregressive generation. It establishes bridges from the generated tokens to textual regions, images, and intervals in the audio/video modality, and aggregates attention across modalities into explanations that correspond to semantically interpretable linguistic units. This type of analysis can answer the important question of whether or not the model effectively uses audio-related signals. The state of the art shows a clear mismatch between the rapid growth of omnimodal architectures and the effectiveness with which they fuse diverse modalities: the same input may carry various signal modalities, yet no comparable progress has been made on the diagnostic side, where benchmarks may instead favor a specific signal, language-based shortcuts, dataset co-occurrences, or extraneous noisy signals. A central unresolved limitation is the construction of benchmarks in which audio and visual information are explicitly supervised. Assessing genuine audio–visual comprehension requires items in which audio contributes relevant information and the correct answer cannot be determined from a single modality alone. Such a design would distinguish

complementarity, dominance, substitution, and interference, while evaluating whether the model selects the appropriate evidence and combines it coherently across modalities. The present study follows this direction. Building on the competence-oriented MAIA framework and the recent findings of *Lost in Translation*, it investigates model behavior on a dataset specifically designed to require a genuine contribution from audio in order to answer the questions correctly. The central hypothesis is that, when audio and video provide partial and complementary evidence, multimodal comprehension depends on the model’s ability to reconstruct the missing information through cross-modal integration, avoiding both unimodal collapse and linguistic shortcuts based on plausibility.

2. Data Collection

2.1 From MAIA to MAIA-AV: Dataset Expansion and Curation

Video Selection

The dataset adopted in MAIA-AV follows the guidelines of Multimodal AI Assessment (MAIA) and comprises 100 short videos, approximately 30 seconds each, selected from YouTube Italia to represent everyday scenarios of Italian culture. Content domains include sports, culinary arts, urban life, and cultural activities, with particular attention to human presence and visually rich scenes. In the original formulation, video selection and annotation were oriented exclusively toward visual comprehension, with audio content excluded entirely. Building on this initial set, we collected an additional 100 videos, each uniquely identified from `video001` to `video100`. These identifiers map each video to its corresponding audio stream and text transcription. Each video was trimmed to isolate the segments relevant to the analysis, producing a corpus of clips consistent in length and content. Selection prioritized cases where the dialogue corresponds directly to elements visible in the scene. This choice makes it possible to study multimodal interaction in a targeted way: by focusing on scenarios where speech is semantically anchored to the visual content, the audio signal takes on a role that goes beyond mere presence, becoming necessary to answer the question correctly. This is the main departure from the original dataset: MAIA-AV (Audio-Visual) is a benchmark designed to combine and compare distributed evidence across the visual, auditory, and textual modalities.

Implementation Details The video acquisition pipeline leverages `yt_dlp` to extract specific temporal segments from chosen YouTube Italia videos. Instead of downloading entire videos for later editing, each clip is identified before acquisition using three key parameters: the source URL, the start timestamp, and the duration. These parameters define

an exact temporal interval, directly controlling the retrieval process. The conversion of timestamps into seconds is managed by `parse_time`, which translates HH:MM:SS or MM:SS formatted strings into integer values.

```
1 def parse_time(time_str):
2     """Converts HH:MM:SS or MM:SS string to seconds."""
3     parts = [int(p) for p in time_str.split(':')]
4
5     if len(parts) == 3:
6         return parts[0] * 3600 + parts[1] * 60 + parts[2]
7     elif len(parts) == 2:
8         return parts[0] * 60 + parts[1]
9     return parts[0]
```

Adding the requested duration to the start position determines the upper boundary of the segment:

```
1 start_sec = parse_time(start_time_str)
2 end_sec = start_sec + int(durata_str)
```

The resulting bounds, `start_sec` and `end_sec`, define the exact temporal window extracted from the source. Output paths are structured prior to downloader configuration, ensuring MP4 container compliance and consistent dataset indexing under the target directory:

```
1 filename_input = input(
2     "Enter output filename (e.g., clip): "
3 ).strip()
4 if not filename_input.endswith(".mp4"):
5     filename_input += ".mp4"
6 output = f"Dataset/video{filename_input}"
```

Prefixing filenames with `video` maintains a unified identifier convention across all dataset instances. Stream selection, temporal trimming, and media normalization are encapsulated within the `yd1_opts` dictionary. Format selection prioritizes high-resolution video streams

alongside optimal audio feeds:

```
1 ydl_opts = {
2     'format': 'bestvideo[height<=2160]+bestaudio/best',
3 }
```

This filter selects video streams up to 2160p paired with the highest quality audio stream, falling back to pre-merged streams (/best) if separate tracks are unavailable. Temporal constraints are applied during acquisition via `download_range_func`:

```
1     'download_ranges': download_range_func(
2         None,
3         [(start_sec, end_sec)]
4     ),
```

Embedding segment bounds directly into the retrieval mechanism eliminates separate media cropping steps and reduces bandwidth overhead. To maintain frame-accurate cutting across keyframe-based compressed video formats, the pipeline enforces keyframe generation at boundary points:

```
1     'force_keyframes_at_cuts': True,
```

Forcing keyframes ensures temporal alignment at cut points, preserving action continuity for downstream multimodal evaluation in MAIA-AV. Output formats and post-processing codecs are explicitly declared to standardize the corpus:

```
1     'merge_output_format': 'mp4',
2     'postprocessor_args': {
3         'ffmpeg': [
4             '-c:v', 'libx264',
5             '-preset', 'fast',
6             '-c:a', 'aac'
7         ]
8     },
```

FFmpeg re-encodes visual streams via H.264 (libx264, fast preset) and acoustic streams

via AAC, eliminating encoding heterogeneity across raw YouTube sources. The unified configuration dictionary combines acquisition parameters, formatting rules, output paths, and authentication credentials:

```
1 ydl_opts = {
2     'format': 'bestvideo[height<=2160]+bestaudio/best',
3     'download_ranges': download_range_func(
4         None,
5         [(start_sec, end_sec)]
6     ),
7     'force_keyframes_at_cuts': True,
8     'merge_output_format': 'mp4',
9     'postprocessor_args': {
10         'ffmpeg': [
11             '-c:v', 'libx264',
12             '-preset', 'fast',
13             '-c:a', 'aac'
14         ]
15     },
16     'outtmpl': output,
17     'cookiefile': 'src/youtube.com_cookies.txt',
18 }
```

Passing a session cookie file (cookiefile) enables retrieval of age-restricted or session-gated content without altering core extraction logic. Execution is wrapped in exception handling blocks, logging errors and returning non-zero exit codes upon failure:

```
1 try:
2     with yt_dlp.YoutubeDL(ydl_opts) as ydl:
3         ydl.download([url])
4 except Exception as e:
5     print(f"\nError: {str(e)}")
6     sys.exit(1)
```

Audio Conversion

In the MAIA-AV framework, the audio track of each video is treated as an independent source of information rather than a supporting element to the visual narrative. Speech, ambient sound, and other acoustic cues can carry evidence of their own, sometimes clarifying what the visual channel leaves ambiguous. Isolating the audio track makes it possible to examine this modality on its own terms and to track its contribution through every later processing stage. For each video, the original audio track is extracted and stored as a separate file, keeping the same timestamps and the same identifier as the corresponding video. This one-to-one correspondence keeps the visual, acoustic, and textual representations aligned to the same underlying event. The extracted audio files are the input for later stages of the pipeline, including automatic speech transcription and dedicated audio analysis. Extraction does not alter the original content: it simply separates the audio channel so that it can be processed independently while remaining synchronized with the source video.

Implementation Details The audio extraction stage transforms each audiovisual dataset instance into a standalone acoustic representation, maintaining multimodal alignment with the source video. This process is implemented using Python’s subprocess module, which calls FFmpeg to automate batch processing across the entire video collection. Input and output paths are clearly specified during the execution setup:

```
1 INPUT_FOLDER = Path("data/video")
2 OUTPUT_FOLDER = Path("data/audio")
3 VIDEO_EXTENSIONS = {".mp4"}
```

Decoupling the input and output directories serves to isolate raw video resources while also setting up an aligned audio mirror directory. By restricting VIDEO_EXTENSIONS to MP4, it ensures compatibility with the output format generated in the preceding acquisition stage. The main conversion process is encapsulated within mp4_to_mp3:

```
1 def mp4_to_mp3(
```

```

2     input_path: Union[str, Path],
3     output_path: Optional[Union[str, Path]] = None,
4     bitrate: str = "192k",
5 ) -> Path:

```

When an explicit destination path is omitted, the function retains the base filename and replaces the container extension:

```

1     input_path = Path(input_path)
2     output_path = (
3         input_path.with_suffix(".mp3")
4         if output_path is None
5         else Path(output_path)
6     )

```

This mapping strategy guarantees cross-modal alignment across representations (e.g., `video001.mp4` yields `video001.mp3`). Prior to invoking FFmpeg, parent directory existence is enforced:

```

1     output_path.parent.mkdir(
2         parents = True,
3         exist_ok = True,
4     )

```

This step allows nested input folder hierarchies to be automatically replicated within the output directory.

Extraction is executed using a single `subprocess.run` call:

```

1     subprocess.run(
2         [
3             "ffmpeg",
4             "-y",
5             "-i",
6             str(input_path),
7             "-vn",

```

```

8         "-c:a",
9         "libmp3lame",
10        "-b:a",
11        bitrate,
12        str(output_path),
13    ],
14    check=True,
15    )

```

The `-vn` flag discards video streams, ensuring that FFmpeg isolates the acoustic signal rather than performing generic media transcode operations. Acoustic streams are normalized using the `libmp3lame` codec at a constant target bitrate of 192 kbit/s (`-b:a 192k`). Passing `-y` overwrites pre-existing destination files without prompt interruption, facilitating seamless pipeline re-execution. Setting `check=True` ensures runtime errors raise a `CalledProcessError` rather than failing silently. Upon completion, the function returns the confirmed target path. Batch processing begins by recursively indexing and sorting compatible media files:

```

1 def find_videos(folder: Path) -> list[Path]:
2     """Return all supported video files found recursively."""
3     return sorted(
4         path
5         for path in folder.rglob("*")
6         if path.is_file()
7         and path.suffix.lower() in VIDEO_EXTENSIONS
8     )

```

Sorting file paths guarantees deterministic execution order across operating environments. Directory validity is verified before initializing conversion loops:

```

1 if not INPUT_FOLDER.exists():
2     raise FileNotFoundError(
3         f"Input folder not found: {INPUT_FOLDER}"
4     )

```

```

5 videos = find_videos(INPUT_FOLDER)
6 if not videos:
7     print(f"No videos found in: {INPUT_FOLDER}")
8     return

```

Validating the input directory early avoids redundant FFmpeg invocations when processing empty or invalid paths. For each video, relative path resolution preserves source subdirectory structures under OUTPUT_FOLDER:

```

1 for index, video_path in enumerate(videos, start=1):
2     relative_path = video_path.relative_to(INPUT_FOLDER)
3     audio_path = (
4         OUTPUT_FOLDER
5         / relative_path.with_suffix(".mp3")
6     )

```

Replicating source directory trees ensures structural parity between visual and acoustic datasets. Each file conversion is executed inside localized exception blocks:

```

1     try:
2         mp4_to_mp3(
3             video_path,
4             audio_path,
5         )
6         print("Audio extraction completed.")
7     except subprocess.CalledProcessError as error:
8         print(f"FFmpeg error: {error}")

```

Capturing CalledProcessError exceptions locally prevents single-file execution failures from interrupting broader batch operations while logging specific processing errors.

Transcription Extraction

Another key representation in our dataset is the audio transcription, produced following the methodology established in our previous work [4]. Transcriptions were generated

using the *Faster-Whisper* model (large-v3), a reimplementation of *OpenAI's Whisper* powered by the *CTranslate2* inference engine for Transformer architectures. The model processes raw audio directly, with the language parameter set to Italian. A *Voice Activity Detection* (VAD) filter is applied to isolate segments containing actual speech, reducing transcriptions of non-speech audio. The transcription procedure converts spoken dialogue into text that downstream models can process directly. This differs from the approach used in [4], which jointly employed *Whisper-1* and *GPT-4o-transcribe* to compare and align outputs for greater accuracy. The current pipeline instead relies on a single, locally executed model, which simplifies the workflow and makes it easier to reproduce.

Implementation Details The transcription stage was implemented as an independent component of the data preparation pipeline to associate each video with a textual representation of its spoken content. Powered by *Faster-Whisper* and the large-v3 model executed locally on GPU, this setup ensures full reproducibility while eliminating dependencies on external cloud services. The execution environment specifically configures a chosen CUDA device and sets up a local cache for model files. By using a local cache, model weights are confined to the working directory, which helps prevent repetitive downloads across different system locations. The transcription model is subsequently instantiated on CUDA using float16 precision:

```
1     model = WhisperModel(  
2         "large-v3", #MODEL_NAME  
3         device="cuda",  
4         device_index=0,  
5         compute_type="float16",  
6         download_root=str(HF_CACHE / "hub"), )
```

Prior to processing, the pipeline recursively retrieves and filters all .mp4 files within the input directory, sorting the resulting collection to guarantee deterministic execution order:

```
1     def find_videos(folder: Path) -> list[Path]:  
2         return sorted(  
3             path
```

```

4         for path in folder.rglob("*")
5             if path.is_file()
6                 and path.suffix.lower() in VIDEO_EXTENSIONS
7         )

```

The core transcription workflow relies on `model.transcribe()`. By directly parsing the audio embedded within the video container, Faster-Whisper eliminates intermediate extraction steps and yields chronologically ordered transcription segments:

```

1     segments, _ = model.transcribe(
2         str(video_path),
3         language="it",
4         beam_size=5,
5         vad_filter=True,
6         condition_on_previous_text=False,
7     )

```

Setting `language="it"` explicitly conditions decoding on Italian. This parameter bypasses automatic language identification and directly aligns model execution with the linguistic profile of the dataset. Three key hyperparameters govern decoding precision and stability:

- `beam_size=5`: Activates beam-search decoding, evaluating multiple competing hypotheses during generation to balance accuracy and computational overhead.
- `vad_filter=True`: Enables Voice Activity Detection to strip environmental noise, music, and unvoiced pauses, preventing non-linguistic signals from reaching the ASR engine.
- `condition_on_previous_text=False`: Disables cross-segment context conditioning, isolating individual speech segments and mitigating error propagation across the audio stream.

The resulting segment iterator is aggregated into a continuous transcript by stripping whitespace, discarding empty segments, and concatenating valid text chronologically:

```

1     return " ".join(
2         segment.text.strip()
3         for segment in segments
4         if segment.text.strip()
5     )

```

Omitting downstream linguistic post-processing or semantic filtering preserves an unadulterated representation of raw ASR performance. The corpus is processed sequentially, using each video filename as a unique primary key:

```

1     transcriptions = {}
2
3     for index, video_path in enumerate(videos, start=1):
4         try:
5             transcription = transcribe_video(video_path)
6             transcriptions[video_path.name] = transcription
7         except Exception as error:
8             transcriptions[video_path.name] = (
9                 f"ERROR: {error}"
10            )

```

Robust error handling isolates processing failures on individual files, preventing execution halts and cleanly decoupling unvoiced media from runtime exceptions.

Finally, extracted transcripts are exported to a structured CSV file indexed by `video_name` and `transcription`:

```

1     writer = csv.DictWriter(
2         csv_file,
3         fieldnames=[
4             "video_name",
5             "transcription"
6         ],

```

```

7         delimiter=";",
8     )
9     writer.writeheader()
10    writer.writerows(
11        {
12            "video_name": video_name,
13            "transcription": transcription,
14        }
15        for video_name, transcription
16        in transcriptions.items()
17    )

```

2.2 Expanding the Questions and Answers Dataset

The construction of linguistic data in MAIA-AV follows the same methodological framework as the original MAIA benchmark, where the question remains the basic unit for collecting answers and building the evaluation dataset. In the initial MAIA benchmark, each video had 2,400 questions, written by expert annotators following category-specific guidelines. Questions were written in an open-ended format to avoid simple yes-or-no answers. Since the original benchmark was designed for models based only on visual input, annotators were explicitly instructed to ignore all acoustic data when formulating questions. MAIA-AV expands the original framework by changing the information base used to generate questions. This version includes 300 questions in total, three per video, addressing spatial, temporal, and causal aspects of the video content. These categories match the structure of the original MAIA benchmark. The spatial category concerns locations or entities, whether visible or only referenced, and the spatial relationships between them. The temporal category concerns the sequence of events: when they occur and how they relate to each other, whether preceding, following, or simultaneous. The causal category concerns the reasons behind, or the consequences of, the actions, events, or behaviors shown in the video. The main difference between the original MAIA and MAIA-AV is the

inclusion of audio information. Some questions specifically target audio elements in the video, such as character dialogue, voice commands to infotainment systems, and ambient sounds, tying linguistic content to both the visual and auditory channels. The procedure for collecting answers follows the MAIA guidelines. Answers to the new set of 300 questions were collected through a questionnaire administered to four volunteers. All four are native Italian speakers, born in Italy, with at least a high school diploma. They watched the 100 videos and answered the corresponding 300 questions, producing four independent responses for each answer. The overall structure is unchanged: videos are the source of multimodal information, questions define what needs to be extracted from them, and the human-written answers serve as the linguistic reference for the benchmark components built on top of them.

Caption-Foil Generation

Building upon the human-annotated responses, the construction of *MAIA-AV* involves the automated generation of *caption-foil* pairs for the Visual Statement Verification (VSV) task. This process follows the methodology established in *MAIA*: each question-answer pair is first converted into a declarative true statement (*True Statement*, TS) and then modified to create a minimally contrastive false statement (*False Statement*, FS). Consequently, the caption and foil share maximum linguistic structure, differing mainly in question-relevant information. While the original benchmark collects eight human responses per question to yield eight TS-FS pairs, *MAIA-AV* adopts a scaled-down approach: four independent responses per question yield four captions and four corresponding foils. With three questions per video (spatial, temporal, and causal), each video comprises twelve caption-foil pairs. The automated generation uses GPT-4o-2024-08-06, matching the exact model version used in the original *MAIA* pipeline. Generation progresses in two distinct stages. In the first stage, the model receives only the question and a single human answer. Instead of creating novel content, the prompt directs the model to synthesize the question-answer pair into a single declarative statement, preserving the original semantic content. The instructions mandate minimal alterations to annotators’ phrasing while requiring the conversion of first-person expressions into an impersonal stance. As a result,

phrases like "I think that" or "I believe that" are removed from the caption without discarding any underlying uncertainty. Restricting outputs to concise declarative sentences leverages natural annotator variance while avoiding reliance on the model for factual core generation. This initial phase employs the prompt used for True Statement generation in *MAIA*:

Sei un assistente che crea affermazioni in italiano sui video.
Data una domanda Q e una risposta A relative a un video, devi creare un'affermazione S basata su A.
Mentre generi S, cerca di non alterare le parole che compongono A.
Se A include verbi o frasi in prima persona (ad es., 'penso', 'credo'), riformula S in modo impersonale, evitando la prima persona.
L'affermazione deve essere una frase breve e dichiarativa.¹

The instructions include few-shot examples demonstrating question–answer fusion, with specific coverage of negative responses, non-deducible information, and expressions of uncertainty. Appended to the prompt, the target question and annotator response appear in Q: and A: format, conditioning the model to output solely the statement S:. The second stage processes the generated caption to yield the corresponding foil. Unlike the initial phase, this prompt adapts to the target question category. The core methodology relies on constructing *minimal pairs*: the model modifies only the information necessary to invalidate the caption, avoiding extraneous additions and preserving sentence structure. Consequently, distinguishing between caption and foil tests a model's ability to ground claims in audiovisual content rather than relying on surface linguistic cues. This design adheres to *MAIA*, directing GPT-4o to alter exclusively elements relevant to the specified semantic domain. Across all categories, the system prompt assigns a uniform role:

Sei un assistente progettato per creare foil a partire dalle caption.²

¹English translation: *You are an assistant that creates statements in Italian about videos. Given a question Q and an answer A concerning a video, you must create a statement S based on A. While generating S, try not to alter the words composing A. If A includes first-person verbs or phrases (e.g., 'I think,' 'I believe'), rephrase S to be impersonal, avoiding a first-person perspective. The statement should be a concise, declarative sentence.*

²English translation: *You are an assistant designed to create foils based on captions.*

Subsequent instructions vary by targeted semantic dimension. For **spatial** questions, foil generation alters strictly the entity’s position or location:

Data una caption in italiano (C) riguardante la posizione o la localizzazione di qualcuno o qualcosa, il tuo compito è creare il suo foil (F) modificando solo l’informazione spaziale.
Non aggiungere altre informazioni rispetto a quanto dichiarato in C.³

An accompanying example relocates a woman standing in a poppy field to a school setting. Preserving the subject and surrounding text isolates spatial information as the sole determinant of truth value. For **temporal** questions, instructions dictate altering the temporal relationships within the caption:

Data una caption in italiano (C) riguardante gli eventi e quando avvengono, il tuo compito è creare il suo foil (F) modificando solo l’informazione temporale.
Non aggiungere altre informazioni rispetto a quanto dichiarato in C.⁴

An illustrative example converts the chronological order of two events into an alternative sequence, preserving core semantic content while modifying event timing or order. For **causal** questions, foil construction alters the cause-and-effect relationship expressed in the caption:

Data una caption in italiano (C) riguardante le cause o gli effetti degli eventi, il tuo compito è creare il suo foil (F) modificando le cause o l’effetto dell’evento principale.

³English translation: *Given an Italian caption (C) regarding the position or location of someone or something, your task is to create its foil (F) by changing only the spatial information. Don’t add other information with respect to what is stated in C.*

⁴English translation: *Given an Italian caption (C) regarding events and when they happen, your task is to create its foil (F) by changing only the temporal information. Don’t add other information with respect to what is stated in C.*

Non aggiungere altre informazioni rispetto a quanto dichiarato in C.⁵

The provided example retains the principal action, such as a person beginning to run, while altering solely the underlying cause. This focuses the contrast on the causal link rather than arbitrary scene modifications. In the final task setup, option ordering is randomized, preventing positional bias and ensuring performance metrics reflect semantic comprehension rather than structural pattern exploitation.

Implementation Details The *caption-foil* generation pipeline is implemented in `caption_foil.py`. The script relies on `pandas` for tabular data processing and the OpenAI Python client for model inference. Both caption and foil generation utilize `gpt-4o-2024-08-06`, with authentication managed via the `OPENAI_API_KEY` environment variable. Questions are loaded from `data/vsv/question_classification.csv`. Only the video identifier, principal dimension, and question text are required for this stage:

```
1 questions = pd.read_csv(  
2     QUESTIONS_FILE,  
3     sep=";",  
4     encoding="utf-8-sig",  
5 )  
6 questions = questions[  
7     ["video_name", "principal_dimension", "question_text"]  
8 ].copy()
```

Prior to processing, the script verifies column availability, discards missing values, normalizes category labels, and confirms the presence of only the three expected macroareas. Video identifiers are standardized using `normalize_video_name()`, which extracts the numeric identifier and converts it to the format used across the dataset:

```
1 def normalize_video_name(value):
```

⁵English translation: *Given an Italian caption (C) regarding the causes or the effects of events, your task is to create its foil (F) by changing the causes or the effect of the main event. Don't add other information with respect to what is stated in C.*

```

2     value = str(value).strip()
3     match = re.search(r"(\d+)", value)
4     if not match:
5         raise ValueError(
6             f"Impossibile ricavare il numero del video da: {value}"
7         )
8
9     return f"video{int(match.group(1)):03d}.mp4"

```

Additionally, the script ensures that each video is associated with exactly one spatial, one temporal, and one causal question. Duplicate video–macroarea pairs are rejected prior to generation. Human responses are automatically loaded from all CSV files within `data/vsv/raw_human_answers`. The file count is validated against the expected number of annotators:

```

1 files = sorted(RAW_ANSWERS_DIR.glob("*.csv"))
2 if len(files) != EXPECTED_ANNOTATORS:
3     raise ValueError(
4         f"Attesi {EXPECTED_ANNOTATORS} file di annotatori in "
5         f"{RAW_ANSWERS_DIR}, trovati {len(files)}"
6     )

```

Each file must contain the columns `titolo_video`, `macroarea`, and `risposta`. For backward compatibility, the alternative column name `risposte` is also accepted and automatically renamed. Annotator identifiers are extracted directly from the filenames:

```

1 if "risposta" not in df.columns and "risposte" in df.columns:
2     df = df.rename(columns={"risposte": "risposta"})
3 df = df[["titolo_video", "macroarea", "risposta"]].copy()
4 df["video_name"] = df["titolo_video"].map(
5     normalize_video_name
6 )
7 df["macroarea"] = (

```

```

8     df["macroarea"]
9     .astype(str)
10    .str.lower()
11    .str.strip()
12 )
13 df["annotator"] = file.stem.removeprefix("risposte_")

```

Questions and responses are subsequently aligned via the shared `video_name` and `macroarea` fields. The merge is constrained as *many-to-one*, matching multiple annotator responses to a single question:

```

1 data = answers.merge(
2     questions,
3     on=["video_name", "macroarea"],
4     how="left",
5     validate="many_to_one",
6 )

```

Additional consistency checks validate question keys against annotator responses. Missing answers or orphan responses interrupt execution prior to invoking any API requests. Responses corresponding to the same video and macroarea are grouped into a common pool. The pool identifier is formed by concatenating the video identifier with the corresponding macroarea:

```

1 data["video_id"] = (
2     data["video_name"].str.removesuffix(".mp4")
3 )
4 data["pool_id"] = (
5     data["video_id"] + "_" + data["macroarea"]
6 )

```

Each pool must contain exactly four responses. Individual entries are assigned a sequential position and a unique identifier:

```

1 data["pool_item"] = (

```

```

2     data.groupby("pool_id").cumcount() + 1
3 )
4 data["id"] = (
5     data["pool_id"]
6     + "_"
7     + data["pool_item"].astype(str).str.zfill(2)
8 )

```

API calls are centralized within `call_openai()`, which serves both caption and foil generation functions:

```

1 response = client.chat.completions.create(
2     model=MODEL,
3     messages=messages,
4     max_tokens=max_tokens,
5 )
6 return response.choices[0].message.content.strip()

```

The function executes up to five attempts for recoverable API failures. Rate-limit errors trigger a 30-second sleep interval, whereas generic API errors pause execution for 10 seconds. Authentication errors terminate execution immediately:

```

1 except RateLimitError:
2     if attempt == retries - 1:
3         raise
4     time.sleep(30)
5 except AuthenticationError:
6     raise RuntimeError(
7         "Errore di autenticazione: controlla OPENAI_API_KEY."
8     )
9 except APIError:
10    if attempt == retries - 1:
11        raise
12    time.sleep(10)

```

For each annotated response, caption and foil generation proceed sequentially. The former accepts the question and human answer as input, whereas the latter accepts the generated caption along with its associated macroarea:

```
1 caption = generate_caption(  
2     row["question_text"],  
3     row["risposta"],  
4 )  
5 foil = generate_foil(  
6     caption,  
7     row["macroarea"],  
8 )
```

Outputs are post-processed via `clean_generation()`, which strips optional S: or F: prefixes returned by the model:

```
1 def clean_generation(text, prefix):  
2     text = text.strip()  
3     if text.startswith(f"{prefix}: "):  
4         text = text[3:]  
5     elif text.startswith(f"{prefix}:"):  
6         text = text[2:]  
7  
8     return text.strip()
```

To accommodate long-running executions, the script implements a checkpointing mechanism. Cached entries are reused only if their question, annotator, and human answer match the current source data and both generated fields are non-null:

```
1 checkpoint_is_valid = (  
2     saved is not None  
3     and str(saved["question"]) == str(row["question_text"])  
4     and str(saved["annotator"]) == str(row["annotator"])  
5     and str(saved["human_answer"]) == str(row["risposta"])  
6     and pd.notna(saved["caption"])
```

```

7     and pd.notna(saved["foil"]))
8 )

```

Newly generated entries are appended incrementally to `caption_foil_checkpoint.csv`:

```

1 pd.DataFrame(generated_rows).to_csv(
2     CHECKPOINT_FILE,
3     index=False,
4     encoding="utf-8-sig",
5 )

```

Following generation, caption and foil positions are randomized independently within each question category. A fixed random seed guarantees reproducibility while maintaining a balanced distribution of options across positions:

```

1 rng = random.Random(RANDOM_SEED)
2 for category, indices in (
3     df.groupby("question_category").groups.items()
4 ):
5     indices = list(indices)
6     n = len(indices)
7     labels = [0] * (n // 2) + [1] * (n // 2)
8
9     if n % 2:
10         labels.append(rng.randint(0, 1))
11
12     rng.shuffle(labels)

```

The resulting target variable indicates the position of the correct caption: a value of 0 places the caption in `answer1`, whereas a value of 1 assigns it to `answer2`:

```

1 df["answer1"] = df.apply(
2     lambda row:
3     row["caption"]
4     if row["target"] == 0

```

```

5     else row["foil"],
6     axis=1,
7 )
8 df["answer2"] = df.apply(
9     lambda row:
10    row["foil"]
11    if row["target"] == 0
12    else row["caption"],
13    axis=1,
14 )

```

The final dataset is saved to `data/vsv/caption-foil/caption_foil.csv`. Alongside the randomized options and target labels, the output retains all metadata required to trace each pair back to its originating video, question, annotator, and human response:

```

1 columns = [
2     "id",
3     "video_id",
4     "video_name",
5     "pool_id",
6     "pool_item",
7     "question_category",
8     "question",
9     "annotator",
10    "human_answer",
11    "caption",
12    "foil",
13    "answer1",
14    "answer2",
15    "target",
16 ]
17 final_df[columns].to_csv(

```

```
18     OUTPUT_FILE,  
19     index=False,  
20     encoding="utf-8-sig",  
21 )
```

Caption–Foil Quality Validation

The automatic generation of a foil does not, by itself, guarantee that the resulting sentence constitutes a valid negative counterpart of the corresponding caption. Although generation is constrained through category-specific prompts, the model may introduce changes that are semantically unrelated to the targeted reasoning phenomenon, alter more information than necessary, or produce a sentence that remains compatible with the original caption. For this reason, the generated *caption–foil* pairs undergo an additional semi-automatic validation stage before being included in the final VSV evaluation set. The validation procedure follows the two-step strategy adopted in *MAIA*, adapted here to the three reasoning dimensions considered in *MAIA-AV*. The first stage is a *structural check*, whose purpose is not simply to determine whether the caption and foil differ, but to verify that their difference concerns the semantic dimension the corresponding question is meant to evaluate. This distinction is fundamental for preserving the diagnostic nature of the benchmark. If, for instance, a temporal foil modified the identity of an entity rather than the ordering of two events, a model could distinguish the two statements through entity recognition alone, without any temporal reasoning. Similarly, a causal pair would stop providing a controlled test of causal understanding if the foil simultaneously changed the cause, the participants, and the location of the event. Structural verification is what keeps the contrast introduced during foil generation aligned with the competency each question is meant to probe. For **causal** instances, validation requires the foil to modify the cause, the effect, or the causal relation expressed by the caption, while preserving the remaining content as much as possible. For **temporal** instances, the contrast must concern temporal information, including event ordering, before–after relations, or the moment at which an event occurs. The **spatial** category calls for a slightly broader criterion. Unlike the original *MAIA* taxonomy,

where spatial reasoning is split into Total and Partial subcategories, the spatial questions in *MAIA-AV* combine phenomena from both settings. The validation prompt therefore accepts changes involving locations, positions, directions, and relative spatial relations, including relations that hold only during a particular phase of the video. This broader formulation avoids forcing the new dataset into a distinction its annotation scheme does not retain, while keeping the core requirement that the foil differ from the caption specifically in its spatial content. Pairs that satisfy this first criterion move on to a second validation step based on *Natural Language Inference* (NLI), where the caption is treated as the premise and the foil as the hypothesis. The model assigns one of three standard NLI labels: *Entailment*, *Contradiction*, or *Neutral*. The expected relation is *Contradiction*, since the foil is designed to offer an alternative statement that should not hold under the same interpretation as the caption. This step addresses a different source of error than the structural check. A foil may correctly manipulate the intended semantic dimension while still failing to create a sufficiently distinct alternative. Conversely, a sentence may differ substantially from the caption yet fail the structural test because the modification does not isolate the phenomenon under evaluation. The NLI result is nevertheless treated as a diagnostic signal rather than an infallible automatic decision, particularly for causal and, to a lesser extent, spatial statements. This limitation is most visible when the contrast is carried by the answer-bearing element itself, such as a change of location in a spatial question or a change in the participant occupying a specific temporal position: the resulting statement may not read as an outright contradiction to a general-purpose NLI model even though it targets exactly the information the question is meant to probe. A related case arises for causal pairs, where two alternative causes may be logically compatible in isolation even though only one is supported by the video; here too, a general-purpose NLI model may classify the relation as *Neutral* rather than *Contradiction*. A similar pattern was observed during the original *MAIA* validation, where qualitative inspection showed that many pairs labelled as neutral were nevertheless valid caption–foil contrasts. For this reason, only pairs that pass the structural check and receive the *Contradiction* label are accepted automatically. Structurally valid pairs classified as *Neutral* are marked for manual review, while *Entailment* cases are given higher review priority, since they provide stronger evidence that the foil

may remain semantically compatible with the caption. Category-sensitive structural verification and category-independent semantic inference thus provide two complementary safeguards. The former preserves the correspondence between each pair and the reasoning ability it is meant to test. The latter checks whether the resulting negative statement establishes an adequate semantic opposition with the corresponding positive statement. Neither stage is used to silently discard uncertain data: every instance flagged as structurally invalid, entailed, neutral, or affected by execution errors is isolated in a dedicated review set, manually inspected, and, whenever the validation reveals an actual problem, corrected so that the intended semantic contrast is restored. The corrected pair is then run back through the same validation procedure. Human verification at this final stage keeps automatic classifier errors from propagating into the benchmark, while the pipeline as a whole remains predominantly automated.

Automatic Validation and Manual Revision

Building on this procedure, the complete caption–foil file was revised manually in full, using the automatic labels as a prioritization signal rather than a final decision. Each automatically flagged instance was inspected together with its question, human answer, semantic category, caption, and foil, since the automatic labels alone do not always capture whether a pair is acceptable once considered in context. When a pair was judged semantically valid and sufficiently consistent with the intended contrast, it was retained without changes. Whenever the inspection revealed an actual issue, such as a foil targeting the wrong semantic dimension, an ambiguous or insufficiently discriminative contrast, a formulation that did not adequately reflect the original human answer, or a mismatch between the question and the generated statements, the pair was manually corrected instead. The final dataset thus combines automatic quality-control signals with human verification, with corrections applied only where direct inspection confirmed an actual issue.

Implementation Details The validation procedure is implemented as a separate post-processing stage operating on the output of the caption–foil generation pipeline. Its input is the complete dataset stored in `data/vsv/caption-foil/caption_foil.csv`. Valida-

tion is performed on the canonical caption and foil fields rather than on the randomized answer1 and answer2 columns used by the final VSV task, since option randomization affects only presentation order and must not interfere with the semantic relation being assessed. Before any API request is issued, the script verifies the availability of the fields required for validation and checks that all category labels belong to the three supported macroareas:

```
1 REQUIRED_COLUMNS = {
2     "id",
3     "question_category",
4     "caption",
5     "foil",
6 }
7 missing_columns = (
8     REQUIRED_COLUMNS - set(df.columns)
9 )
10 if missing_columns:
11     raise ValueError(
12         "Mancano le seguenti colonne nel CSV: "
13         + ", ".join(sorted(missing_columns))
14     )
```

The first validation stage selects a category-specific prompt according to the value stored in question_category. All three prompts share the same decision format but define validity according to different semantic constraints. Causal pairs, for instance, explicitly require the modification to concern causal information:

```
1 if category == "causale":
2     return (
3         "Given an Italian caption (C) dealing "
4         "with causal relations between events "
5         "and its foil (F), your task is to assess "
6         "whether F is a valid causal foil of C.\n\n"
```

```

7
8     "To be valid, F must modify the cause, "
9     "the effect, or the causal relation "
10    "expressed in C while preserving the "
11    "remaining relevant content as much "
12    "as possible.\n"
13
14    "If F is a valid causal foil, output "
15    "'correct'. Otherwise output "
16    "'not correct'.\n\n"
17
18    f"C: {caption}\n"
19    f"F: {foil}\n"
20    "Output:"
21 )

```

The temporal prompt follows the same structure but constrains the acceptable difference to event ordering or other temporal properties:

```

1  if category == "temporale":
2      return (
3          "To be valid, F must modify the temporal "
4          "information expressed in C, such as "
5          "the ordering of events, before/after "
6          "relations, the moment at which an event "
7          "occurs, or another temporally relevant "
8          "relation.\n"
9      )

```

For spatial instances, the prompt is deliberately formulated to cover both stable locations and time-dependent spatial configurations:

```

1  if category == "spaziale":
2      return (

```

```

3         "The spatial category considered here "
4         "includes both locations and spatial "
5         "relations, including cases in which "
6         "spatial information refers to a specific "
7         "moment or phase of the video.\n\n"
8
9         "To be valid, F must modify spatial "
10        "information expressed in C, such as "
11        "the location, position, direction, or "
12        "spatial relation of a person, object, "
13        "or event.\n"
14    )

```

The model is constrained to produce only one of two labels, correct or not correct.

Outputs are normalized before being stored, to remove minor formatting variations:

```

1 def normalize_structural_label(text):
2     value = clean_text(text).lower()
3
4     if value.startswith("output:"):
5         value = value[len("output:"):].strip()
6
7     value = value.rstrip(".").strip()
8
9     if value == "not correct":
10        return "not correct"
11
12    if value == "correct":
13        return "correct"
14
15    return "invalid"

```

The structural result is saved for every caption–foil pair in the `structural_eval` field.

Incremental serialization is performed after each processed instance, so that completed evaluations survive an interruption:

```
1 df.at[index, "structural_eval"] = result
2
3 df.to_csv(
4     STRUCTURAL_OUTPUT,
5     index=False,
6 )
```

The second stage applies only to pairs that pass structural validation. Caption and foil are mapped respectively to the NLI premise S1 and hypothesis S2:

```
1 prompt = (
2     "Your task is to determine the natural "
3     "language inference (NLI) relationship "
4     "between S1 and S2.\n\n"
5
6     "The possible labels are:\n"
7     "- Entailment: S2 logically follows "
8     "from S1.\n"
9     "- Contradiction: S2 contradicts S1.\n"
10    "- Neutral: S1 and S2 are related, "
11    "but S2 neither follows from nor "
12    "contradicts S1.\n\n"
13
14    f"S1: {caption}\n"
15    f"S2: {foil}\n"
16    "Output:"
17 )
```

The resulting output is normalized to one of the three expected NLI labels:

```
1 if value == "contradiction":
2     return "Contradiction"
```

```

3
4 if value == "neutral":
5     return "Neutral"
6
7 if value == "entailment":
8     return "Entailment"

```

The two validation signals are then combined into an explicit status variable, with automatic acceptance limited to pairs that are structurally valid and classified as contradictions:

```

1 def get_validation_status(
2     structural_eval,
3     nli_eval,
4 ):
5     if structural_eval != "correct":
6         return "FAIL_STRUCTURAL"
7
8     if nli_eval == "Contradiction":
9         return "PASS"
10
11     if nli_eval == "Neutral":
12         return "REVIEW"
13
14     if nli_eval == "Entailment":
15         return "REVIEW_HIGH_PRIORITY"
16
17     return "ERROR"

```

Pairs that fail the structural check do not proceed automatically to semantic acceptance and are marked as FAIL_STRUCTURAL. Structurally valid pairs classified as Neutral are assigned REVIEW, while Entailment produces REVIEW_HIGH_PRIORITY. This hierarchy keeps uncertain model judgements from being treated as definitive dataset errors. At the end of the procedure, all non-passing cases are extracted into a dedicated manual-review

file:

```
1 review_mask = df[
2     "validation_status"
3 ].isin(
4     [
5         "REVIEW",
6         "REVIEW_HIGH_PRIORITY",
7         "FAIL_STRUCTURAL",
8         "ERROR",
9     ]
10 )
11
12 review_df = df.loc[
13     review_mask
14 ].copy()
15
16 review_df.to_csv(
17     OUTPUT_DIR /
18     "caption_foil_manual_review.csv",
19     index=False,
20 )
```

This file serves as the interface between automatic and human validation. Each flagged instance is manually inspected by comparing its caption, foil, question, and semantic category. Cases where the automatic validator is overly conservative but the contrast is in fact appropriate are retained unchanged. When inspection instead confirms that the foil modifies the wrong semantic property, remains compatible with the caption, introduces unrelated information, or otherwise fails to provide the intended minimal contrast, the foil is corrected by hand. Corrected instances are then re-evaluated through the same structural and NLI checks before being incorporated into the final dataset. The complete validation output retains the original metadata alongside `structural_eval`, `nli_eval`,

and `validation_status`. This preserves full traceability from each final VSV item back to its generated caption, foil, question, annotator response, automatic validation decisions, and any manual correction.

3. Study Design and Experimental Setup

3.1 Overview of the Evaluated Models

Qwen Family

The Qwen models considered in this work share an omnimodal architecture designed to process heterogeneous information within a unified computational framework. Both Qwen2.5-Omni and Qwen3-Omni accept textual, visual, and acoustic inputs, with modality-specific encoders transforming raw signals into representations compatible with a shared language-model backbone. Multimodal integration takes place after perceptual encoding, and the whole family places particular weight on the temporal coordination of audio and video [26, 27].

Qwen2.5-Omni-3B

Qwen2.5-Omni introduces the *Thinker-Talker* architecture, separating multimodal understanding from speech generation. The *Thinker* processes the input modalities and produces textual representations and responses, while the *Talker* takes the Thinker’s internal representations and generates speech from them. Perceptual understanding and spoken generation thus remain functionally distinct components of an otherwise end-to-end pipeline [26]. Text is processed through the Qwen tokenizer. Audio is resampled to 16 kHz and converted into a 128-channel mel-spectrogram, with the encoder producing representations at a temporal resolution of roughly 40 ms. Visual information is handled by a Vision Transformer derived from Qwen2.5-VL and shared across image and video inputs; frames are sampled dynamically to preserve relevant visual content while matching the temporal granularity of the acoustic stream [26]. The modality-specific encoders are limited to perceptual extraction: their output is passed to the *Thinker*, where multimodal information is processed within the shared Transformer. Both audio and visual encoders additionally support block-wise processing, allowing long multimodal sequences to be handled incre-

mentally [26]. Multimodal alignment relies on *Time-aligned Multimodal Rotary Position Embedding* (TMRoPE), which represents positional information along temporal, vertical, and horizontal dimensions. Text is assigned identical positional values across all three; images retain a constant temporal coordinate while preserving spatial position along height and width; audio receives temporal positions tracking its progression through the acoustic stream [26]. Video combines both axes: each frame keeps its two-dimensional spatial layout while carrying a temporal coordinate tied to its position in the clip, on a scale coordinated with the acoustic representation. Audio and video can thus be placed within a common temporal reference system, which is what allows the *Thinker* to relate acoustic and visual evidence occurring at corresponding moments. Encoded representations from both streams are arranged within a single multimodal sequence and processed jointly by the language model’s attention mechanism; the model therefore integrates audio and video before generating a response, rather than combining independent modality-specific predictions late in the pipeline [26]. The *Talker* is a dual-track autoregressive Transformer decoder, conditioned on the high-dimensional representations it receives from the *Thinker*; the resulting discrete speech units are converted into a waveform through a sliding-window diffusion Transformer [26]. The technical report describes this organization at the level of the Qwen2.5-Omni family as a whole, without a fusion mechanism specific to the 3B variant [26].

Qwen3-Omni-30B-A3B-Instruct

Qwen3-Omni keeps the Thinker-Talker organization of its predecessor while redesigning several components. It handles text, images, audio, and video, and supports both textual and spoken generation; in the 30B-A3B configuration, both the *Thinker* and the *Talker* adopt a Mixture-of-Experts architecture, increasing overall capacity while limiting the parameters activated at each inference step [27]. The *Thinker* still carries out multimodal understanding and text generation by combining representations from dedicated perceptual encoders. The *Talker*, however, is more loosely coupled to it than in Qwen2.5-Omni: rather than depending directly on the Thinker’s high-level textual hidden states, it conditions speech generation on textual context together with high-dimensional visual and

acoustic representations [27]. Audio is processed by the *Audio Transformer* (AuT). Raw audio is resampled to 16 kHz and represented as a 128-channel mel-spectrogram; convolutional downsampling brings the resulting representations to roughly 12.5 Hz, or about 80 ms per acoustic frame. AuT uses attention windows of different sizes to support both global audio understanding and block-wise processing [27]. Visual input is processed by an encoder derived from Qwen3-VL and initialized from SigLIP2-So400m, shared across images and video; as in Qwen2.5-Omni, frames are sampled dynamically to preserve visual content while matching the temporal granularity of the audio stream [27]. TMRoPE is retained for multimodal alignment, again along temporal, vertical, and horizontal dimensions: text keeps identical coordinates across all three, images hold a constant temporal coordinate while preserving spatial position, and audio now receives absolute temporal identifiers at roughly 80 ms resolution [27]. Video frames are assigned temporal identifiers taken from their actual timestamps, on the same scale used for audio, so acoustic and visual representations share an absolute temporal coordinate system while spatial information is still carried separately through height and width [27]. This replaces the fixed temporal chunking of Qwen2.5-Omni with alignment based on absolute temporal identifiers, producing a finer correspondence between acoustic and visual events and removing the dependence on a predefined chunk size [27]. The aligned representations are then processed by the Mixture-of-Experts *Thinker*, meaning that audio and visual evidence are combined within the shared Transformer computation only after perceptual encoding and temporal alignment; the pre-training strategy reinforces this from the outset, drawing on unimodal and cross-modal data spanning image-text, audio-text, video-text, video-audio, and video-audio-text combinations [27]. Speech generation is likewise redesigned. The *Talker* predicts a multi-codebook acoustic representation based on residual vector quantization: the first codebook is generated autoregressively, while a Multi-Token Prediction module estimates the remaining residual codebooks. The final waveform is produced by *Code2Wav*, a causal convolutional renderer built for incremental synthesis [27].

Gemma Family

The Gemma models considered in this work implement multimodality through two distinct strategies. Gemma 3n E4B relies on dedicated perceptual encoders that convert acoustic and visual signals into representations processed by the language-model backbone. Gemma 4 12B instead adopts a unified, encoder-free architecture, mapping raw perceptual information directly into the decoder’s embedding space through lightweight projection modules [20, 7].

Gemma 3n E4B IT

Gemma 3n accepts text, images, audio, and video and generates textual responses; separate perceptual encoders convert acoustic and visual signals into token representations for the central language model [20]. The E4B configuration is built on *MatFormer* (Matrioshka Transformer), which nests a smaller E2B computational path within the larger E4B structure; both are trained jointly, allowing the model to operate under different computational budgets without requiring separate weights [20]. Gemma 3n also introduces *Per-Layer Embeddings*, storing a substantial portion of the embedding parameters outside accelerator memory. As a result, E4B keeps only around four billion core parameters resident on the accelerator despite a larger overall parameter count, a gain in computational efficiency rather than a change to how modalities are integrated [20]. Audio is processed by an encoder based on the Universal Speech Model, producing roughly one acoustic token every 160 ms (about six tokens per second) that feed directly into the language model [20]. Visual input is handled by MobileNet-V5-300M, supporting resolutions of 256×256 , 512×512 , and 768×768 pixels across both image and video content [20]. A *Multi-Scale Fusion VLM adapter* processes the resulting visual representations before they reach the language model; its role is limited to combining features at different scales and does not amount to a general audiovisual fusion mechanism [20]. The language-model backbone remains the point where modalities actually meet: audio and visual encoders produce modality-specific representations that become available to the Transformer during generation, and the available documentation does not describe a

shared temporal alignment mechanism between acoustic and visual signals comparable to TMRoPE in the Qwen models [20]. Gemma 3n further employs *KV Cache Sharing*, letting keys and values from intermediate attention layers be reused by subsequent ones, which reduces redundant computation during the prefill stage, a benefit that is particularly relevant for the long sequences produced by audio and video input [20].

Gemma 4 12B IT

Gemma 4 is a natively multimodal family built on a decoder-only Transformer, and the 12B configuration departs from the rest of the family architecturally. Rather than relying on perceptual encoders, it is trained from scratch under a unified encoder-free paradigm, with visual and acoustic input mapped directly into the language model’s embedding space through lightweight projection modules [7]. This sets it apart from the other Gemma 4 configurations: the smaller E2B and E4B models retain distinct vision and audio encoders, and even the larger variants keep dedicated visual processing. Gemma 4 12B alone drops both the vision and the audio encoder, replacing them with a direct projection of raw perceptual information [7]. For visual input, the model operates on RGB patches of $48 \times 48 \times 3$ values. A large linear projection of roughly 35 million parameters takes the place of the Vision Transformer used elsewhere in the family, mapping the resulting representations directly into the language model’s embedding space [7]. Spatial information is preserved by adding two-dimensional coordinate-based positional embeddings to the projected patches, followed by a final LayerNorm before the visual representations enter the backbone, enough to keep spatial structure available without a separate visual encoder [7]. The audio pathway follows the same principle. Raw audio, sampled at 16 kHz, is divided into 40 ms segments of 640 values each; the USM-based Conformer encoder used by the other Gemma 4 variants is removed entirely, and each 640-dimensional segment is projected directly into the language model’s embedding space [7]. No additional positional encoding is introduced for these acoustic representations, since their sequential order already preserves the temporal structure of the audio stream. Visual and acoustic information thus enter the same embedding space through modality-specific linear projections rather than through independent high-capacity encoders [7]. In this design, integra-

tion happens directly inside the decoder-only Transformer: the projection modules only establish representational compatibility between raw perceptual input and the language model’s embedding space, while the Transformer itself carries out the actual interaction between perceptual and linguistic information, narrowing the separation between perceptual encoding and language processing found in more conventional multimodal systems [7]. The 12B model is dense and combines local and global self-attention, with five local attention blocks per global block. Local layers use RoPE, while global layers use p -RoPE with $p = 0.25$; the architecture also reuses keys as values in global attention layers, reducing the memory footprint of the global KV

3.2 Preliminary Analysis as a Diagnostic Tool for Multimodal Assessment

The preliminary analysis aims to determine which information relevant to the subsequent resolution of *MAIA-AV* questions can be autonomously recovered by each model from the visual content of the video alone. The analysis is accordingly conducted in a *question-independent* setting: the model receives solely the video input, with no access to the questions, corresponding answers, audio streams, or transcriptions. The output is a structured and temporally grounded representation of the visual content, which is only later compared with the informational requirements of individual questions. This separation distinguishes the model’s capacity to extract visual information from its ability to retrieve relevant evidence when guided by a linguistic query. While the preliminary analysis does not directly measure performance on the VSV task, its purpose is to characterize what the model can represent prior to the introduction of a specific question-answering objective. Isolating these two abilities addresses a key limitation of evaluations based solely on final accuracy. A correct prediction does not guarantee that the model has correctly identified relevant entities, reconstructed events, or established the relations necessary to support its output. Conversely, an incorrect answer fails to reveal which stage of the process caused the failure. The preliminary analysis thus introduces an intermediate diagnostic layer between the raw video content and the final task outcome. Methodologically, this approach relates

to *VILMA* [11], specifically regarding its distinction between proficiency tests and main tests. In *VILMA*, proficiency tests evaluate fundamental capabilities deemed necessary for resolving complex temporal phenomena via a zero-shot caption–foil paradigm. While our method does not replicate this specific evaluation protocol, it adopts the same underlying principle: complex performance must be interpreted relative to the prerequisite capabilities that underpin it. In *MAIA-AV*, these prerequisites are examined through a structured representation of the visual content. For each video, four multimodal and omnimodal models, subsequently evaluated on the VSV task, independently analyze temporally localized video segments. Outputs are initially kept separate to enable cross-representation comparison without prematurely merging predictions. Inter-model agreement is interpreted as an indicator of **prediction stability** rather than absolute ground truth, since multiple models may converge on an identical yet unsupported interpretation. The representation is constructed progressively across three levels: **entities**, **events**, and **relations between events**. This hierarchy separates directly observable perceptual information from information that requires temporal integration or higher-level inference. The **entity level** covers the objects, individuals, and other scene components identified by the model, together with their observable properties and the time intervals in which they are present. The **event level** captures actions, interactions, and state changes tied to specific temporal intervals. Because overlapping video segments may yield duplicate descriptions of the same event, extracted outputs are consolidated to eliminate redundancy while preserving genuinely recurrent events. Temporal anchoring maintains a direct link between each event and the video segment providing its supporting evidence. The **relation level** organizes extracted events according to temporal and inferential dependencies. Temporal relations encompass configurations such as *before*, *after*, *during*, and *overlaps*, supporting event reordering and the tracking of changes across the video timeline. Causal relations are treated independently of temporal succession: simple temporal precedence does not by itself establish a cause-and-effect relationship. The analysis accordingly differentiates causal information directly supported by visual evidence from relations that require further inference over events, states, or intentions, keeping observable content separate from information introduced by the model’s own inferential process. The distinction between representational

levels matters because questions assigned to the same nominal category can demand substantially different types and volumes of evidence. Identifying an object’s location within a single frame, for instance, requires fundamentally different operations than determining its position following a sequence of movements. Similarly, recognizing an isolated event differs from reconstructing a sequence of multiple events or inferring causal links between them. To capture these distinctions systematically, questions are analyzed along three operational dimensions: **spatial**, **temporal**, and **causal**. A unified three-level scale is applied across each dimension. Level 0 denotes cases where the required information can be directly retrieved. Level 1 involves integrating an explicit relation or a limited number of evidence units. Level 2 demands broader reconstruction across multiple events, temporal phases, state changes, or implicit relations. These criteria are summarized in Table 3.1. This classification reflects the informational requirements of the questions rather than their linguistic complexity, accounting for properties such as the temporal distribution of evidence, the number of relevant evidence units, the presence of inter-event relations, and whether the required information is explicit or implicit. Once constructed independently, the visual representation can be compared against the question requirements to determine whether the entities, events, and relations necessary to answer a question are already captured within the vision-only representation. The evaluation therefore does not depend directly on the model’s capacity to generate the correct answer. Instead, it assesses whether the supporting information is present within its visual representation. This comparison is particularly relevant for *MAIA-AV*, since it establishes a visual baseline prior to introducing additional modalities. Prior findings [4] showed that, in the original *MAIA* setting, incorporating acoustic information yielded no systematic performance gains, while including transcriptions led to a slight yet statistically robust decline in accuracy, consistent with a dataset originally designed around visual content rather than information explicitly conveyed via audio. In this sense, the preliminary analysis establishes *what can be seen* before evaluating the contribution of *what can be heard*. It clarifies whether the visual channel inherently contains the required evidence and highlights categories where visual information alone proves insufficient, so that the incremental contributions of audio and transcriptions can be assessed against an explicit visual baseline rather than solely

through shifts in downstream task accuracy.




Figure 3.1: Conceptual schema of the preliminary analysis.

Taken together, these components turn *MAIA-AV* from a purely result-driven evaluation into a diagnostic, competence-oriented framework. Final performance can thus be analyzed across specific sub-capabilities, including entity recognition, event identification, temporal integration, spatial reconstruction, and causal inference, disentangling evidence availability failures from errors in downstream selection or integration during the VSV task.

Category	Level 0	Level 1	Level 2
Spatial (S)	Direct retrieval of an entity’s location or spatial property from a single frame.	Explicit spatial relation or event-conditioned localization involving at most two evidence units.	Spatial reconstruction across multiple phases, including movement, relocation, or state change.
Temporal (T)	Recognition of a single event or state without temporal comparison.	Temporal localization or pairwise ordering between two explicit events.	Multi-event temporal reconstruction, including duration, repetition, onset, or state evolution.
Causal (C)	Retrieval of an explicitly stated cause without inference.	Explicit cause-and-effect relation involving at most two evidence units.	Implicit or multi-event causal reconstruction requiring contextual or intentional inference.

Table 3.1: Complexity levels for spatial, temporal, and causal questions.

Video Preprocessing: Temporal Segmentation and Frame Sampling

Prior to the entity extraction phase, each video is preprocessed by dividing it into four-second temporal windows with a stride of three seconds. Consecutive spans share one second of visual content. This overlap reduces the risk that an action or state change occurring near the boundary between two segments is only partially observed, while maintaining shared context across consecutive observations without altering the chronological order of the video. For each segment, the `start_time` and `end_time` are preserved, keeping each segment associated with its original temporal position in the video. This information is maintained throughout all subsequent stages of the pipeline, allowing observations to be reconnected with their exact position in the video. The windowing process serves exclusively to provide temporal indexing and structure visual evidence, leaving content and semantic interpretation to subsequent phases. Redundant information is deliberately retained at this stage to prioritize continuity of evidence, with redundancy resolution de-

ferred to the output consolidation step. The result of this preprocessing step is an ordered sequence of temporally localized visual units, used as the common input for every later stage of the analysis.

Implementation Details The preprocessing stage is implemented in `video_preprocessing.py` to convert each source video into a set of temporally indexed visual representations. The script first retrieves core metadata via `ffprobe`, including duration, frame rate, spatial resolution, codec, and audio stream presence. Although audio metadata is recorded, the acoustic signal is excluded from visual representation construction. Frames are extracted using OpenCV by seeking directly to target temporal positions, and each is saved as a JPEG image with its original timestamp embedded in both the filename and the JSON manifest, preserving temporal provenance independently of downstream processing. A core output of the script is the dense visual stream, sampled by default at 4 FPS:

```
1 parser.add_argument("--dense-fps", type=float, default=4.0)
```

Extracted frames are saved in the `dense_frames` directory for each video. Simultaneously, the script generates 32 uniform global frames and performs shot-boundary detection via `PySceneDetect`. These boundaries define variable-length segments with a target duration of 4 seconds, ranging from 2 to 6 seconds, and a default overlap of 0.5 seconds, extracting eight frames per segment. The preprocessing output is serialized into `metadata.json`, `frames.json`, `shots.json`, and `segments.json`. For the subsequent preliminary analysis, semantic extraction scripts operate directly on the dense stream rather than shot-based segments. Utility functions locate each video's `dense_frames` directory, parse timestamps from filenames, and sort frames chronologically, so that the model receives temporally ordered visual observations instead of isolated or semantically segmented images. The inference temporal windows are configured as follows:

```
1 parser.add_argument("--window-seconds", type=float, default=4.0)
2 parser.add_argument("--stride-seconds", type=float, default=3.0)
```

A 4-second window with a 3-second stride yields a 1-second overlap between consecutive

observations. The pipeline selects dense frames within these nominal boundaries, and if frame counts exceed the allowed maximum, frames are sampled uniformly across the window. At 4 FPS, a 4-second interval contains 16 frames, remaining well below the default 32-frame ceiling. Duplicate consecutive windows are automatically omitted. Each temporal unit is encapsulated in a `TemporalWindow`, mapping a unique segment ID to its start time, end time, and frame paths. Recorded boundaries reflect the timestamps of the first and last selected frames rather than nominal window limits, with millisecond precision maintained to ensure exact alignment with the source video. The transition to semantic extraction relies strictly on *ordered dense frames*. For each window, the model receives selected RGB images alongside a prompt specifying the video ID, temporal interval, and frame-to-timestamp map, with no questions, answers, transcriptions, or acoustic signals provided. Temporal segmentation thus serves purely as an indexing mechanism for visual evidence, deferring semantic interpretation to the subsequent extraction stage.

Semantic Representation Extraction

Building upon the defined temporal windows, the pipeline proceeds with semantic representation extraction. This stage generates a structured semantic description for each window, independently of the others. The extraction procedure is repeated across all four evaluated models. Although the models employ distinct architectures and inference modalities, each receives identical inputs and is constrained to produce a representation that follows the same common schema, shown in Table 3.2. Consequently, performance differences can be attributed directly to model capabilities rather than annotation formatting. For each window, the representation spans seven semantic categories:

entities : Operating at the most fundamental descriptive level, each entity receives a local identifier, a concise label, visually observable attributes, and its role within the scene when discernible. The representation tracks the temporal occurrence of the entity and the specific frames providing visual evidence. These local identifiers enable consistent references across other descriptive categories within the same segment.

actions : Actions represent visually observable occurrences, specifying the performing

agent and target object whenever feasible.

events : Events denote broader semantic units linked to one or more participants across a given temporal span. This distinction separates the detection of an isolated action from the wider context in which it occurs.

spatial relations : Spatial configurations between entities or scene elements are defined using a subject-relation-object structure. Permissible relations include configurations such as *in front of*, *behind*, *above*, *below*, *inside*, *outside*, *left of*, *right of*, *near*, and *on*. These spatial descriptors are temporally localized, meaning a spatial configuration observed in one segment is not assumed to hold throughout the video.

state changes : State changes capture observable transformations affecting an entity, represented as the contrast between a prior state and a subsequent state within the segment. This level is essential when answering a question requires tracking changes in an entity’s properties, position, or condition over time.

temporal relations : The model describes temporal links such as *before*, *after*, *overlaps*, *during*, *starts*, *finishes*, and *simultaneous*. These relations, however, remain strictly localized to the current window: global video temporal structure is not reconstructed at this stage, nor are event correspondences across different segments systematically resolved.

causal hypotheses : Causal candidate relations are explicitly separated from directly observed facts. The framework requires causal hypotheses to be flagged as inferred evidence, so that temporal precedence is never automatically interpreted as a cause-and-effect relationship.

The distinction between observed and inferred evidence is preserved through the *evidence_type* field, accompanied by a confidence score between 0 and 1. Each element also includes an *evidence_frames* list identifying supporting frames. During output normalization, frame references are retained only if they fall within the processed window, and generated temporal intervals are similarly bounded by segment limits, so that every semantic element stays traceable to the visual evidence that supports it. Categories may

remain empty if a model detects insufficient evidence within a segment. At the segment level, the pipeline preserves *segment_id*, *start_time*, *end_time*, and *input_frames*. At this initial stage, a unified global video description is not yet constructed. Because consecutive windows overlap, identical entities, actions, or events may be extracted repeatedly across adjacent segments. This redundancy is intentionally preserved during extraction to maintain a direct mapping between each prediction and its supporting local evidence, leaving entity reconciliation and global temporal aggregation to subsequent pipeline stages.

Field	Information represented
entities	Identified entities, observable attributes, roles, and temporal occurrence.
actions	Local actions, their actors, and involved objects.
events	Temporally extended local events and their participants.
spatial_relations	Spatial configurations between entities or scene elements.
state_changes	Observable transitions between an entity’s previous and subsequent states.
temporal_relations	Temporal relations between events identified within the same local context.
causal_hypotheses	Candidate causal relations explicitly marked as inferred evidence.

Table 3.2: Semantic information extracted from each temporal window.

Implementation Details Entity and semantic analysis is implemented through a unified extraction pipeline shared across all evaluated models. Model-specific wrapper scripts load the corresponding architectures and format selected input frames, while temporal framing, prompt generation, schema definitions, validation logic, and serialization are centralized in `utils.py`, so that every model is evaluated against an identical semantic protocol regardless of underlying architectural differences. The common semantic space is defined across seven categories:

```

1 SEMANTIC_CATEGORIES = (
2     "entities",
3     "actions",
4     "events",

```

```

5     "spatial_relations",
6     "state_changes",
7     "temporal_relations",
8     "causal_hypotheses",
9 )

```

The categories illustrated in Table 3.2 specify the structured information required from each temporal window. For each temporal window, the pipeline constructs a standardized extraction prompt via `build_semantic_prompt`, which supplies the video ID, window boundaries, and an explicit frame-to-timestamp map alongside the category definitions:

```

1 frame_map = "\n".join(
2     f'- frame {index}: "{path.name}", '
3     f'timestamp={get_timestamp(path):.3f}s'
4     for index, path in enumerate(
5         window.frame_paths,
6         start=1,
7     )
8 )

```

Rather than eliciting free-form descriptions, the prompt¹ enforces strict visual grounding: every extracted element must reference supporting frames, temporal boundaries, a confidence score, and an `evidence_type`. Directly visible details are classified as `observed`, whereas inferential claims are labeled as `inferred`, and the prompt explicitly instructs the model not to treat simple temporal succession as causality:

```

1     """Non inventare dettagli non visibili.
2     Usa evidence_type="observed"
3     per fatti direttamente visibili.

```

¹English translation: "Do not invent details that are not visible. Use `evidence_type="observed"` for directly visible facts. Use `evidence_type="inferred"` only for unavoidable inferences or causal hypotheses. For each element, indicate confidence between 0 and 1 and the exact frame names that support it. Do not confuse temporal succession with a causal relationship. If a category contains no elements, return an empty list. Return exclusively a valid JSON object."

```

4     Usa evidence_type="inferred"
5     soltanto per inferenze inevitabili
6     o ipotesi causali.
7     Per ogni elemento indica confidence
8     tra 0 e 1 e i nomi esatti dei frame
9     che lo supportano.
10    Non confondere una successione
11    temporale con una relazione causale.
12    Se una categoria non contiene elementi,
13    restituisci una lista vuota.
14    Restituisci esclusivamente
15    un oggetto JSON valido. """

```

Models are constrained to output structured JSON objects containing standardized meta-data, including `start_time`, `end_time`, `evidence_frames`, `evidence_type`, and `confidence`.

Entity identifiers generated within a window can be cross-referenced across actions, events, and relations in the same segment. The main execution loop is encapsulated in `process_video_with_inf`. After dense frames are partitioned into temporal windows, each window is processed sequentially:

```

1  for window_index, window in enumerate(
2      windows,
3      start=1,
4  ):
5      prompt = build_semantic_prompt(
6          video_directory.name,
7          window,
8      )
9      response = inferencer(
10         window.frame_paths,
11         prompt,
12     )
13

```

```

14     parsed = parse_model_json(response)
15
16     segment = normalize_semantic_output(
17         parsed,
18         window,
19     )
20
21     segments.append(segment)

```

The inferencer component is the sole model-dependent step in this loop, accepting frame paths and prompts and returning textual output. All downstream parsing and validation operations remain identical across architectures. To manage output formatting inconsistencies from generative models, the parsing stage strips Markdown wrappers and extracts the primary JSON object. If standard decoding fails, the pipeline attempts recovery using `json_repair` before marking the response invalid:

```

1  try:
2      parsed = json.loads(candidate)
3  except json.JSONDecodeError as first_error:
4      try:
5          from json_repair import repair_json
6          parsed = json.loads(
7              repair_json(candidate)
8          )
9      except Exception as repair_error:
10         raise SemanticOutputError(
11             f"JSON non valido: {first_error}"
12         ) from repair_error

```

Structural constraints are subsequently enforced by `normalize_semantic_output`, which clips temporal boundaries to window limits, restricts confidence scores to the interval $[0, 1]$, and filters out invalid frame references. Evidence types are validated against observed, inferred, or uncertain, with causal hypotheses automatically reclassified as inferred

if initially marked as observed:

```
1 confidence = _safe_float(  
2     item.get("confidence"),  
3     0.0,  
4 )  
5 item["confidence"] = round(  
6     max(0.0, min(1.0, confidence)),  
7     4,  
8 )  
9 evidence_type = str(  
10     item.get(  
11         "evidence_type",  
12         "observed",  
13     )  
14 ).lower()  
15 if evidence_type not in {  
16     "observed",  
17     "inferred",  
18     "uncertain",  
19 }:  
20     evidence_type = "uncertain"  
21 if (  
22     category == "causal_hypotheses"  
23     and evidence_type == "observed"  
24 ):  
25     evidence_type = "inferred"
```

Frame-level grounding validation retains only filenames present in the input temporal window:

```
1 valid_evidence = set(  
2     window.evidence_frames  
3 )
```

```

4 item["evidence_frames"] = [
5     str(name)
6     for name in evidence_frames
7     if str(name) in valid_evidence
8 ]

```

Normalized segments are enriched with temporal metadata from the preprocessing stage:

```

1 normalized = {
2     "segment_id": window.segment_id,
3     "start_time": window.start_time,
4     "end_time": window.end_time,
5     "input_frames": window.evidence_frames,
6 }

```

Outputs are saved for each video in a `_semantic.json` file alongside model metadata and window configuration settings. Window failures are logged in an `errors` field containing segment identifiers, temporal bounds, frame lists, and error descriptions, preserving optional raw responses for diagnostic review. This phase ultimately transforms temporally indexed visual evidence into structured local representations. Cross-window deduplication is deliberately omitted at this step to preserve local grounding, leaving element consolidation and full video reconstruction to subsequent analysis stages.

Event Analysis

Spatial Relations

Spatial information captures entity configurations within the scene and tracks positional changes across localized visual observations. Unlike temporal relations, spatial features are extracted directly during the semantic representation phase, while visual windows are still accessible. Spatial relations encode explicit configurations between entities, or between an entity and a scene element. These include positional, proximity, containment, and orientation relations such as *left of*, *right of*, *above*, *below*, *inside*, *outside*, *in front of*,

behind, *near*, and *on*. Because these descriptors are tied to localized segments, a single entity pair may exhibit distinct spatial arrangements across different temporal intervals. Displacements and state transitions also drive spatial variation. A positional change is represented by combining the spatial relations observed at different timestamps with the corresponding action and any recorded `state_changes`, distinguishing static scene layouts from dynamic rearrangements. Event-dependent localization further evaluates spatial configurations relative to an event interval, before, during, or after, isolating how a prior action alters an entity’s final position. Methodologically, co-occurrence does not imply a spatial relation. The presence of two entities within the same temporal segment provides insufficient grounds to assume proximity, orientation, containment, or contact: a relation is retained only when explicitly captured in the semantic analysis. Spatial evidence currently remains bound to these local semantic descriptions. The event extraction stage draws on `spatial_relations` for event reconstruction but does not construct a global spatial graph.

Temporal Relations

Temporal information is formalized during event extraction, where local semantic descriptions are consolidated into a chronological timeline covering the full video. Each consolidated event retains a `start_time` and an `end_time` defining its temporal boundary and, implicitly, its duration. Relations between events are modeled explicitly through pairs of `event_id` references. Supported relation types are *before*, *after*, *overlaps*, *during*, and *simultaneous*, each linking two existing events in the consolidated representation together with an associated confidence score. *Before* and *after* express temporal precedence and sequence, *during* denotes temporal containment, *overlaps* indicates partial interval sharing, and *simultaneous* marks concurrent occurrences. This structure captures chronological succession, repetition, containment, and overlap across the video timeline while keeping distinct occurrences separate. Generation order within the JSON output does not itself constitute temporal evidence: dependencies derive strictly from segment timestamps and relational metadata rather than from the model’s output sequence. The temporal schema adopts the following structure:

```

1  "temporal_relations": [
2  {
3    "first_event": "E0001",
4    "relation": "before|after|overlaps|during|simultaneous",
5    "second_event": "E0002",
6    "confidence": 0.0
7  }
8  ]

```

This consolidated event structure also serves as input for the subsequent causal analysis, though chronological organization is kept distinct from causal inference: temporal precedence does not by itself establish a cause-and-effect relationship.

Implementation Details The event analysis is implemented as a second-stage transformation over the structured semantic representations produced by the previous phase. No visual input is processed again at this stage. For each model, the corresponding semantic output is loaded and used as the only source of information for event reconstruction. The event analyzer can reorganize and consolidate previously extracted information, but it cannot introduce evidence obtained from a new inspection of the video, preserving the separation between perceptual extraction and subsequent structural reasoning. The input therefore consists of the temporally indexed semantic segments generated for the same model. Each segment retains its identifier and temporal boundaries together with the semantic categories extracted from the corresponding visual window. The information made available to the event analysis includes entities, actions, local events, spatial relations, state changes, temporal relations, and causal hypotheses:

```

1  keys = (
2    "segment_id",
3    "start_time",
4    "end_time",
5    "entities",
6    "actions",

```

```

7  "events",
8  "spatial_relations",
9  "state_changes",
10 "temporal_relations",
11 "causal_hypotheses",
12 )

```

These fields are considered jointly because an event may not be represented exclusively in the local events field. An action or an observable state transition, for instance, may provide relevant information for reconstructing an event even when the local semantic description does not encode it as an independent event. The event analysis therefore operates over the complete local semantic representation rather than simply concatenating the previously extracted event lists. Before inference, the semantic segments are arranged according to their temporal position and serialized into a structured textual input. The model is explicitly instructed to treat this representation as a closed source of evidence, with the common prompt constraining the reconstruction process through a set of rules designed to prevent additional perceptual or commonsense information from being introduced:

```

1
2  * usa esclusivamente le informazioni presenti
3    nell'analisi semantica
4  * non inventare azioni, oggetti, intenzioni,
5    emozioni o cause mancanti
6  * unisci soltanto i duplicati causati dalla
7    sovrapposizione delle finestre
8  * mantieni separati eventi realmente distinti
9    o ripetuti
10 * ricava i tempi soltanto dai segmenti forniti

```

A central function of this phase is the consolidation of local descriptions produced by overlapping temporal windows. Since the semantic extraction uses 4-second windows with a 3-second stride, adjacent segments share part of the same visual evidence, and the same

physical event can consequently appear in multiple semantic outputs. During event reconstruction, descriptions referring to the same occurrence are merged into a single event, an operation intended to remove redundancy introduced by the preprocessing strategy rather than repetitions that genuinely occur in the video. The prompt explicitly enforces this distinction. Events with similar descriptions are not automatically treated as duplicates: two occurrences are merged only when their temporal context and supporting segments indicate that they represent the same event observed through overlapping windows. Repeated actions occurring at different moments, by contrast, remain represented as distinct events, so that repetition is preserved as a meaningful temporal property of the video. The consolidated representation follows a common JSON schema. Each reconstructed event contains a unique identifier, a concise description, temporal boundaries, the participants involved, the semantic segments supporting the event, the type of evidence, and a confidence value:

```
1 {  
2   "event_id": "E0001",  
3   "description": "string",  
4   "start_time": 0.0,  
5   "end_time": 0.0,  
6   "participants": ["entity or participant"],  
7   "evidence_segments": ["segment_0001"],  
8   "evidence_type":  
9   "observed|inferred|uncertain",  
10  "confidence": 0.0  
11 }
```

The `event_id` field provides a stable reference for subsequent relational analyses. Events are assigned progressive identifiers such as E0001, E0002, and E0003, associated with the consolidated chronological representation rather than inherited from the local semantic outputs, so that events can be referred to consistently once multiple local descriptions have been merged. Temporal localization is inherited from the semantic analysis, since the event extractor does not estimate event boundaries by accessing the original frames. Instead, `start_time` and `end_time` must remain compatible with the temporal intervals

of the semantic segments supporting the event. When the same occurrence is represented across multiple overlapping windows, its consolidated interval can reflect the temporal evidence distributed across those segments. The `evidence_segments` field preserves provenance after consolidation. An event extracted from a single semantic window may contain only one segment identifier, whereas an event observed across overlapping windows may be associated with several:

```
1 "evidence_segments": [  
2   "segment_0004",  
3   "segment_0005"  
4 ]
```

Consolidation removes duplicated event descriptions without discarding the evidence from which they originated, keeping the resulting event representation traceable to the local observations produced in the previous phase, even though it is more compact than the semantic representation itself. Participants are derived from entities or participants already represented in the semantic analysis and are not generated as independent new entities at this stage, so that the event representation preserves semantic continuity between entity extraction and event reconstruction. The epistemic status of each event is represented through `evidence_type`. The analysis distinguishes three values:

```
1 "observed"  
2 "inferred"  
3 "uncertain"
```

Observed is used for events supported by information previously represented as directly visible. Inferred is reserved for information that already requires an inferential interpretation of the available semantic evidence. Uncertain represents cases in which the available evidence does not support a stronger classification. The event extraction stage is not intended to transform directly observed information into unsupported inference. Confidence is represented numerically in the $[0, 1]$ interval and subsequently normalized to ensure a consistent representation. Normalization also verifies the structural validity of the generated output before serialization. In addition to individual events, the same infer-

ence produces an initial global temporal organization. Temporal relations connect events through their newly assigned identifiers rather than through their textual descriptions:

```
1  "temporal_relations": [  
2  {  
3    "first_event": "E0001",  
4    "relation":  
5    "before|after|overlaps|during|simultaneous",  
6    "second_event": "E0002",  
7    "confidence": 0.0  
8  }  
9  ]
```

The allowed relations distinguish precedence, succession, containment, partial temporal overlap, and simultaneity, providing an explicit representation of the temporal structure that emerges once local events have been consolidated. Temporal relations cannot be derived solely from the order in which events are written in the generated JSON. The prompt explicitly requires temporal organization to depend on the intervals and relations available in the semantic representation: the order in which the text is generated is not treated as temporal evidence, and temporal succession is likewise kept distinct from causal interpretation. The model output is parsed as structured JSON rather than retained as free-form text. Responses are cleaned of possible Markdown delimiters, and the JSON object is extracted before decoding. The same robust parsing strategy used in the semantic stage allows malformed but recoverable outputs to be repaired when possible:

```
1  text = (  
2  text.replace("`json", "")  
3      .replace("`", "")  
4  .strip()  
5  )  
6  
7  start = text.find("{")  
8  end = text.rfind("}")
```

```

9
10 if start < 0 or end < start:
11     raise ValueError(
12         "La risposta non contiene un oggetto JSON"
13     )
14
15 payload = json.loads(
16     text[start:end + 1]
17 )

```

This phase primarily changes the granularity of the representation. The semantic analysis describes evidence locally within individual overlapping windows, while the event analysis reorganizes those observations into a coherent event-based representation of the complete video, removing redundancies caused by temporal overlap, preserving genuinely repeated occurrences, assigning stable identifiers, maintaining temporal and evidential provenance, and establishing the first explicit relations among events, all without any new inspection of the visual input.

Causal Inference

Causal inference operates exclusively on the consolidated event representation, independently of perceptual extraction. This phase receives no visual frames or raw multimodal input, relying solely on previously reconstructed events and their established temporal relations. Isolating causal reasoning in this way prevents the recovery of omitted visual details or the insertion of unobserved events. Before processing, the event representation is pruned to its essential fields: `event_id`, `description`, `temporal interval`, `participants`, `supporting semantic segments`, `evidence_type`, and `confidence`. The analysis considers four types of causal dependency:

- **causes**: one event directly produces the outcome represented by another.
- **enables**: one event establishes the conditions required for another to occur.

- **motivates:** the available evidence supports an intentional or motivational link between events.
- **prevents:** one event inhibits or counteracts another.

Each entry links a `cause_event` to an `effect_event`, optionally supplemented by `supporting_events` where additional events are needed to substantiate the inference. No new events are introduced during this phase. The resulting structure follows this schema:

```

1  "causal_relations": [
2    {
3      "cause_event": "E0001",
4      "relation": "causes|enables|motivates|prevents",
5      "effect_event": "E0002",
6      "evidence_type": "direct|inferred|uncertain",
7      "supporting_events": ["E0003"],
8      "explanation": "string",
9      "confidence": 0.0
10   }
11 ]

```

The epistemic status of each relation is indicated by `evidence_type`. Dependencies grounded in explicit physical interactions or state transitions are labeled `direct`. Connections that require contextual interpretation are labeled `inferred`, while partially supported links receive `uncertain`. Motivational dependencies require explicit support within the event representation rather than speculative plausibility. Temporal succession does not imply causation: a chronological link such as *before* or *after* remains insufficient on its own to establish a causal relation. Dependencies are introduced only when the semantic content explicitly supports causation, enabling, motivation, or prevention. Keeping these stages separate ensures that temporal ordering is never mistaken for causal evidence, preserving the distinction between extracted facts and inferential reasoning.

Implementation Details The causal inference stage is a text-only processing step applied to the consolidated event representation produced in the preceding phase. For each

video, the pipeline loads the corresponding `_events.json` file generated by the model. No raw video frames, audio signals, transcriptions, or downstream text prompts are supplied, so the causal analyzer operates exclusively on a closed representation of previously extracted visual information. Prior to inference, the event file is pruned by the `compact` function, which retains only the fields required for causal reasoning:

```
1 def compact(payload):
2     return {
3         "id_video": payload.get("id_video"),
4         "events": [
5             {
6                 "event_id": event.get("event_id"),
7                 "description": event.get("description"),
8                 "start_time": event.get("start_time"),
9                 "end_time": event.get("end_time"),
10                "participants": event.get("participants", []),
11                "evidence_segments": event.get("evidence_segments", [])
12            },
13            {
14                "evidence_type": event.get("evidence_type"),
15                "confidence": event.get("confidence"),
16            }
17            for event in payload.get("events", [])
18        ],
19        "temporal_relations": payload.get("temporal_relations", []),
20    }
```

This filtering isolates causal reasoning from the original perceptual features, requiring the model to derive dependencies strictly from the consolidated events and their temporal layout. The allowable relation types and evidence categories are bounded explicitly in code:

```
1 ALLOWED_RELATIONS = {
2     "causes",
```

```

3     "enables",
4     "motivates",
5     "prevents",
6 }
7 ALLOWED_EVIDENCE = {
8     "direct",
9     "inferred",
10    "uncertain",
11 }

```

These definitions separate direct cause-and-effect links (causes) from enabling conditions (enables), intentional drivers (motivates), and inhibitory actions (prevents). Inference is guided by a standardized prompt that instructs the model to rely solely on the provided events and temporal relations, explicitly forbidding it from treating simple chronological succession as causal evidence:

```

1
2 * usa esclusivamente gli eventi e le relazioni
3 temporali presenti nell'input
4 * non usare domande, caption, foil o altre
5 informazioni esterne
6 * la semplice successione temporale NON implica causalita'
7 * crea una relazione soltanto quando il contenuto
8 degli eventi supporta una dipendenza causa-effetto,
9 una condizione abilitante, una motivazione oppure
10 una prevenzione
11 * cause_event ed effect_event devono essere ID
12 di eventi presenti nell'input
13 * non inventare eventi, oggetti, intenzioni o cause
14 non ricavabili dalla rappresentazione fornita
15 * se non esiste evidenza sufficiente, restituisci
16 causal_relations come lista vuota

```

Within this schema, relationships are classified by epistemic confidence: `direct` for explicit physical or state-change interactions, `inferred` for contextual or intentional deductions, and `uncertain` for weakly supported cues. Output formatting enforces the same structured JSON schema used elsewhere in the pipeline:

```
1 {
2   "causal_relations": [
3     {
4       "cause_event": "E0001",
5       "relation": "causes|enables|motivates|prevents",
6       "effect_event": "E0002",
7       "evidence_type": "direct|inferred|uncertain",
8       "supporting_events": ["E0003"],
9       "explanation": "string",
10      "confidence": 0.0
11    }
12  ]
13 }
```

The fields `cause_event` and `effect_event` reference existing event IDs, while `supporting_events` optionally logs contextually relevant secondary events. Output parsing strips Markdown delimiters and decodes the outer JSON object. If formatting errors occur, `json_repair` attempts structural recovery, and if parsing fails entirely, a secondary generation call requests strictly formatted JSON without altering the input prompt:

```
1 text = text.replace("json", "").replace("\n", "").strip()
2 start = text.find("{")
3 end = text.rfind("}")
4 candidate = text[start : end + 1]
5 try:
6     return json.loads(candidate)
7 except json.JSONDecodeError:
8     from json_repair import repair_json
```

```
9     return json.loads(repair_json(candidate))
```

A normalization step then applies a sequence of structural checks before the relations are accepted. It first collects the set of valid event identifiers parsed from the input:

```
1 valid_ids = {  
2     event.get("event_id")  
3     for event in payload.get("events", [])  
4     if event.get("event_id")  
5 }
```

Candidate relations that reference an unknown event, or that link an event to itself, are discarded:

```
1 if (  
2     cause not in valid_ids  
3     or effect not in valid_ids  
4     or cause == effect  
5 ):  
6     continue
```

Relations using an unrecognized type are dropped, while evidence labels outside the allowed set default to inferred:

```
1 if rel not in ALLOWED_RELATIONS:  
2     continue  
3 if evidence not in ALLOWED_EVIDENCE:  
4     evidence = "inferred"
```

Relations sharing the same cause, relation, and effect are then deduplicated, keeping only the first occurrence of each triplet:

```
1 key = (cause, rel, effect)  
2 if key in seen:  
3     continue  
4 seen.add(key)
```

Supporting-event references are filtered in the same way, excluding invalid identifiers as well as the cause and effect events themselves:

```
1 supporting = [  
2     event_id  
3     for event_id in relation.get("supporting_events", [])  
4     if event_id in valid_ids and event_id not in {cause, effect}  
5 ]
```

Finally, confidence scores are clipped to the $[0, 1]$ interval:

```
1 try:  
2     confidence = float(relation.get("confidence", 0.0))  
3 except (TypeError, ValueError):  
4     confidence = 0.0  
5 confidence = max(0.0, min(1.0, confidence))
```

Final outputs are serialized per video into `_causal.json` files containing metadata, model parameters, and the resulting relations. Unresolvable errors yield empty relation arrays rather than speculative fallbacks.

NLI-Based Visual Proficiency Assessment

The final stage of the preliminary analysis integrates the information produced at the semantic, event, temporal, spatial, and causal levels into a consolidated representation of the visual evidence recovered by each model. The procedure operates exclusively on the outputs generated during the preceding phases, preserving the separation between visual analysis and question-dependent evaluation without performing any further inspection of the video. The integrated representation combines complementary information produced at different levels of the pipeline. Entity descriptions, spatial relations, and state changes originate from the semantic analysis. Consolidated events provide a global event-level representation together with their participants, temporal boundaries, supporting segments, evidence type, and confidence. Temporal relations connect these events through their identifiers, while causal relations encode the dependencies established during causal in-

ference. These components are treated not as independent observations, but as progressively structured views of the same visual evidence. Finalization does not amount to a simple concatenation of all intermediate outputs. Redundancies introduced by overlapping temporal windows have already been resolved during event consolidation, where multiple descriptions of the same occurrence are merged while genuinely repeated events remain distinct, and duplicated or structurally invalid causal relations have already been removed during causal normalization. The epistemic status assigned during the previous stages is preserved as well: information classified as observed, inferred, or uncertain remains distinguishable in the consolidated representation, together with the corresponding confidence values. Uncertain information is not automatically promoted to positive evidence, nor are conflicting descriptions resolved through an additional interpretation introduced at this stage. The final representation thus captures the entities, events, spatial configurations, temporal organization, and causal dependencies recovered by the model, without expanding the evidence beyond what the visual analysis actually produced. Crucially, **finalization does not constitute an answer to any VSV question**. Its purpose is to fix a representation of what the model has recovered from the visual modality before the linguistic content of the evaluation item is introduced. Questions, captions, and foils therefore play no role in constructing the representation, and once it has been finalized, its content remains fixed throughout the evaluation of the corresponding question: the question may provide context for interpreting the linguistic alternatives, but it cannot trigger the extraction of new entities, events, relations, or visual evidence. The proficiency assessment itself is formulated as a Natural Language Inference task. The finalized visual representation acts as the **premise**, while the caption and the corresponding foil are evaluated independently as alternative hypotheses. GPT-4o is employed exclusively as an NLI judge, not to answer the original question, generate a new description of the video, or select directly between the two alternatives. Its role is restricted to determining the logical relation between the evidence contained in the finalized representation and each hypothesis. For both caption and foil, the judge assigns one of three labels: entailment, contradiction, or neutral. Entailment indicates that the hypothesis is supported by the available visual representation. Contradiction indicates that it is incompatible with

the represented evidence. `Neutral` is used when the available information is insufficient to establish either relation, so that missing evidence is not interpreted as contradiction and absence of information is not treated as negative evidence. The caption and foil are judged separately. A caption–foil pair receives `PASS` only when the finalized representation entails the caption and contradicts the foil. Every other combination produces `NOT PASS`. This rule adopts a deliberately conservative criterion. Entailment of the caption alone is insufficient, since the same representation must also provide enough evidence to reject the minimally contrastive foil. Contradiction of the foil, in turn, cannot compensate for a caption classified as `neutral`. A model is therefore considered proficient on a given pair only when its visual representation supports both sides of the required distinction. The procedure is repeated independently for the four caption–foil pairs associated with each question. Question-level proficiency requires consistency across the complete set: a question receives `PASS` only when all four pairs satisfy the pair-level criterion, and is classified as `NOT PASS` if at least one pair fails. This strict aggregation reduces the influence of individual linguistic formulations and requires the recovered visual evidence to remain sufficient across multiple, minimally different realizations of the same underlying item. For each pair, the procedure records the classification assigned to the caption, the classification assigned to the foil, and the resulting binary decision, which are then aggregated into the final proficiency label associated with the question. The resulting measure is accordingly distinct from conventional question-answering accuracy: rather than evaluating whether the model can generate or select the correct answer after being exposed to the question, it evaluates whether the information independently recovered from the visual channel is sufficiently complete and discriminative to support the correct statement while rejecting its minimally contrastive alternative, connecting question-independent visual analysis to the competence-oriented evaluation objective of the preliminary analysis.

Implementation Details The final evaluation is implemented in `final_evaluation.py`, which operates on the finalized representations produced for the four evaluated models. These representations are stored as JSON files in `data/preliminar_analysis/final_results` and constitute the fixed visual premises used throughout the proficiency assessment. The

script expects one finalization file for each model and verifies their presence before starting the evaluation:

```
1 EXPECTED_MODELS = (  
2     "gemma",  
3     "gemma-4-12B",  
4     "qwen",  
5     "qwen-30B",  
6 )  
7  
8 FINAL_DIR = Path(  
9     "data/preliminar_analysis/final_results"  
10 )  
11  
12 final_files = [  
13     FINAL_DIR / f"{model}.json"  
14     for model in EXPECTED_MODELS  
15 ]
```

Finalized representations are loaded independently for each model and indexed by normalized video identifiers. The loading procedure accepts different top-level JSON organizations, including lists of video representations, dictionaries containing videos or results, individual video objects, and dictionaries directly indexed by video ID, allowing the subsequent NLI procedure to retrieve the complete representation associated with a video without modifying its content. Question metadata and caption–foil pairs are loaded separately from `question_classification.csv` and `caption_foil.csv`. Video identifiers and question categories are normalized before the two sources are joined on the video and principal question dimension. The question is therefore introduced only after the finalized visual representation has been loaded, and is retained as contextual information for the NLI judge rather than incorporated into the premise itself. The NLI evaluation uses `gpt-4o-2024-08-06` with deterministic generation:

```
1 MODEL = "gpt-4o-2024-08-06"
```

```

2  risposta = client.chat.completions.parse(
3  model=MODEL,
4  temperature=0,
5  messages=[
6  {
7  "role": "system",
8  "content":
9  f"{PROMPT}\n\n{REGOLE[categoria]}",
10 },
11 {
12 "role": "user",
13 "content": contenuto,
14 },
15 ],
16 response_format=RispostaNLI,
17 )

```

Each request contains three distinct fields: the question, the finalized representation used as the premise, and one hypothesis. The caption and foil are never evaluated jointly. Instead, two independent calls are performed using the same premise and question context:

```

1  nli_caption = classifica(
2  premessa,
3  row["domanda"],
4  row["categoria"],
5  row["caption"],
6  )
7
8  nli_foil = classifica(
9  premessa,
10 row["domanda"],
11 row["categoria"],
12 row["foil"],

```

13)

The system prompt explicitly restricts the judge to the information contained in the premise. External knowledge, completion of missing information, and plausibility-based reasoning are excluded, and incomplete, weak, or ambiguous evidence must be classified as *neutro*. Additional constraints are introduced according to the question category: spatial relations cannot be inferred from simple co-occurrence, temporal relations cannot be inferred from the textual order of events, and causal relations cannot be established from temporal succession alone. The model response is constrained through a Pydantic schema rather than parsed from unrestricted natural-language output:

```
1 class RispostaNLI(BaseModel):
2     etichetta: Literal[
3         "implicazione",
4         "contraddizione",
5         "neutro",
6     ]
```

This design restricts each inference to one of the three admissible NLI labels, preventing the judge from generating explanations or free-form answers. A missing or structurally invalid classification raises an error rather than being converted into an evaluation decision. Pair-level proficiency is subsequently computed through a deterministic rule implemented outside GPT-4o:

```
1 def valuta_coppia(caption, foil):
2     return (
3         "PASS"
4         if caption == "implicazione"
5         and foil == "contraddizione"
6         else "NOT PASS"
7     )
```

GPT-4o's role is thus limited to assigning the two NLI labels: it does not directly decide whether a proficiency test is passed. The binary outcome is derived programmatically, re-

ceiving PASS only when the finalized representation entails the caption and contradicts the corresponding foil. All other label combinations result in NOT PASS. Results are written incrementally to a model-specific checkpoint in the `nli` subdirectory. Before evaluating a pair, the script checks whether its identifier has already been processed, allowing interrupted executions to resume without repeating completed API calls. Each stored row contains the model, pair identifier, video, category, question, caption, foil, the two NLI classifications, and the resulting pair-level outcome. Once all model-specific evaluations have been completed, pair-level records are concatenated into `nli_risultati.csv`. Question-level proficiency is then computed by grouping results according to model, video, category, and question:

```
1 df = risultati.groupby(  
2 [  
3     "modello",  
4     "video_id",  
5     "categoria",  
6     "domanda",  
7 ]  
8 ).agg(  
9     totale_coppie=("esito", "size"),  
10    coppie_pass=("pass", "sum"),  
11 ).reset_index()  
12  
13 df["proficiency"] = [  
14     "PASS" if passate == totale else "NOT PASS"  
15     for passate, totale in zip(  
16         df["coppie_pass"],  
17         df["totale_coppie"],  
18     )  
19 ]
```

The aggregation adopts the same conservative principle used at pair level: proficiency is

assigned only when every caption–foil pair associated with the question passes. Since the experimental data contain four pairs for each question, this corresponds to the required 4/4 criterion, and a single NOT PASS is sufficient to classify the question as NOT PASS. The script additionally stores the number of evaluated pairs, the number of successful pairs, and their ratio in `question_proficiency.csv`. The implementation keeps the responsibilities of the two components distinct: the finalized JSON representation remains unchanged and functions solely as the visual premise, GPT-4o performs constrained NLI classification over each caption and foil, and the final proficiency decision is produced by deterministic aggregation rules rather than by an additional generative judgment.

3.3 Core Reasoning Assessment Framework

Experimental Logic and Hypotheses

The experimental design of MAIA-AV goes beyond a simple comparison of model accuracy across conditions. Its primary objective is to analyze the role played by individual modalities in the decision process and verify whether correct answers are genuinely derived from the available evidence. Experiments are structured around a **Modality-Ablation Design**: the test item is held constant while the input modalities are systematically varied, so that shifts in prediction can be observed as a specific information source is added or removed. This control enables **paired comparisons**, attributing behavioral variance more precisely to the selected modality configuration rather than to differences between items. The design further distinguishes between **Modality Sufficiency** and **Modality Necessity**. A modality is sufficient when it independently contains all the information required to solve the task, though sufficiency does not imply that the modality is uniquely useful, nor does it preclude further benefit from additional sources. Conversely, supplementing a sufficient modality with another does not guarantee a performance gain: the added input may be redundant, introduce irrelevant noise, or even interfere with a previously correct prediction. Our experimental setup thus shifts the focus from overall model capability toward how effectively a system integrates the available modalities, and whether

distinct modality preferences emerge. This perspective motivates the concepts of **Modality Dominance**, **Modality Gain**, **Repair**, **Damage**, and **Cross-Modal Interaction**, formally defined as evaluation metrics in subsequent sections. This framework is rooted in the principle of **controlled information removal** established in *Lost in Transcription* [4], where the underlying task remains fixed while the information accessible to the model is systematically varied, making it possible to evaluate whether an added modality compensates for the absence of another or provides an advantage when partial evidence is already available. That earlier study focused on assessing acoustic contributions within a benchmark built around visual evidence, asking whether the acoustic or textual signal could compensate for absent visual evidence and how this contribution shifted as visual evidence was added. MAIA-AV extends this paradigm through a more articulated manipulation of input modalities. The framework begins with four conditions that establish how much relevant information is retrievable from a single source at a time: complete absence of input, and the isolated acoustic, visual, and transcribed modalities. Two multimodal configurations are then introduced: video with native audio, the full audiovisual clip, and video with audio transcriptions, muted video paired with text. In these multimodal settings, the question extends beyond whether audio merely assists a primarily visual decision. The goal instead is to understand how the availability and representation of information across different modalities reshape the model’s overall reasoning, turning the core question from a simple *does audio help* into a broader investigation of how distinct information sources jointly contribute to multimodal reasoning.

Visual Statement Verification

All experimental conditions adopt the same task formulation, **Visual Statement Verification** (VSV). Each decision is evaluated on a pair consisting of a caption, a correct description, and a foil, a minimally contrastive false description. The two alternatives are presented to the model as a binary choice, and the system must return exclusively 0 or 1. This constrained formulation minimizes the confounds associated with open-ended text generation and enables direct comparison across models and input configurations, since the same experimental unit can be observed under different modality configurations with-

out any change to prompt structure or evaluation criteria. This task structure supports two analytical levels: **Item-Level Decisions**, corresponding to single caption-foil choices, and **Aggregate Accuracy**, which evaluates sets of items linked to the same underlying conceptual category to assess whether the model maintains consistent behavior across different formulations. Both metrics are formally defined in subsequent sections.

Domanda	
Dove cammina la bambina alla fine della scena?	
Caption	Foil
Alla fine della scena, la bambina cammina su un tappeto .	Alla fine della scena, la bambina cammina su un sentiero .
Question	
Where does the little girl walk at the end of the scene?	
Caption	Foil
At the end of the scene, the little girl walks on a carpet .	At the end of the scene, the little girl walks on a path .

Table 3.3: Example of a minimally contrastive caption–foil pair, reported in Italian with its English translation.

Experimental Setup

The protocol is operationalized through a standardized pipeline, with model-specific adapters covering only the interface differences between architectures. This pipeline centralizes item preparation, the definition of experimental conditions, prompt construction, input management, and output storage, so that the same evaluation logic applies uniformly across models. Inputs consist of caption-foil pairs, original videos, muted versions without audio, separate audio files, and precomputed transcripts. The pre-established order of alternatives is preserved for every item, and the target label retains the information needed

to verify decision accuracy. For each item, the raw response generated by the model is preserved alongside the binary label subsequently extracted from it, allowing valid responses to be distinguished from unparsable output and any execution errors to be documented. A lack of support for a given input type is not treated as an incorrect response. When an architecture cannot process a specific modality, the corresponding configuration is marked as unavailable, and the model is excluded only from comparisons requiring that condition. The protocol should therefore be read as a model-by-mode matrix, in which the availability of conditions varies across architectures.

Cross-Architecture Input Standardization The experimental protocol remains uniform across model families. Input delivery is adapted to each architecture’s native multi-modal interface, including the different processors and formats used by Qwen and Gemma, ensuring an identical information payload across architectures without altering the experimental design.

Video Processing Visual content follows a standardized sampling policy: each video is represented using 32 frames evenly distributed across its duration, keeping the amount of visual information comparable across architectures while accommodating the preprocessing requirements of each processor. Visual-only conditions use a version of the video with audio removed, while audiovisual configurations pair the visual stream with the acoustic track specified by the protocol. Sampling is therefore a shared experimental choice, while differences among the visual encoders themselves are addressed separately in the architecture-specific descriptions.

Audio and Transcription Processing The preparation of the acoustic component is tailored to meet the specifications of each experimental condition. In the audio-only scenario, the model is provided with a preprocessed audio file corresponding to the associated video. Conversely, in the audio-visual condition, the original, unmuted video is supplied. When necessary, the model-specific adapter extracts the acoustic stream, converts it into the required format for the corresponding audio frontend, and applies resampling if a 16kHz signal is mandated. Transcriptions adhere to a distinct methodological framework. They

are not generated during the primary experiment, nor do they serve as an intermediate output of the model under evaluation. Instead, they are precomputed inputs maintained consistently across relevant conditions and integrated directly into the textual context. This methodological distinction is crucial, as transcription captures the recognized linguistic content, while the audio signal retains additional acoustic properties. By comparing these two forms of representation, one can study not only the presence of information but also the manner in which it becomes accessible to the model.

Inference Protocol Inference parameters are configured to minimize stochastic variability. Because the task requires a binary classification choice rather than open-ended generation, deterministic decoding is enforced through greedy search, setting `temperature = 0` or `do_sample=False` depending on the model interface, and the maximum output length is strictly constrained to a minimal token budget, limiting non-compliant formatting and keeping outputs consistent across models and conditions. Inference runs in `bfloat16` precision where supported, across execution backends including `vLLM` and `Hugging Face Transformers`. These infrastructure choices preserve the underlying experimental logic, while hardware-dependent parameters, such as GPU allocation and framework-specific workarounds, are detailed in the technical appendix for reproducibility.

Shared Evaluation Infrastructure A small set of components is shared unchanged across every experimental condition, and is described here once rather than repeated for each mode below. Media resources are resolved for each item from its video identifier:

```
1 def paths(row):
2     stem = Path(row["video_name"]).stem
3     return {
4         "audio": AUDIO_DIR / f"{stem}.mp3",
5         "mute": MUTE_DIR / f"{stem}.mp4",
6         "video": VIDEO_DIR / f"{stem}.mp4",
7     }
```

Each condition draws on a different subset of these paths, or on none at all, as detailed

below. Prompts are built centrally, before any request reaches a model adapter:

```
1 def prompt(row, mode, transcript=""):  
2     return PROMPTS[mode].format(  
3         answer1=row["answer1"],  
4         answer2=row["answer2"],  
5         transcript=transcript,  
6     )
```

The transcript argument is populated only for the two conditions that require it:

```
1 transcript = (  
2     transcripts.get(row["video_name"], "")  
3     if mode in (  
4         "only_transcription",  
5         "transcript_video",  
6     )  
7     else ""  
8 )
```

Precomputed transcriptions are loaded once from a dedicated file and indexed by video name:

```
1 TRANSCRIPTIONS = Path(  
2     "data/input/transcription/transcriptions.csv"  
3 )  
4  
5 def transcriptions():  
6     with TRANSCRIPTIONS.open(  
7         encoding="utf-8-sig",  
8         newline=""  
9     ) as f:  
10         return {  
11             r["video_name"]: r["transcription"]  
12             for r in csv.DictReader(f, delimiter=";")
```

```
13     }
```

Which modalities reach the model in a given condition is determined by two flags, defined once for the whole protocol:

```
1 use_video = mode in {
2     "only_video",
3     "video_audio",
4     "transcript_video",
5 }
6
7 use_audio = mode in {
8     "only_audio",
9     "video_audio",
10 }
```

video_audio is the only mode in which both flags are active simultaneously, providing a direct, implementation-level distinction between the joint audiovisual configuration and the unimodal ablations. Every condition delegates inference to the corresponding model adapter through the same call:

```
1 raw = infer(
2     mode,
3     row,
4     prompt(row, mode, transcript),
5     paths(row),
6 )
7 pred = prediction(raw)
```

Regardless of internal multimodal interface, every model is required to return a binary decision. The generated text is normalized through a shared parser that extracts the first isolated occurrence of 0 or 1:

```
1 def prediction(text):
2     m = re.search(r"(?<!\d)[01](?!\\d)", str(text))
```

```

3     return int(m.group()) if m else None

```

If neither label can be recovered, the output is classified as `invalid` rather than mapped automatically to one of the two alternatives. Valid predictions are compared against the stored target through the same rule in every condition:

```

1 result = (
2     "correct" if pred == int(row["target"])
3     else "wrong" if pred is not None
4     else "invalid"
5 )

```

Sampling is disabled throughout, so that inference is deterministic and variation between predictions reflects the information available in each condition rather than sampling noise. For every item, the pipeline logs a fixed set of fields, including item and video metadata, model identifier, experimental condition, raw output, parsed prediction, and correctness status:

```

1 FIELDS = [
2     "id", "video_id", "video_name", "pool_id",
3     "pool_item", "question_category", "question",
4     "caption", "foil", "answer1", "answer2",
5     "target", "model", "mode", "raw_model_output",
6     "predicted_label", "is_correct", "result", "error",
7 ]

```

Results are written to a mode-specific directory with a separate file per model. Existing item identifiers are cached before execution, so that interrupted runs resume without re-computing completed instances, and outputs are flushed immediately after writing to avoid data loss. The system instruction is likewise shared across the entire pipeline, and is not repeated for each condition below:

Sei un assistente binario per un task di Visual Statement Verification.
 Rispondi esclusivamente con 0 oppure 1.²

²English translation: *You are a binary assistant for a Visual Statement Verification task. Respond exclu-*

Model-specific code is accordingly limited to weight loading, chat template formatting, and generation execution. Prompt construction, decision formatting, output parsing, correctness evaluation, error handling, and serialization are identical across models, so that performance differences reflect model capability rather than inconsistencies in the experimental protocol. What follows for each condition is therefore limited to what changes: the modality flags in effect, the media or text actually attached to the request, and the condition-specific instruction.

Zero-input Condition

The no-input condition removes perceptual information entirely, providing the model exclusively with the textual caption-foil pair required to complete the task. No video frames, audio signal, or transcription is supplied. Rather than assessing audiovisual comprehension, its function is strictly diagnostic: it establishes the extent to which decisions can be derived from the linguistic plausibility of the options, structural regularities in the dataset, or systematic parametric biases of the model. Performance well above chance in this condition is a strong indicator of linguistic shortcuts operating independently of genuine content understanding.

Implementation Details Neither flag introduced above is active in this condition, and no transcript is retrieved: the model receives only the two textual alternatives, with the original question, video frames, audio, and transcription entirely omitted. The condition-specific prompt is:

```
Basandoti esclusivamente sulle due alternative testuali, scegli.3
```

```
0: {answer1}
```

```
1: {answer2}
```

```
Rispondi esclusivamente con 0 oppure 1.
```

where `answer1` and `answer2` are the two members of the caption-foil pair. Their numerical ordering carries no semantic meaning, since the dataset’s target label independently

³English translation: *Based exclusively on the two textual alternatives, choose.*

indicates whether the correct caption occupies position 0 or 1, so the model performs a binary selection between competing linguistic realizations rather than open-ended generation. Disabling sampling is especially important here, since it removes output variance unrelated to the underlying linguistic representations. Results are written to a dedicated `no_input` directory, one file per model, following the shared logging and resumability scheme described above.

Audio-only Condition

In the audio-only condition, the model receives exclusively the raw acoustic signal, with no visual frames or transcription provided. This isolates which caption-foil pairs can be resolved from spoken content, environmental sound, music, or prosody, and comparing performance here against other configurations reveals whether acoustic information retains independent utility once combined with vision, or is instead diminished or overshadowed by a dominant modality.

Implementation Details Only the `use_audio` flag is active: the model receives the audio resolved via `paths(row) ["audio"]` together with the two textual alternatives, while the original question is kept only as dataset metadata and no transcript is retrieved. The condition-specific prompt is:

```
Ascolta esclusivamente l'audio fornito.  
Scegli quale delle due descrizioni è corretta sulla base delle  
informazioni ricavabili esclusivamente dall'audio.  
  
0: {answer1}  
1: {answer2}  
  
Rispondi esclusivamente con 0 oppure 1.4
```

How this resource reaches the model is architecture-dependent: models that accept audio files directly receive it through their native processor, while architectures requiring raw

⁴English translation: *Listen exclusively to the provided audio. Choose which of the two descriptions is correct based on the information available exclusively from the audio. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

waveforms decode the signal into the format expected by their inference backend. Either way, this concerns only input formatting and introduces no textual or visual information beyond what the shared infrastructure already provides. Results are written to a dedicated `only_audio` directory, one file per model. This condition differs from the no-input baseline solely in the availability of the acoustic signal, so the comparison between them isolates whether speech, environmental sound, music, or prosody carry enough information to discriminate captions from foils, independently of visual evidence or its textual transcription.

Video-only Condition

In the video-only condition, the model receives exclusively the visual stream, with the audio track removed. This is the primary baseline for evaluating which decisions can be sustained by visual evidence alone, and it plays a central role in every paired comparison that follows: introducing audio or transcription afterward measures the marginal contribution of that source relative to a baseline where visual context is already available.

Implementation Details Only the `use_video` flag is active. To guarantee that no acoustic information reaches the model, this condition uses the muted video resolved via `paths(row) ["mute"]` together with the two textual alternatives and no transcript. The condition-specific prompt is:

```
Osserva il video fornito.  
Scegli quale delle due descrizioni è corretta sulla base  
esclusivamente delle informazioni visive disponibili.  
  
0: {answer1}  
1: {answer2}  
  
Rispondi esclusivamente con 0 oppure 1.5
```

⁵English translation: *Observe the provided video. Choose which of the two descriptions is correct based exclusively on the available visual information. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

The video is represented with the standard budget of 32 frames, sampled at uniformly spaced indices across its full duration rather than from a single local interval:

```
1 N_FRAMES = 32
2
3 frame_indices = np.linspace(
4     0,
5     total_frames - 1,
6     min(N_FRAMES, total_frames),
7     dtype=int,
8 )
```

Videos with fewer than 32 frames retain all available frames. Extracted frames are converted to RGB before being passed to the multimodal processor, and depending on the native interface of the architecture, the resulting sequence is represented either as a video input or as an ordered sequence of images, a difference confined to the model-specific adapter that does not alter the visual evidence available in the experiment. Results are written to a dedicated `only_video` directory, one file per model. Comparison with the no-input condition quantifies the contribution of visual grounding on its own, while comparison with the audiovisual conditions establishes the reference point needed to determine whether additional modalities improve, preserve, or interfere with visually grounded predictions.

Transcription-only Condition

In the transcription-only condition, the model receives exclusively the textual transcription of the audio track. This isolates spoken linguistic content from non-verbal acoustic features such as prosody, speaker characteristics, background noise, and music, evaluating performance driven strictly by explicit verbal information.

Implementation Details Neither modality flag is active for this condition. The model receives the transcript retrieved via `transcriptions()` together with the two textual alternatives. No video frames, video representation, or raw audio signal are provided, and

the original question is kept only as dataset metadata. The condition-specific prompt is:

Leggi esclusivamente la trascrizione dell'audio fornita.
Scegli quale delle due descrizioni è corretta sulla base
esclusivamente delle informazioni testuali disponibili nella
trascrizione.

Trascrizione dell'audio: {transcript}

0: {answer1}

1: {answer2}

Rispondi esclusivamente con 0 oppure 1.⁶

This condition differs from the no-input setting in a single controlled respect: the caption-foil pair is now accompanied by the textual representation of the audio track, so the comparison between them isolates whether speech-derived linguistic information contributes to the decision even without direct perceptual access to the original video or acoustic signal. Because no multimodal resource is attached, the transcription is tokenized as ordinary text alongside the task instruction and the two alternatives, without ever passing through an audio encoder. This is central to the condition's design: it evaluates access to the semantic content extracted from audio rather than to the acoustic modality itself, so prosody, speaker characteristics, environmental sound, and other non-textual properties remain unavailable unless they were explicitly captured in the transcription. Results are written to a dedicated `only_transcription` directory, one file per model. Comparison with the no-input condition measures the contribution of the transcription itself, comparison with the audio-only condition distinguishes the semantic information preserved by textualization from what is available in the original acoustic signal, and comparison with the conditions that include visual evidence tests whether the same linguistic information remains useful once combined with the primary visual channel.

⁶English translation: *Read exclusively the provided audio transcription. Choose which of the two descriptions is correct based exclusively on the textual information available in the transcription. Audio transcription: {transcript}. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

The Audio-Visual (AV) Condition

The audiovisual condition represents the natural configuration of information available to the model, comprising both the visual sequence and the native acoustic signal from the original video, provided as a single integrated stream rather than as the separately combined channels of the video-plus-transcription condition described below. This is the core configuration for testing whether the simultaneous presence of multiple modalities yields an advantage over unimodal conditions or instead leads to *unimodal collapse*. Its value extends beyond aggregate performance, since it also allows item-level observation of whether the acoustic channel repairs errors made in the video-only condition or introduces new incorrect decisions.

Implementation Details Both `use_video` and `use_audio` are active here, the only condition in which both are true simultaneously. The model receives the original, non-muted video resolved via `paths(row) ["video"]`, ensuring both modalities originate from the same source stimulus, together with the two textual alternatives and no transcript. The condition-specific prompt is:

```
Osserva il video e ascolta il relativo audio.  
Scegli quale delle due descrizioni è corretta sulla base  
dell'insieme delle informazioni visive e acustiche disponibili.  
  
0: {answer1}  
1: {answer2}  
  
Rispondi esclusivamente con 0 oppure 1.7
```

When an architecture requires the acoustic stream as a separate input, the waveform is extracted directly from the same original file used for the visual channel, rather than from an independent recording, preserving the correspondence between the two modalities at the stimulus level:

⁷English translation: *Observe the video and listen to its audio. Choose which of the two descriptions is correct based on the combined visual and acoustic information available. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

```

1 def extract_audio(video):
2     out = (
3         Path(tempfile.gettempdir())
4         / "maia_av_audio_cache"
5         / f"{video.stem}.wav"
6     )
7
8     out.parent.mkdir(exist_ok=True)
9
10    if not out.exists():
11        subprocess.run(
12            [
13                "ffmpeg",
14                "-loglevel", "error",
15                "-y",
16                "-i", str(video),
17                "-vn",
18                "-ac", "1",
19                "-ar", "16000",
20                str(out),
21            ],
22            check=True,
23        )
24
25    return out

```

producing a mono waveform at 16kHz. The video itself follows the same 32-frame uniform sampling used in the video-only condition, with indices additionally deduplicated:

```

1 N_FRAMES = 32
2
3 frame_indices = np.unique(

```

```

4     np.linspace(
5         0,
6         total_frames - 1,
7         min(N_FRAMES, total_frames),
8         dtype=int,
9     )
10 )

```

Both streams are then exposed to the inference backend as explicit multimodal inputs where the adapter requires it, or recovered directly from the original video by adapters offering native audiovisual processing, a difference confined to the model-specific adapter that does not affect the evidence available during inference:

```

1 multi_modal_data = {}
2
3 if video is not None:
4     multi_modal_data["video"] = video
5
6 if audio is not None:
7     multi_modal_data["audio"] = (
8         audio,
9         AUDIO_SAMPLE_RATE,
10    )
11
12 request["multi_modal_data"] = multi_modal_data

```

Results are written to a dedicated `video_audio` directory, one file per model. This condition is best understood not as an independent task but as the joint-modality counterpart of the two unimodal ablations: comparison with video-only provides a controlled measure of whether adding raw acoustic information improves, preserves, or interferes with decisions already supported by vision, while comparison with the transcription-based multimodal condition below distinguishes the contribution of native acoustic processing from that of the linguistic information preserved in a transcription.

The Visual-Language (VL) Configuration

The video-plus-transcription condition combines the visual component with a textual representation of speech. This is the second multimodal configuration, allowing direct audio-visual integration to be compared against textualized linguistic information: comparison with the audiovisual condition tests whether the informative content of the acoustic track is exploited more effectively once converted to text, while comparison with video-only clarifies the benefit of audio-derived linguistic information when visual context is already available.

Implementation Details Only `use_video` is active, since this condition does not expose the raw acoustic signal. The model receives the muted video resolved via `paths(row) ["mute"]`, together with the transcript retrieved via `transcriptions()` and the two textual alternatives. The condition-specific prompt combines the visual instruction with the transcription:

```
Osserva il video fornito e leggi la trascrizione dell'audio.
Scegli quale delle due descrizioni è corretta sulla base
dell'insieme delle informazioni visive e testuali disponibili.

Trascrizione dell'audio: {transcript}

0: {answer1}
1: {answer2}

Rispondi esclusivamente con 0 oppure 1.8
```

The video follows the same 32-frame uniform sampling described for the video-only condition. No audio extraction takes place here. Since only the transcription of the acoustic channel is exposed to the model, never the raw signal itself, the muted video is sufficient and the native waveform is never decoded. Results are written to a dedicated `transcript_video` directory, one file per model. Because this condition and the audiovisual condition differ only in how the acoustic content reaches the model, native waveform

⁸English translation: *Observe the provided video and read the audio transcription. Choose which of the two descriptions is correct based on the combined visual and textual information available. Audio transcription: {transcript}. 0: {answer1}. 1: {answer2}. Respond exclusively with either 0 or 1.*

versus transcription, while the visual channel and the underlying decision remain identical, their comparison isolates the effect of textualization itself on top of an already visually grounded prediction.

Metrics Selection and Aggregation Strategy

Linguistic and Structural Bias Assessment under Zero-Input

The no-input condition calls for different analytical treatment than the other experimental settings. Since neither visual nor acoustic information is available, performance in this condition cannot be interpreted as evidence of multimodal understanding, and is therefore unsuitable for direct inclusion in standard performance comparisons. Instead, dedicated metrics are used to quantify **linguistic and structural shortcuts**, that is, whether the model can exploit linguistic regularities, structural properties of the alternatives, or systematic biases in the dataset to solve the task without perceptual evidence. Let N denote the number of evaluated items. For each item i , let $y_i \in \{0, 1\}$ be the target label and $\hat{y}_i \in \{0, 1\}$ the prediction produced by the model.

Raw Accuracy The first measure is *Raw Accuracy*, the proportion of items for which the predicted label matches the target:

$$\text{Acc}_{\text{NoInput}} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = y_i), \quad (3.1)$$

where $\mathbb{I}(\cdot)$ is the indicator function, equal to 1 when the condition holds and 0 otherwise. Since each item requires a binary choice between two alternatives, a model selecting uniformly at random has expected accuracy

$$\mathbb{E}[\text{Acc}_{\text{random}}] = \frac{1}{2} = 0.5. \quad (3.2)$$

An accuracy close to 0.5 is therefore compatible with random guessing, while a substantially higher value indicates that the model identifies the target more often than chance would predict despite the absence of perceptual evidence, suggesting that some property of the alternatives' linguistic form may be predictive of the correct

answer. Raw Accuracy should not be read as a measure of language comprehension: the model is not required to understand the underlying video event, nor does it have access to the evidence needed for grounded reasoning. The metric only captures its ability to recover the target from the two textual alternatives available in this condition.

Language-Only Advantage To express the deviation from the random baseline more directly, we compute the *Language-Only Advantage* (LOA), obtained by subtracting the expected chance level from the observed no-input accuracy:

$$\text{LOA} = \text{Acc}_{\text{NoInput}} - 0.5. \quad (3.3)$$

LOA = 0 indicates performance exactly at the random baseline, positive values indicate performance above chance, and negative values indicate performance below chance. The main advantage of this transformation is interpretability: rather than reporting the raw accuracy value, LOA directly expresses the amount of performance exceeding what chance would predict, and can be used as a compact indicator of the potential contribution of linguistic or structural shortcuts. A positive LOA does not identify the nature of the shortcut. It only establishes that the textual alternatives contain sufficient regularities for the model to perform above chance, so additional diagnostics are required to determine which properties may contribute to this effect.

Matthews Correlation Coefficient Accuracy alone can conceal degenerate decision strategies, since a model with a systematic preference for one output label may still obtain non-negligible accuracy when predictions are unevenly distributed across the two positions. For this reason, the analysis also includes the *Matthews Correlation Coefficient* (MCC). For a binary task, predictions can be summarized through the four entries of the confusion matrix, true positives (TP), true negatives (TN), false positives (FP), and false negatives (FN):

$$\text{MCC} = \frac{\text{TP} \cdot \text{TN} - \text{FP} \cdot \text{FN}}{\sqrt{(\text{TP} + \text{FP})(\text{TP} + \text{FN})(\text{TN} + \text{FP})(\text{TN} + \text{FN})}}. \quad (3.4)$$

The numerator compares concordant and discordant predictions, correct decisions contributing through $\text{TP} \cdot \text{TN}$ and the two error types through $\text{FP} \cdot \text{FN}$, while the

denominator normalizes this difference against the marginal distributions of targets and predictions. MCC ranges between -1 and 1 . Values close to 1 indicate strong agreement between predictions and targets, values close to 0 indicate little or no association, and negative values indicate a systematic inverse association. MCC complements Raw Accuracy by weighing all four components of the confusion matrix rather than only the diagonal, distinguishing genuine target-sensitive performance from apparently successful behavior produced by strongly unbalanced prediction strategies.

Position Bias *Position Bias* measures directly whether the model systematically favors one candidate position independently of its content. Let

$$p_0 = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = 0) \quad (3.5)$$

denote the empirical probability of selecting the first alternative, with $p_1 = 1 - p_0$ the probability of selecting the second. Under perfectly balanced behavior, $p_0 = p_1 = 0.5$, so we define the magnitude of the positional deviation as

$$\text{PB} = |p_0 - 0.5|, \quad (3.6)$$

and the preferred position as

$$\text{PreferredPosition} = \underset{j \in \{0,1\}}{\operatorname{argmax}} p_j. \quad (3.7)$$

A Position Bias of 0 indicates perfectly balanced selection, and the maximum possible value, 0.5 , is reached when the model always selects the same position. This diagnostic matters because above-chance Raw Accuracy does not by itself demonstrate the exploitation of meaningful linguistic information: part of the observed result may instead reflect a systematic preference for the first or second alternative, which Position Bias measures directly.

Lexical Attraction *Lexical Attraction* evaluates whether the model preferentially selects the alternative that is lexically closer to the original question. For each item i , let q_i

denote the question and a_{i0}, a_{i1} the two alternatives, each represented as a TF-IDF vector. The cosine similarity between the question and alternative j is

$$s_{ij} = \cos(\mathbf{v}_{q_i}, \mathbf{v}_{a_{ij}}) = \frac{\mathbf{v}_{q_i}^\top \mathbf{v}_{a_{ij}}}{\|\mathbf{v}_{q_i}\|_2 \|\mathbf{v}_{a_{ij}}\|_2}, \quad j \in \{0, 1\}, \quad (3.8)$$

and the lexically closest alternative is

$$\ell_i = \underset{j \in \{0, 1\}}{\operatorname{argmax}} s_{ij}. \quad (3.9)$$

The *Lexical Attraction Rate* measures how often the model selects this alternative:

$$\text{LAR} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = \ell_i). \quad (3.10)$$

A high value indicates a preference for the candidate sharing greater lexical overlap with the question, though this behavioral tendency alone does not establish that lexical overlap is a dataset bias. To evaluate the structural relationship between lexical similarity and the target, we separately compute the *Lexical Target Alignment*:

$$\text{LTA} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(y_i = \ell_i), \quad (3.11)$$

the frequency with which the correct alternative is also the one with the highest lexical similarity to the question. The two measures capture different phenomena, LAR reflecting model behavior and LTA reflecting dataset structure:

$$\text{LAR} \rightarrow \text{model behavior}, \quad (3.12)$$

$$\text{LTA} \rightarrow \text{dataset structure}. \quad (3.13)$$

A high LAR is only informative as a shortcut if lexical similarity is genuinely correlated with the target:

$$\text{LAR} \gg 0.5 \quad \text{and} \quad \text{LTA} \gg 0.5 \quad (3.14)$$

jointly provide evidence for both the availability of a lexical shortcut in the dataset and the model's tendency to exploit it.

Length Bias A parallel analysis applies to the length of the two alternatives. Let $L(a_{ij})$ denote the length of alternative j for item i , and let

$$h_i = \operatorname{argmax}_{j \in \{0,1\}} L(a_{ij}) \quad (3.15)$$

identify the longer candidate. The model’s preference for longer alternatives is the *Length Attraction Rate*:

$$\text{LenAR} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = h_i), \quad (3.16)$$

a behavioral property of the model, with values well above 0.5 indicating a tendency to select the longer alternative. As with lexical similarity, a separate measure establishes whether this heuristic is actually rewarded by the dataset, the *Length Target Alignment*:

$$\text{LenTA} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(y_i = h_i), \quad (3.17)$$

the frequency with which the correct alternative is also the longer one. The distinction again separates model behavior from dataset structure:

$$\text{LenAR} \rightarrow \text{preference of the model for longer alternatives}, \quad (3.18)$$

$$\text{LenTA} \rightarrow \text{association between alternative length and the target}. \quad (3.19)$$

A model may prefer longer alternatives even when length carries no information about correctness, just as the dataset may contain a length-correctness correlation that a given model fails to exploit. Considering both quantities together keeps these two cases from being conflated.

No-Input Category Breakdown The main no-input diagnostics are additionally computed separately for each reasoning category. Let c denote a category and

$$I_c = \{i \mid \text{category}(i) = c\} \quad (3.20)$$

the subset of items belonging to it, with $N_c = |I_c|$. Category-specific Raw Accuracy and Language-Only Advantage follow directly:

$$\text{Acc}_c = \frac{1}{N_c} \sum_{i \in I_c} \mathbb{I}(\hat{y}_i = y_i), \quad \text{LOA}_c = \text{Acc}_c - 0.5, \quad (3.21)$$

and MCC and Position Bias are computed analogously by restricting the confusion matrix and prediction frequencies to I_c , with

$$p_{0,c} = \frac{1}{N_c} \sum_{i \in I_c} \mathbb{I}(\hat{y}_i = 0), \quad \text{PB}_c = |p_{0,c} - 0.5|. \quad (3.22)$$

This decomposition makes it possible to determine whether no-input solvability varies systematically across reasoning categories, without treating a higher category-level accuracy as evidence that the model is more competent at the corresponding reasoning process: since no perceptual evidence is available in this condition, such differences instead indicate that some groups of items expose stronger linguistic or structural regularities than others. This distinction matters directly when no-input results are later compared against performance under perceptual conditions, since differences observed between spatial, temporal, and causal questions should not be attributed automatically to multimodal reasoning ability when part of the variation may already be present before any perceptual modality is introduced. The category breakdown therefore provides the baseline needed to separate genuine grounding effects from differences rooted in the linguistic structure of the task.

Item-Aligned Modality Influence Metrics

Let c denote an experimental condition and $C_i^{(c)} \in \{0, 1\}$ indicate whether the model answers item i correctly under condition c :

$$C_i^{(c)} = \mathbb{I}(\hat{y}_i^{(c)} = y_i), \quad (3.23)$$

where y_i is the target label and $\hat{y}_i^{(c)}$ the prediction obtained under c . For every comparison between a baseline condition b and a condition n obtained by adding a new modality, only items available in both are retained, giving the paired evaluation set

$$\mathcal{I}_{b,n} = \mathcal{I}_b \cap \mathcal{I}_n, \quad N_{b,n} = |\mathcal{I}_{b,n}|. \quad (3.24)$$

All modality-contribution metrics are computed on this common subset, which preserves item-level comparability and keeps differences in data availability from affecting the estimated contribution of the added modality.

Conditional Modality Gain The *Conditional Modality Gain* measures the change in performance produced by adding a new source of information to a fixed baseline, for the two primary comparisons

$$\text{Video Only} \rightarrow \text{Audio} + \text{Video} \quad \text{and} \quad \text{Video Only} \rightarrow \text{Video} + \text{Transcription}. \quad (3.25)$$

Let b denote the baseline condition and n the condition after the additional modality has been introduced. Accuracy on the paired subset is

$$\text{Acc}_b = \frac{1}{N_{b,n}} \sum_{i \in \mathcal{J}_{b,n}} C_i^{(b)}, \quad \text{Acc}_n = \frac{1}{N_{b,n}} \sum_{i \in \mathcal{J}_{b,n}} C_i^{(n)}, \quad (3.26)$$

and the Conditional Modality Gain is their difference, which, since both accuracies are computed on the same items, can equivalently be written as an item-level average:

$$\Delta_{b \rightarrow n} = \text{Acc}_n - \text{Acc}_b = \frac{1}{N_{b,n}} \sum_{i \in \mathcal{J}_{b,n}} \left(C_i^{(n)} - C_i^{(b)} \right). \quad (3.27)$$

This second formulation makes the paired nature of the metric explicit. Each item contributes $+1$ when the additional modality turns an incorrect response into a correct one, -1 when it turns a correct response into an incorrect one, and 0 when correctness is unchanged, so $\Delta_{b \rightarrow n} > 0$ indicates a positive net contribution of the added modality and $\Delta_{b \rightarrow n} < 0$ an overall detrimental effect. For the two main comparisons, the corresponding gains are

$$\Delta_{\text{Audio}|\text{Video}} = \text{Acc}_{\text{Audio}+\text{Video}} - \text{Acc}_{\text{VideoOnly}}, \quad \Delta_{\text{Transc.}|\text{Video}} = \text{Acc}_{\text{Video}+\text{Transc.}} - \text{Acc}_{\text{VideoOnly}}. \quad (3.28)$$

These quantities measure the marginal contribution of acoustic or transcribed information when visual evidence is already available, directly extending the controlled information-addition principle adopted in *Lost in Transcription*.

Repair Rate Aggregate gain does not reveal how individual decisions change. The *Repair Rate* isolates items answered incorrectly in the baseline condition and measures how often the additional modality converts them into correct predictions:

$$N_{\text{repair}} = \sum_{i \in \mathcal{J}_{b,n}} \mathbb{I} \left(C_i^{(b)} = 0 \wedge C_i^{(n)} = 1 \right), \quad N_{\text{base error}} = \sum_{i \in \mathcal{J}_{b,n}} \mathbb{I} \left(C_i^{(b)} = 0 \right), \quad (3.29)$$

$$\text{RepairRate}_{b \rightarrow n} = \frac{N_{\text{repair}}}{N_{\text{base error}}} = P\left(C^{(n)} = 1 \mid C^{(b)} = 0\right). \quad (3.30)$$

The metric conditions exclusively on baseline failures, measuring the corrective contribution of the newly available modality rather than its effect on the complete item set. For audio added to visual evidence,

$$\text{RepairRate}_{A|V} = P\left(C^{(AV)} = 1 \mid C^{(V)} = 0\right), \quad (3.31)$$

where a high value indicates that acoustic evidence frequently resolves decisions that cannot be solved from visual evidence alone.

Damage Rate The *Damage Rate* captures the opposite transition, the proportion of baseline-correct items that become incorrect once the additional modality is introduced:

$$N_{\text{damage}} = \sum_{i \in \mathcal{I}_{b,n}} \mathbb{I}\left(C_i^{(b)} = 1 \wedge C_i^{(n)} = 0\right), \quad N_{\text{base correct}} = \sum_{i \in \mathcal{I}_{b,n}} \mathbb{I}\left(C_i^{(b)} = 1\right), \quad (3.32)$$

$$\text{DamageRate}_{b \rightarrow n} = \frac{N_{\text{damage}}}{N_{\text{base correct}}} = P\left(C^{(n)} = 0 \mid C^{(b)} = 1\right). \quad (3.33)$$

For the addition of audio to video,

$$\text{DamageRate}_{A|V} = P\left(C^{(AV)} = 0 \mid C^{(V)} = 1\right), \quad (3.34)$$

a transition that provides a direct behavioral measure of cross-modal interference: the additional source does not merely fail to contribute useful information, it actively overturns a previously correct decision. Conditional Modality Gain alone cannot distinguish repair from damage. When $N_{\text{repair}} = N_{\text{damage}}$, the net gain $\Delta_{b \rightarrow n} = 0$ even though a substantial fraction of individual predictions has changed, a situation fundamentally different from $N_{\text{repair}} = N_{\text{damage}} = 0$, where the additional modality produces no change in correctness at all. Paired decisions are accordingly partitioned into four mutually exclusive transition classes:

$$\begin{aligned} N_{\text{stable correct}} &= \sum_i \mathbb{I}\left(C_i^{(b)} = 1 \wedge C_i^{(n)} = 1\right), \\ N_{\text{repair}} &= \sum_i \mathbb{I}\left(C_i^{(b)} = 0 \wedge C_i^{(n)} = 1\right), \\ N_{\text{damage}} &= \sum_i \mathbb{I}\left(C_i^{(b)} = 1 \wedge C_i^{(n)} = 0\right), \\ N_{\text{stable wrong}} &= \sum_i \mathbb{I}\left(C_i^{(b)} = 0 \wedge C_i^{(n)} = 0\right), \end{aligned} \quad (3.35)$$

satisfying

$$N_{b,n} = N_{\text{stable correct}} + N_{\text{repair}} + N_{\text{damage}} + N_{\text{stable wrong}}, \quad (3.36)$$

the item-level decomposition underlying the aggregate gain.

Modality Interdependence Analysis

Factorial Modality Interaction Analysis The joint contribution of audio and video is analyzed through a 2×2 factorial design, with $A \in \{0, 1\}$ indicating the availability of the acoustic modality and $V \in \{0, 1\}$ the availability of the visual modality:

	$V = 0$	$V = 1$
$A = 0$	No Input	Video Only
$A = 1$	Audio Only	Audio + Video

(3.37)

The *No Input* condition is included solely to complete the factorial structure, representing the cell $(A, V) = (0, 0)$ in which neither perceptual modality is present, without implying that it constitutes a multimodal comprehension condition. Let $Y_i \in \{0, 1\}$ indicate whether item i is answered correctly. Correctness probability is modeled through logistic regression,

$$\text{logit}[P(Y_i = 1)] = \beta_0 + \beta_A A_i + \beta_V V_i + \beta_{AV} (A_i V_i), \quad \text{logit}(p) = \log\left(\frac{p}{1-p}\right). \quad (3.38)$$

Writing $p_{av} = P(Y = 1 \mid A = a, V = v)$ for the four cell probabilities, the coefficients admit a direct interpretation:

$$\beta_0 = \text{logit}(p_{00}), \quad \beta_A = \text{logit}(p_{10}) - \text{logit}(p_{00}), \quad \beta_V = \text{logit}(p_{01}) - \text{logit}(p_{00}), \quad (3.39)$$

$$\begin{aligned} \beta_{AV} &= \text{logit}(p_{11}) - \text{logit}(p_{10}) - \text{logit}(p_{01}) + \text{logit}(p_{00}) \\ &= [\text{logit}(p_{11}) - \text{logit}(p_{01})] - [\text{logit}(p_{10}) - \text{logit}(p_{00})]. \end{aligned} \quad (3.40)$$

The interaction thus measures whether the effect of audio changes as a function of visual availability. Under the null hypothesis $H_0 : \beta_{AV} = 0$, the joint audiovisual configuration is compatible with additive modality effects on the logit scale, while $\beta_{AV} > 0$ indicates a joint effect stronger than the separate modality contributions

predict and $\beta_{AV} < 0$ a joint effect weaker than that additive expectation. On the odds-ratio scale, $OR_{AV} = \exp(\beta_{AV})$, so that $OR_{AV} > 1$ corresponds to positive interaction and $OR_{AV} < 1$ to negative interaction. For the estimated coefficient $\hat{\beta}_{AV}$, the Wald statistic

$$z = \frac{\hat{\beta}_{AV}}{SE(\hat{\beta}_{AV})} \quad (3.41)$$

yields the two-sided p -value $p = 2[1 - \Phi(|z|)]$, where $\Phi(\cdot)$ is the cumulative distribution function of the standard normal distribution. The 95% confidence interval for the interaction coefficient is

$$CI_{95\%}(\beta_{AV}) = \hat{\beta}_{AV} \pm 1.96 SE(\hat{\beta}_{AV}), \quad (3.42)$$

and its odds-ratio counterpart is obtained through exponentiation,

$$CI_{95\%}(OR_{AV}) = \left[\exp\left(\hat{\beta}_{AV} - 1.96 SE(\hat{\beta}_{AV})\right), \exp\left(\hat{\beta}_{AV} + 1.96 SE(\hat{\beta}_{AV})\right) \right]. \quad (3.43)$$

Standard errors are estimated using a cluster-robust covariance matrix at the `pool_id` level, letting observations within the same pool be statistically dependent while treating different pools as independent clusters. Only items observed in all four factorial cells are retained, restricting the analysis to

$$\mathcal{I}_{2 \times 2} = \mathcal{I}_{00} \cap \mathcal{I}_{10} \cap \mathcal{I}_{01} \cap \mathcal{I}_{11}, \quad (3.44)$$

a complete-case restriction that preserves the paired structure required for a valid comparison of modality effects.

Joint-Modality Evaluation

Joint-Modality Accuracy Aggregate audiovisual accuracy does not distinguish items that genuinely require multiple modalities from items already solvable from a single source. The *Joint-Modality Accuracy* restricts evaluation to items classified as requiring a specific combination of modalities. Let $J_i^{AV} \in \{0, 1\}$ indicate whether item i requires both audio and visual information, with $\mathcal{I}_{AV} = \{i \mid J_i^{AV} = 1\}$. Audio–

Video Joint-Modality Accuracy is then

$$\text{JMA}_{AV} = \frac{\sum_{i \in \mathcal{J}_{AV}} C_i^{(AV)}}{|\mathcal{J}_{AV}|}. \quad (3.45)$$

For items requiring visual information together with an audio-derived transcription, the analogous definitions $J_i^{VT} \in \{0, 1\}$ and $\mathcal{J}_{VT} = \{i \mid J_i^{VT} = 1\}$ give

$$\text{JMA}_{VT} = \frac{\sum_{i \in \mathcal{J}_{VT}} C_i^{(VT)}}{|\mathcal{J}_{VT}|}. \quad (3.46)$$

This restriction changes the object being measured. General audiovisual accuracy is $\text{Acc}_{AV} = P(C^{(AV)} = 1)$, whereas Joint-Modality Accuracy is $\text{JMA}_{AV} = P(C^{(AV)} = 1 \mid J^{AV} = 1)$. High performance on the first can arise from items solvable through a single modality, while high performance on the second requires successful predictions specifically within the subset for which joint information has been identified as necessary, making it a more selective estimate of multimodal integration.

Alignment Sensitivity

Alignment Sensitivity measures whether model performance depends on the correct correspondence between information sources. Let a denote the correctly aligned multimodal condition and m the corresponding misaligned condition, with evaluation restricted to $\mathcal{J}_{a,m} = \mathcal{J}_a \cap \mathcal{J}_m$. Aligned and misaligned accuracies are

$$\text{Acc}_{\text{aligned}} = \frac{1}{|\mathcal{J}_{a,m}|} \sum_{i \in \mathcal{J}_{a,m}} C_i^{(a)}, \quad \text{Acc}_{\text{misaligned}} = \frac{1}{|\mathcal{J}_{a,m}|} \sum_{i \in \mathcal{J}_{a,m}} C_i^{(m)}, \quad (3.47)$$

and Alignment Sensitivity is their difference which, since both conditions are evaluated on the same items, can equivalently be written item by item:

$$\text{AS} = \text{Acc}_{\text{aligned}} - \text{Acc}_{\text{misaligned}} = \frac{1}{|\mathcal{J}_{a,m}|} \sum_{i \in \mathcal{J}_{a,m}} (C_i^{(a)} - C_i^{(m)}). \quad (3.48)$$

A positive value, $\text{AS} > 0$, indicates that correct cross-modal correspondence improves performance relative to mismatched inputs, while a value near zero indicates limited sensitivity to the relationship between the modalities. For audiovisual inputs,

$$\text{AS}_{AV} = \text{Acc}_{AV}^{\text{aligned}} - \text{Acc}_{AV}^{\text{misaligned}}, \quad (3.49)$$

and for video and transcription,

$$AS_{VT} = \text{Acc}_{VT}^{\text{aligned}} - \text{Acc}_{VT}^{\text{misaligned}}. \quad (3.50)$$

This metric enters the main experimental results only when the corresponding misaligned conditions are part of the executed protocol. If misaligned configurations are supported only by the evaluation infrastructure, AS remains a formally defined extension rather than an observed main-experiment metric.

Stratified Behavioral Analysis

Independent Visual-Evidence Estimation The preliminary analysis provides an estimate of visual evidence availability that is independent of the VSV predictions produced in the main experiments.

For each item i , let R_i^V denote the structured representation extracted exclusively from the visual stream, obtained from the analysis of entities, events, and inferential relations, and let Q_i^* denote the information required to resolve the corresponding question. A Natural Language Inference stage compares R_i^V with Q_i^* and produces an evidence-availability variable

$$E_i^V = f_{\text{NLI}}(R_i^V, Q_i^*) \in \{0, 1\}, \quad (3.51)$$

where $E_i^V = 1$ indicates that the visual representation contains sufficient information according to the adopted NLI criterion, and $E_i^V = 0$ indicates that the required information cannot be established from it.

This variable is not defined from VSV correctness, $E_i^V \neq C_i^{(V)}$: the two quantities sit at distinct analytical levels. E_i^V estimates whether the required information is available in the independently extracted visual representation, while $C_i^{(V)}$ records whether the model actually produces the correct VSV decision under the video-only condition. Evidence availability, in short, is not the same thing as observed task success.

Mapping Preliminary and Experimental Outcomes Preliminary and main-experiment outputs are joined through their shared item identifiers. For each item i , the resulting

analytical representation contains at least the triplet

$$\left(E_i^V, C_i^{(N)}, C_i^{(m)}\right), \quad (3.52)$$

where E_i^V is the independently estimated visual-evidence availability, $C_i^{(N)}$ correctness in the *No Input* condition, and $C_i^{(m)}$ correctness under experimental modality configuration m . The complete modality vector can be written as

$$\mathbf{C}_i = \left(C_i^{(N)}, C_i^{(A)}, C_i^{(V)}, C_i^{(AV)}, C_i^{(T)}, C_i^{(VT)}\right), \quad (3.53)$$

with N, A, V, AV, T , and VT denoting *No Input*, *Audio Only*, *Video Only*, *Audio + Video*, *Transcription Only*, and *Video + Transcription*, respectively.

The preliminary taxonomy assigns each item to an independently defined stratum $S_i \in \{S_0, S_1, \dots, S_K\}$, not derived from main-experiment correctness, $S_i \neq g(\mathbf{C}_i)$, but acting as an external stratification variable. Performance under condition m can therefore be computed separately for each class S_k ,

$$\text{Acc}_{m, S_k} = \frac{\sum_i \mathbb{I}(S_i = S_k) C_i^{(m)}}{\sum_i \mathbb{I}(S_i = S_k)}, \quad (3.54)$$

with the corresponding difference between two modality conditions b and n within the same stratum given by

$$\Delta_{b \rightarrow n}^{(S_k)} = \text{Acc}_{n, S_k} - \text{Acc}_{b, S_k}. \quad (3.55)$$

This stratification lets experimental behavior be interpreted jointly with independently estimated evidence availability. The pair $(E_i^V = 0, C_i^{(V)} = 1)$ identifies an item solved correctly under the video-only condition despite the absence of sufficient evidence in the independently extracted visual representation, while $(E_i^V = 1, C_i^{(V)} = 0)$ identifies a failure despite independently estimated visual evidence being available. The additional No-Input variable further separates perceptual from non-perceptual success. $(E_i^V = 0, C_i^{(N)} = 1)$ is compatible with a decision supported by linguistic or structural regularities rather than by visual grounding, while $(E_i^V = 1, C_i^{(N)} = 0, C_i^{(V)} = 1)$ is consistent with a performance gain that emerges only once visual evidence becomes available. The preliminary classes $S_0, S_1, \dots, S_K$ thus provide an independent

stratification layer for the modality-ablation results. Their role is analytical rather than evaluative: they allow observed VSV decisions to be interpreted against the information independently recoverable from the visual stream, without defining the classes from the same predictions that are subsequently analyzed against them.

3.4 Attention-Based Modality Analysis

The attention analysis complements the behavioural experiments described above. Rather than asking only whether a modality’s presence affects performance, it investigates how the model internally allocates attention across the available sources of information during generation. The design follows directly from the findings discussed earlier for *Lost in Transcription* [4]: the availability of multiple modalities does not by itself guarantee their effective integration, and behavioural comparisons alone cannot show how a model distributes its internal processing when several sources are available at once. A model may receive both video and audio while attending almost entirely to the visual representations, or a transcription may attract substantial decoder attention even where its addition barely changes task accuracy. For this reason, the attention analysis provides an internal counterpart to the behavioural modality-ablation experiments. The behavioural analysis addresses the question:

Does the availability of a given modality change model performance?

The attention analysis instead addresses the complementary question:

When multiple modalities are available, how is attention distributed across them during answer generation?

The distinction matters for questions of modality dominance and unimodal collapse. Strong performance under a multimodal configuration does not demonstrate that every available modality contributes meaningfully to the decision, since a model can reach a correct prediction by relying mainly on a dominant source while largely ignoring the rest. Comparing unimodal and multimodal conditions exposes this behaviour at the output level, but says nothing about whether the internal computation reflects the same imbalance. Tracing

how much attention is directed toward visual, acoustic, and textual representations during generation provides that missing evidence, showing whether the dominance observed behaviourally is also reflected in the model’s internal allocation. The comparison between the *Audio + Video* and *Video + Transcription* conditions is especially informative here. Both settings preserve the visual channel but carry the information associated with the audio track in different representational forms, native acoustic signal in one case, language in the other. Comparing the two allows us to ask whether linguistic reformulation makes audio-derived information more accessible to the language decoder, and whether this in turn changes the attention assigned to the video. The question connects directly to the observations made in *Lost in Transcription*. If transcription-based input attracts substantially more attention than raw audio, this would help explain why multimodal models behave differently once acoustic information is converted into text. If instead the decoder keeps concentrating attention on visual representations in both conditions, that would point to persistent visual dominance regardless of how the acoustic content is represented. The analysis is thus meant to connect three complementary levels: input availability, which measures whether adding or removing a modality changes task performance, semantic accessibility, which evaluates whether the information required by the question can already be recovered from the visual stream alone, and internal allocation, which investigates how the model distributes attention across the available sources while producing the answer.

Modality Availability → Behavioural Contribution

Visual Content → Semantic Accessibility (3.56)

Internal Attention → Modality Allocation

When examined together, the three levels provide a more comprehensive understanding of multimodal integration than considering each one individually. They help differentiate between a modality that is accessible but receives minimal attention, one that receives significant attention but offers little behavioral improvement, and one where both internal focus and performance shifts indicate a true multimodal contribution. Attention, in this context, is not viewed as a direct indicator of causal influence. A high level of attention doesn’t necessarily mean a modality determines the final outcome, just as a low level doesn’t imply its irrelevance. Instead, attention functions as an indicator of how resources

are internally distributed, while causal insights continue to emerge from manipulating modalities in controlled settings and observing the resulting performance advantages. The analysis of attention serves as a supplement to the behavioral ablation framework, rather than a substitute for it.

Setup	Input	Description
No Input	None	The model receives no multimodal information. The setup isolates how far the answer can be inferred from linguistic priors or dataset biases alone.
Audio Only	Audio	The model receives only the raw audio track, while the visual content is removed. It measures the information recoverable from the acoustic modality in isolation.
Video Only	Video	The model receives only the visual content of the video, without audio. It measures how much task-relevant information the visual modality alone provides.
Audio + Video	Audio + Video	The complete audiovisual signal is provided to the model. It represents the fully multimodal condition, assessing the contribution of joint audio-visual processing.
Transcription + Video	Transcription + Video	The visual content is combined with a textual transcription of the audio track. It tests whether linguistic information extracted from audio is exploited more effectively when given explicitly as text.

Table 3.4: Overview of the experimental setups used to evaluate the contribution of different input modalities.

4. Results and Discussions

4.1 Preliminary Analysis Results

Expected Behaviour under the Preliminary Condition

Before examining the experimental outcomes, it is useful to lay out qualitative expectations for model performance under the frame-only condition. This provides a methodological anchor for interpreting distinctions across question categories, rather than deriving explanations retrospectively from the results themselves. For spatial questions, a relatively high presence of visual evidence is expected, given the inherent visibility of entity locations, movements, object configurations, and spatial relationships within the visual stream. Complexity may increase, however, when this information must be reconstructed across multiple temporal intervals, as when tracking an entity through a sequence of movements or deriving a spatial configuration that results from several successive actions. Temporal questions are expected to display a more varied profile. Visually distinctive actions and event sequences can often be pinpointed through timestamp data and organized by the temporal structure produced during Event Analysis, but not all temporal information is necessarily recoverable from visual input alone. Some questions may involve events whose interpretation depends on linguistic context, or on relations between moments that remain ambiguous from visual information alone. The adequacy of the resulting representation therefore hinges on the quality of chronological reconstruction as much as on the inherent nature of the events involved. Causal questions are anticipated to be the most demanding under the visual-only condition. The visual stream can capture event succession, interactions, and observable state changes, but the underlying reasons for these events often depend on spoken dialogue or remain simply indistinguishable from static imagery. Motivations, intentions, and explanatory relationships are less likely to be directly supported by visual evidence alone. This expectation also underlines a methodological distinction maintained throughout the analysis: temporal succession is not treated as evidence of cau-

sation. These expectations are formed independently of the empirical results and serve only as a qualitative reference. No claim is made about the magnitude of category-level differences or about which category will ultimately score highest. Proficiency results are then compared with downstream Visual Statement Verification (VSV) performance, in order to separate evidence availability from task execution. Visual proficiency asks whether the independently constructed representation contains enough information to distinguish a caption from its foil. VSV performance asks whether the model produces the correct response once the question and the condition-specific inputs are directly available. Together, the two measures offer a diagnostic view that plain task accuracy cannot provide on its own. A Pass proficiency result paired with a correct VSV response indicates that the relevant visual evidence was both represented and successfully used downstream. A Pass result paired with an incorrect response instead points to a failure occurring after the evidence was already available, in selection, integration, or later reasoning. When both proficiency and VSV performance are negative, the most plausible explanation is a perceptual or representational shortfall, though this comparison alone cannot pinpoint the exact locus of failure. The most diagnostically informative configuration is the opposite one: proficiency Not Pass paired with a correct downstream answer. In this context, the model opts for the correct alternative even if the visual depiction fails to meet the proficiency standard. One plausible rationale is that the prediction harnesses partial visual indications, parametric priors, heuristic strategies, or random guessing at chance level. These elements are scrutinized more thoroughly in Section 4.2. However, this phenomenon might also suggest effective multimodal integration, where inadequacies in retrieving information from the visual stream are offset by evidence accessible through another input modality. Such interpretations should be approached prudently. The distinction between PASS and NOT PASS arises from a score given by an automatic evaluator based on a defined heuristic criterion and should not be viewed as a flawless evaluation of information accessibility. Instead, it serves as a diagnostic baseline for contrasting the initial representation with the model’s ensuing behavioral performance.

Entity, Event, and Inference Recovery

As discussed in Section 3.2, the preliminary analysis asks whether the visual content of a video alone allows a model to recover the information required to answer a question correctly. It also serves an independent diagnostic role by verifying that the newly constructed questions do not inherit the same limitation as the original MAIA benchmark [24], namely a dependence on visual information strong enough to reduce the possibility of assessing genuine cross-modal integration.

The analysis targets an earlier stage of the reasoning process than VSV itself, namely the semantic construction of what has been observed. This makes it possible to determine whether a model can independently extract from the video the entities, events, and inferential relations required by the question. The resulting representation separates information grounded in visual evidence from answers that may instead be supported by linguistic priors.

Table 4.1 reports the aggregate scores for the four evaluated models.

Model	Entity	Event	Inference	Comprehension
Gemma-4B	27.87	23.13	16.48	20.46
Gemma-12B	33.77	30.07	21.48	25.90
Qwen-3B	2.77	2.07	0.95	2.18
Qwen-30B	34.47	31.53	23.95	27.75

Table 4.1: Average preliminary-analysis scores across entity recovery, event recovery, inference recovery, and overall comprehension.

Entity-level information remains the most accessible representational component for all four models. At this level, the model primarily needs to associate a perceptually available element with its semantic representation, establishing referential grounding between an observed entity and the lexical concept used to describe it. The two larger models obtain the strongest results, with entity-recovery scores of 34.47 for Qwen-30B and 33.77 for Gemma-12B. The smaller Gemma-4B model follows the same pattern at a lower absolute level, reaching 27.87. Qwen-3B represents a clear outlier. Its entity score reaches only 2.77, followed by 2.07 for events and 0.95 for inference, resulting in an overall compre-

hension score of 2.18. The model therefore remains close to the lower bound of the evaluation despite preserving the same entity–event–inference ordering observed in the other architectures. This result indicates that only a very limited portion of the question-relevant visual information is recovered by its preliminary representation. The decline from entity recovery to inference recovery is visible in every model and reflects increasing representational demands. Event recovery requires individually detected entities to be integrated into dynamic structures involving actions, participants, and state transitions. This introduces a predicate–argument structure that extends beyond simple referential grounding. Inference recovery requires an additional transformation, since the model must derive relations or interpretations that do not necessarily correspond to a directly observable perceptual pattern. These may involve temporal ordering, causal interpretation, state dependencies, or relations connecting multiple observed events. Qwen-30B moves from 34.47 for entities to 31.53 for events and 23.95 for inferences. Gemma-12B follows a similar pattern, decreasing from 33.77 to 30.07 and then to 21.48. Gemma-4B decreases from 27.87 to 23.13 and finally to 16.48. The same hierarchy is present in Qwen-3B, although at substantially lower values, moving from 2.77 to 2.07 and finally to 0.95. The analysis highlights two complementary effects pertaining to representational complexity. Firstly, an increase in this complexity systematically diminishes the volume of retrievable information. Accurate entity identification does not necessarily imply that the associated event is represented, nor does the successful recovery of an event ensure the accurate reconstruction of the relations required by the query. Secondly, larger model variants within both families retrieve significantly more information compared to their smaller versions. This disparity is particularly evident in the Qwen model family. However, it is important to interpret the notably low Qwen-3B values with caution, as model size alone does not entirely determine the quality and completeness of the intermediate structured representation. Qwen2.5-Omni-3B demonstrated a heightened sensitivity to prompt formulation and input text handling during experimental analyses. Given that the experimental protocol mandated the consistent use of the same prompt across all four models, adaptations were made as much as possible to accommodate the formatting constraints intrinsic to Qwen2.5-Omni-3B without modifying the shared prompt itself. Despite these measures, a considerable number of

generation errors related to prompts were still encountered. Therefore, the anomalously low scores observed in this model may, in part, be attributed to instabilities in prompt processing and generation, rather than solely indicating a true deficiency in its capacity to recover visual information.

Category-Level Differences

The category breakdown reported only partly matches the qualitative expectations introduced above. Spatial information is particularly accessible to Qwen-30B, for which it produces the highest scores at every representational level, including a comprehension score of 38.85. In contrast, both Gemma-4B models achieve their highest overall scores on causal questions, with comprehension values of 25.79 and 31.75 respectively. Temporal questions constitute the most consistent difficulty across architectures. They produce the lowest comprehension score for every model, with values of 10.60 for Gemma, 15.05 for Gemma-12B, 1.30 for Qwen-3B, and 13.45 for Qwen-30B. The weakness is particularly evident at the inference level, where the corresponding scores fall to 7.35, 10.50, 0.50, and 9.30. This regularity indicates that extracting individual visual elements is not sufficient for reconstructing the temporal structure required by the questions. Causal questions display a different profile. Entity and event recovery remain comparatively high for the three stronger configurations, while inference recovery decreases substantially. This distinction is important because a causal question may involve visually accessible entities and events even when the causal relation connecting them is not directly observable. High entity or event recovery therefore does not imply successful causal reasoning. The preliminary analysis separates these stages and shows that much of the difficulty emerges precisely when the model must move from perceptual description to relational interpretation. Overall, the category-level results reinforce the distinction between perceptual availability and structured reasoning. Spatial configurations can often be recovered directly from visual evidence, especially by Qwen-30B. Temporal relations require the integration of information distributed across multiple moments and remain systematically difficult. Causal questions often provide sufficiently salient entities and events, but the inferential link between them remains substantially less recoverable. Qwen-3B reproduces the same general hierarchy

at markedly lower absolute values, indicating that the structure of the task remains visible even when the amount of successfully recovered information is minimal.

Preliminar Spatial Information

Spatial questions are the most successful category for Qwen-30B, reaching 41.50 on entity level, 40.50 on event level, and 37.05 on inference level, for a total comprehension score of 38.85. The relatively small drop from one level to the next is notable in itself: once the required entities have been retrieved, the model retains enough relational information to reconstruct the spatial configuration the question asks for, suggesting that spatial relations are often grounded directly in observable image content. The same trend appears for Gemma-12B, at lower absolute values: 35.90 for entities, 33.20 for events, and 28.25 for inferences. Gemma-4B reaches 29.10, 26.10, and 22.60 respectively. This category appears to benefit from the perceptual nature of its target relations: locations, relative positions, trajectories, and object configurations can often be detected within a limited temporal window, keeping the semantic relation the question targets close to perceptual structure. Spatial questions require relating lexicalized spatial concepts, location, direction, containment, proximity, movement, to geometrical or relational structures in the scene, which in turn requires correctly retrieving the referents and the relation between them. The stronger performance of Qwen-30B in this category indicates that its visual representations retain more of this relational structure. Spatial grounding is nonetheless not always successful: Qwen-30B’s overall comprehension score stays under 40, and correctly detecting an object does not guarantee that the required relation is also retrieved, particularly when it depends on tracking the entity over time.

Preliminar Temporal Information

Of all categories, temporal questions show the most degraded performance for every model tested: 10.60 for Gemma, 15.05 for Gemma-12B, and 13.45 for Qwen-30B. This decline is not confined to inference. Entity and event representation are already low for Gemma, at 15.50 and 12.50, while Gemma-12B reaches 20.10 and 19.40, and Qwen-30B reaches

18.30 and 17.60. Inference scores are lower still: 7.35, 10.50, and 9.30 respectively. Temporal difficulty, in other words, is not confined to the final reasoning step. The initial representations already fail to recover one or more of the entities or events needed to reconstruct the temporal relation itself. A temporal question typically depends on at least two event anchors that must be identified individually and then arranged relative to each other; missing one of them prevents the relation from being restored even when the other has been correctly identified. Qualitative inspection confirms this pattern: the model often recognizes the entities involved but fails to anchor one of the two events precisely enough to support relations such as *before*, *after*, *first*, or *second*. This points to a demand beyond simple lexical grounding: events need to be represented in their structural semantic form and placed within a shared temporal context, which is exactly what the low inference scores reflect. This is consistent with the original MAIA study, where temporal phenomena were especially difficult to resolve under visual-only evaluation [24], and with benchmarks such as VILMA, which treat temporal grounding as a capability distinct from static visual recognition.

Preliminar Causal Information

The causal category produces a more varied picture than expected at first glance. Causal questions do not receive the lowest scores; Gemma-4B and Gemma-12B in fact achieve their highest scores in this category, at 25.79 and 31.75 respectively, and Qwen-30B reaches 30.95, only slightly below its spatial score. This is not evidence that causality is visually easy, however. The score distribution tells a different story, with both models scoring far higher on entities and events than on inference. Gemma-12B reaches an entity score of 45.30 and an event score of 37.60, while its inference score drops to 25.70. Qwen-30B shows the same pattern, with scores of 43.60, 36.50, and 25.50 across the three levels. Gemma-4B shows the sharpest proportional drop, from an entity score of 39.00 to 30.80 for events and 19.50 for inference. The visual stream, in short, usually contains enough information to identify the participants and actions involved in a causal situation, but representing *why* something happens is considerably harder than representing *what* happens. This gap matters for how causal reasoning should be understood, since it depends on a re-

lation between propositions or events rather than a property of a single event on its own. A model can correctly lexicalize an action and identify its participants without representing the explanatory link between that action and an antecedent state, intention, or motivation, especially when that link was conveyed only through speech and never shown visually. The visual representation may therefore contain the observable outcome while lacking the information needed for causal interpretation. This experiment accordingly distinguishes between *causal event availability* and *causal relation recovery*.

Qualitative Distribution of Comprehension Levels

The continuous proficiency scores are complemented by a qualitative classification into four levels, *insufficient*, *partial*, *sufficient*, and *complete*. Table 4.2 reports the global distribution.

Model	Insufficient	Partial	Sufficient	Complete
Gemma-4B	70.0%	21.3%	4.7%	4.0%
Gemma-12B	60.7%	25.7%	7.7%	6.0%
Qwen-3B	99.0%	1.0%	0.0%	0.0%
Qwen-30B	61.7%	21.0%	8.7%	8.7%

Table 4.2: Distribution of qualitative comprehension levels in the preliminary analysis.

The analysis indicates that most models exhibit insufficient evidence for question-specific tasks. Gemma-4B shows a 70.0% success rate, which decreases to 60.7% for Gemma-12B and 61.7% for Qwen-30B. Complete representations are uncommon, with Qwen-30B achieving 8.7%, Gemma-12B at 6.0%, and Gemma-4B at 4.0%. In contrast, Qwen-3B presents a unique pattern, with 99.0% of its representations deemed insufficient and only 1.0% classified as partial, lacking any sufficient or complete cases. While these statistics offer a snapshot of model outputs, they should be interpreted with caution. Earlier evaluations showed that Qwen2.5-Omni-3B was sensitive to prompt and input formatting, necessitating uniform prompt use across four different models. As a result, there was a need to strictly follow the Qwen2.5-Omni-3B’s formatting rules while keeping prompts consistent. Despite these measures, issues related to prompt formatting persisted, affecting

the evaluators’ available information and implying that the low distribution for Qwen-3B wasn’t solely due to visual evidence deficiencies. Qwen-3B remains widely available but is excluded from detailed category analyses to prevent misinterpretation of potential generation errors as differences in spatial, temporal, or causal capabilities. Consequently, the focus is on the three other models where initial structured representations were stable enough for analysis. Data suggests that greater representational capacity is linked to a higher frequency of initial representation richness, though insufficiency is still predominant even among the strongest models. This underscores that representations are created without direct regard for future questions, highlighting the models’ ability to retrieve task-relevant evidence naturally rather than through targeted questioning. The category breakdown, reported in Table 4.3, further distinguishes the behaviour of spatial, temporal, and causal information. Temporal questions carry the highest concentration of insufficient represen-

Model	Dimension	Insufficient	Partial	Sufficient	Complete
Gemma-4B	Causal	58%	37%	5%	0%
Gemma-4B	Spatial	66%	15%	7%	12%
Gemma-4B	Temporal	86%	12%	2%	0%
Gemma-12B	Causal	48%	37%	14%	1%
Gemma-12B	Spatial	55%	22%	7%	16%
Gemma-12B	Temporal	79%	18%	2%	1%
Qwen-30B	Causal	51%	34%	14%	1%
Qwen-30B	Spatial	52%	12%	12%	24%
Qwen-30B	Temporal	82%	17%	0%	1%

Table 4.3: Distribution of qualitative comprehension levels by model and principal question dimension. Qwen-3B is excluded from the category-level comparison because the high incidence of prompt-related generation faults prevents a sufficiently reliable fine-grained interpretation across dimensions.

tations across the three models included in the category-level comparison. The proportion reaches 86% for Gemma, 79% for Gemma-12B, and 82% for Qwen-30B. Complete

temporal representations are almost nonexistent, with none observed for Gemma-4B and only 1% for Gemma-12B and Qwen-30B. This pattern indicates that the reconstruction of temporally distributed information remains particularly difficult when the representation is generated without knowledge of the subsequent question. Spatial questions show a markedly different distribution. Qwen-30B reaches complete comprehension in 24% of spatial cases and sufficient comprehension in a further 12%. Gemma-12B reaches 16% complete and 7% sufficient, while Gemma-4B reaches 12% and 7%, respectively. Spatial information therefore appears to be more readily preserved in the preliminary representation. In many cases, the required evidence can be grounded in directly observable entities, locations, surfaces, and local spatial relations without requiring the reconstruction of an extended sequence of events. The causal category occupies an intermediate position. Partial representations account for 37% of causal cases for both Gemma-4B and Gemma-12B and 34% for Qwen-30B. Gemma-12B and Qwen-30B also reach 14% sufficient representations, compared with 5% for Gemma. Complete causal representations remain rare. This distribution is consistent with the score-level pattern discussed above: models frequently recover some of the relevant participants, objects, and observable actions without reconstructing the complete causal dependency required by the question. The distinction between the global and category-level analyses is particularly important for Qwen-3B. Its global values document the outputs actually obtained during the experiment and are therefore retained in the overall comparison. However, the prevalence of prompt-related faults makes it difficult to determine whether differences between individual semantic dimensions would reflect genuine variation in visual recovery or instability in the generation interface. Excluding the model from the dimension-level table therefore avoids assigning a semantic interpretation to differences that the present experimental setup cannot reliably disentangle.

4.2 Experimental Results and Discussion

The experimental conditions offer a controlled view of how visual, auditory, and transcription-derived information contribute to Visual Statement Verification. Rather than reading mul-

timodal performance purely in terms of absolute accuracy, the analysis below treats each condition as evidence of the informational resources available to a model and of its ability to exploit them consistently. Throughout, it draws on the preliminary analysis in Section 4.1 to separate what a model can perceive from what it can do once perception is paired with a question. The comparison is particularly relevant against the original MAIA results [24] and against the experimental setting introduced in *Lost in Transcription* [4]. MAIA was designed as a visually grounded benchmark, in which the distinction between a true statement and its minimally contrastive foil is expected to depend primarily on information extracted from the video. The original evaluation showed that increasing visual evidence generally improves performance, though gains vary across reasoning categories and temporal phenomena remain particularly difficult. *Lost in Transcription* extended this setting to audio and transcription, finding that raw audio contributed substantially when visual evidence was absent but became almost neutral once dense visual information was available, and that adding textualized audio in the transcription-based setting even produced a small but systematic degradation. The results reported here present a very different scenario across model families and scales.

Overall Performance Across Experimental Conditions

Table 4.4 reports item-level accuracy across the five conditions involving perceptual or transcription-derived evidence. All models benefit from the availability of multimodal information, although the relative contribution of each channel varies substantially. Whether

Model	Audio Only	Transc. Only	Video Only	Audio + Video	Video + Transc.
Gemma-4B	74.58%	70.25%	69.83%	77.17%	75.08%
Gemma-12B	75.92%	69.25%	73.08%	80.50%	84.75%
Qwen-3B	73.50%	66.50%	72.50%	78.92%	75.50%
Qwen-30B	79.17%	77.00%	73.25%	83.17%	84.08%

Table 4.4: Item-level accuracy across unimodal and multimodal experimental conditions. Bold values indicate the best-performing multimodal configuration for each model.

these differences reflect systematic behaviour rather than descriptive fluctuation was tested

directly. Each model was evaluated on 1,200 items distributed across 300 pools, and since predictions under different conditions refer to the same underlying items, all comparisons below treat the six conditions as paired observations. An omnibus Cochran’s Q test¹, applied separately to each model, returns a statistically significant effect of experimental condition in every case, both at item level and at the stricter pool level defined by MAIA’s Aggregate Accuracy criterion. The null hypothesis of equal performance across conditions

Model	Item-level Q		p	Pool-level Q		p
Gemma-4B	194.45	< .001		113.63	< .001	
Gemma-12B	361.29	< .001		217.78	< .001	
Qwen-3B	392.89	< .001		193.52	< .001	
Qwen-30B	349.32	< .001		195.28	< .001	

Table 4.5: Cochran’s Q tests comparing performance across all six experimental conditions, at item level and at pool level.

is therefore rejected for every model, confirming that modality availability systematically affects predictions. This licenses the pairwise comparisons developed throughout the rest of the section, each tested with the exact McNemar test² and Holm-corrected³ for multiple comparisons within each model. A first descriptive observation concerns the relatively limited variability of the video-only condition, ranging from 69.83% for Gemma-4B to 73.25% for Qwen-30B, a spread of only 3.42 points. The two Qwen-3B configurations differ by even less, 72.50% against 73.25%. Such a compressed range contrasts sharply

¹Cochran’s Q test is a statistical method used to assess whether there are differences in a binary dependent variable across three or more related groups.

²The McNemar Test is a non-parametric statistical test applied to 2 x 2 contingency tables. It is commonly used to assess whether there are differences in dichotomous data—such as presence/absence or positive/negative—before and after a certain change, event, or treatment. By doing this, it evaluates the effectiveness of that treatment when frequency data is available. The purpose of this test is to determine if the marginal frequencies of the rows and columns are equal.

³In statistics, the Holm–Bonferroni method, also known as the Holm method or the Bonferroni–Holm method, addresses the issue of multiple comparisons. Its purpose is to manage the family-wise error rate (FWER) and it provides a straightforward test that is generally more powerful than the Bonferroni correction. The method is named after Sture Holm, who formalized it, and Carlo Emilio Bonferroni.

with the larger differences observed once linguistic or auditory information becomes available, and is compatible with the hypothesis, developed further in Section 4.2, that video-only performance is partly constrained by the perceptual representation itself rather than by the capacity of the language backbone. The original MAIA analysis points in the same direction: larger models generally benefited from a 32-frame configuration but continued to show weaknesses in visually demanding categories, particularly temporal duration and other phenomena involving distributed evidence [24]. The present results extend this observation: model capacity has a far clearer effect once an additional linguistic or auditory source becomes available, a pattern examined in detail below.

No-Input Condition and Response Bias

The no-input condition evaluates model behaviour when the question and the two candidate statements are provided without video, raw audio, or transcription. This setting calls for careful interpretation. Achieving a performance level significantly above chance indicates that the decision can be made, at least partially, without relying on perceptual evidence. However, it does not reveal the underlying mechanism responsible for this performance. **Above-chance accuracy is not, by itself, evidence of linguistic plausibility.** Prior work on multiple-choice evaluation supports this caution. Large language models exhibit systematic selection biases toward particular answer identifiers or positions, and their performance can shift substantially when the same candidate answers are reordered [28, 15]. For this reason, no-input accuracy needs to be read alongside the model’s response distribution, not on its own. The present dataset is globally balanced with respect to answer position: 600 items have the correct answer in the first position and 600 in the second. Table 4.6 reports no-input accuracy together with the proportion of valid predictions assigned to the first answer, the recall obtained when each answer position is correct, and the standard deviation of the two recalls ($RStd$). Following [28], $RStd$ is used as a diagnostic of selection bias, with larger values indicating a stronger asymmetry between answer positions. The four models exhibit markedly different behaviours in the no-input condition. Qwen-3B provides the clearest example of a positional shortcut. Its item accuracy remains close to chance at 50.83%, while 99.00% of its valid predictions select the first candidate.

Model	Accuracy	Ans. 1	R_0	R_1	$RStd$
Gemma	57.42%	64.01%	71.50%	43.33%	14.08
Gemma-12B	59.25%	77.35%	86.33%	32.17%	27.08
Qwen	50.83%	99.00%	99.83%	1.83%	49.00
Qwen-30B	61.00%	43.50%	54.50%	67.50%	6.50

Table 4.6: No-input performance and answer-selection diagnostics. Δ_{chance} is the difference between item accuracy and the 50% chance level. Ans. 1 reports the proportion of valid predictions selecting the first candidate answer. R_0 and R_1 denote recall conditional on the correct answer occupying the first or second position, respectively, and $RStd$ is their standard deviation.

Recall reaches 99.83% when the first candidate is correct and falls to 1.83% when the second candidate is correct, resulting in an $RStd$ of 49.00. Its near-chance accuracy therefore does not reflect an unbiased discrimination between the two statements. Rather, it is the consequence of an almost deterministic preference for the first answer position on a position-balanced dataset. Gemma-12B also displays a substantial positional asymmetry. The model selects the first candidate in 77.35% of valid predictions, with recalls of 86.33% and 32.17% for the first and second answer positions respectively ($RStd = 27.08$). Gemma-4B follows the same tendency, although less strongly, with an $RStd$ of 14.08. In both cases, positional preference contributes to the observed no-input performance and prevents aggregate accuracy from being interpreted directly as evidence that the model is discriminating between the semantic content of the two candidates. Qwen-30B presents a different pattern. Its predictions are substantially more balanced across answer positions, with an aggregate $RStd$ of only 6.50. A fixed positional shortcut is therefore unlikely to explain its above-chance performance. However, repeated evaluations carried out with exactly the same experimental configuration and candidate ordering produced noticeably different accuracies, approximately ranging from 55% to 62%. The observed performance should consequently not be regarded as a stable deterministic property of the model. A plausible interpretation is that Qwen-30B exhibits a **model-specific sensitivity to weak structural cues contained in the contrastive candidate pairs**. These cues may corre-

spond to subtle regularities introduced by the construction of true statements and their corresponding foils. This interpretation is distinct from a general linguistic plausibility effect. If the correct candidate could be systematically identified through the same overt linguistic regularities, a comparable advantage would be expected across architectures. Instead, the effect appears strongly model-dependent. The variation observed across repeated runs further suggests that these weak cues do not determine the response consistently, but may bias the decision when perceptual evidence is unavailable. Breaking the results down by reasoning category reinforces the heterogeneous nature of these no-input strategies.

Model	Causal		Spatial		Temporal	
	Acc.	RStd	Acc.	RStd	Acc.	RStd
Gemma	68.25%	0.25	55.50%	12.00	48.50%	30.00
Gemma-12B	68.00%	17.00	57.75%	30.75	52.00%	33.50
Qwen	51.25%	48.25	50.50%	49.50	50.75%	49.25
Qwen-30B	64.75%	18.25	62.75%	0.25	55.50%	1.50

Table 4.7: No-input item accuracy and answer-position recall imbalance across reasoning categories.

Accuracy and answer-position imbalance do not covary consistently across categories. Gemma, for example, reaches its highest no-input performance on causal questions while remaining almost perfectly balanced between answer positions. Qwen-30B shows a different configuration. Its spatial and temporal predictions are nearly position-balanced, whereas causal questions produce a larger positional asymmetry. Qwen, by contrast, preserves essentially the same strong preference for the first answer across all three reasoning categories. These results separate two phenomena that aggregate accuracy alone would conflate. **Answer-position bias explains a substantial part of the no-input behaviour of some models, but it cannot account for all above-chance performance.** Conversely, above-chance accuracy in the absence of a strong position bias does not demonstrate that the model has recovered information relevant to the audiovisual event. In the case of Qwen-30B, the results are instead compatible with a model-specific sensitivity to weaker structural regularities in the contrastive pairs, although dedicated controls would be re-

quired to identify their precise nature. The no-input condition is therefore best interpreted as a diagnostic of *decision behaviour in the absence of perceptual grounding*. Depending on the model, decisions may be dominated by a positional shortcut, influenced by weak structural properties of the candidate pair, or exhibit substantial instability across repeated inferences. No-input accuracy should consequently not be treated as evidence of a single common linguistic bias, but as the observable outcome of different model-specific strategies operating when the perceptual evidence required by the task is removed.

Audio-Only Performance

One of the most notable results concerns the audio-only condition. As shown in Table 4.8, raw audio outperforms video-only input for every model. This result matters

Model	Audio Only	Video Only	Difference
Gemma-4B	74.58%	69.83%	+4.75
Gemma-12B	75.92%	73.08%	+2.83
Qwen-3B	73.50%	72.50%	+1.00
Qwen-30B	79.17%	73.25%	+5.92

Table 4.8: Comparison between audio-only and video-only item accuracy. Differences are expressed in percentage points.

because MAIA was originally developed as a visually grounded benchmark: its audio track played no role in constructing the questions and statements. Strong audio-only performance does not make the benchmark intrinsically audio-based: it shows instead that the acoustic stream carries information strongly correlated with the events and situations the questions query. Raw audio can supply several complementary sources of evidence at once. Spoken language can explicitly identify objects, actions, locations, intentions, or causal relations. Environmental sounds can signal ongoing events. Acoustic dynamics can offer temporal cues about occurrence or transition. The linguistic component of speech is likely central to this advantage. Once speech is encoded, its semantic representation interfaces naturally with the language backbone that compares the two candidate statements, narrowing part of the representational gap between visual patterns and linguistic proposi-

tions. The result is best read as evidence of high *informational accessibility* in the acoustic channel, not as proof that audio is inherently richer than vision: a visual scene may contain more information overall while still being harder for the architecture to convert into a representation aligned with the VSV decision. The pattern differs markedly from *Lost in Transcription*, where Qwen2.5-Omni-7B obtained 82.70% with 32 visual frames and 82.68% once raw audio was added, a negligible -0.02 point difference. There, raw audio became beneficial mainly once visual evidence was removed entirely, raising black-frame accuracy from 67.40% to 73.19% [4]. Audio plays a substantially stronger role across the architectures evaluated here. Model-family effects, differences in the reasoning categories considered, and stronger audio exploitation in more recent omnimodal architectures could all contribute. The present results support the behavioural difference without isolating its architectural cause.

Transcription-Only Performance

Comparing raw audio against its transcription gives a more direct view of the distinction between acoustic and linguistic information. Raw audio outperforms transcription

Model	Audio Only	Transcription Only	Difference
Gemma-4B	74.58%	70.25%	+4.33
Gemma-12B	75.92%	69.25%	+6.67
Qwen-3B	73.50%	66.50%	+7.00
Qwen-30B	79.17%	77.00%	+2.17

Table 4.9: Comparison between raw audio and transcription-only accuracy. Positive differences indicate an advantage for raw audio.

for every model, from 7.00 points for Qwen-3B down to 2.17 for Qwen-30B. The contribution of audio therefore cannot be reduced to a simple speech-to-text transformation.

Transcription preserves lexical and propositional content, but it strips out or compresses non-verbal sound, prosody, overlapping speakers, temporal acoustic events, and background activity. The transcription pipeline can itself introduce recognition errors or drop uncertain material, so the textual representation is a lossy transformation of

the original signal. The gap between the small and large Qwen-3B configurations adds a further, linguistically informative pattern. Qwen-3B reaches 66.50% with transcription alone, Qwen-30B reaches 77.00%, a gap considerably larger than the 0.75-point difference the same two models show in video-only input. Increased model capacity is especially beneficial once information is already represented in linguistic form: a larger backbone exploits lexical semantics, discourse relations, event descriptions, and contextual associations more effectively than a smaller one, so scaling gains look more pronounced in language-mediated reasoning than in isolated visual grounding. This fits a broader distinction between perceptual encoding and linguistic inference. Once evidence has been converted to text, what remains is largely semantic selection and integration. Visual evidence, by contrast, must first be turned into a linguistic-semantic representation before the same kind of reasoning can even begin. The pattern also contrasts with *Lost in Transcription*, where Qwen2.5-VL-7B reached 68.72% in the linguistic-only baseline but only 63.98% with transcription alone: transcription failed to outperform the model’s own linguistic prior [4]. The models evaluated here instead gain substantially from transcription relative to their no-input baselines, particularly Qwen-30B. The ability to convert auxiliary textual evidence into useful task information appears to vary with scale rather than being a fixed property of an architecture.

Video-Only Performance and the Visual Bottleneck

Video-only accuracy ranges from 69.83% to 73.25%, as already shown in Table 4.4. All four models extract substantial task-relevant evidence from the sampled frames, yet the narrow range across architectures suggests a shared bottleneck in visual processing rather than a strongly scale-dependent one. Increasing model size from Qwen-3B to Qwen-30B produces only a 0.75-point gain in video-only accuracy, against a 10.50-point gap between the same two models in transcription-only accuracy. Linguistic scale therefore appears to improve more strongly the operations performed once relevant semantic information is already accessible, including lexical comparison, semantic compatibility assessment, integration of explicit contextual evidence, and reasoning over propositions, without producing an equivalent improvement in temporally distributed visual extraction. The distinction

becomes considerably sharper once video-only accuracy is compared with the preliminary analysis of Section 4.1. Overall visual comprehension there reaches only 20.46 for Gemma, 25.90 for Gemma-12B, and 27.75 for Qwen-30B, well below the 69.83%, 73.08%, and 73.25% obtained by the same models in video-only VSV. This discrepancy reflects different informational regimes rather than a contradiction between the two evaluations. The preliminary analysis constructs a representation without access to the downstream question and measures whether the entities, events, and relations later required by the task are spontaneously recovered from the visual input. In the video-only VSV condition, by contrast, the question and the two minimally contrastive alternatives explicitly identify the semantic distinction to be resolved. Visual processing can therefore be conditioned on a specific linguistic target, directing evidence selection toward the relevant entities, actions, temporal intervals, or relations. The restricted contrastive decision space also means that a complete semantic reconstruction of the video is not always necessary for a correct prediction. The distinction is therefore one between *question-independent representation* and *question-conditioned retrieval*. A direct statistical comparison supports this interpretation. At the video level, correlations between preliminary comprehension and downstream video-only performance are generally weak. Gemma-12B is the only model showing a significant association, with $\rho = 0.280$ for item accuracy ($p = .0048$), while the other models remain close to zero or non-significant. At the question level, however, stronger preliminary representations clearly increase the probability of downstream success. Item accuracy rises from 68.8% to 80.8% for Gemma, from 71.1% to 85.4% for Gemma-12B, and from 70.3% to 87.5% for Qwen-30B when moving from *Not Pass* to *Pass* preliminary representations. Qwen-30B also shows the strongest pool-level effect, increasing from 35.9% to 67.3% ($OR = 3.68$, $p < .001$). These results show that richer question-independent visual evidence facilitates downstream reasoning but is not a necessary condition for correctness. This is particularly evident for Qwen-3B, which achieves 72.5% video-only item accuracy despite having no preliminary representations classified as *Pass*. Its preliminary scores are affected by an extreme floor effect and by the prompt-related generation faults discussed above, so their relationship with downstream performance cannot be interpreted in the same way as for the other models. More

generally, the high accuracy retained under *Not Pass* conditions shows that evidence not sufficiently recovered during the preliminary stage may still become accessible once the question constrains the retrieval process. The original MAIA study offers a useful external reference point. With 32 frames, Qwen2.5-VL-7B obtained a pool-based VSV accuracy of 0.44, while Qwen2.5-VL-72B reached 0.54 [24]. This indicates a scaling advantage, but not full consistency: spatial, causal, and implicit categories benefited from richer visual evidence, while temporal reasoning remained comparatively weak. Direct numerical comparison with the present models must remain cautious because the current evaluation covers a different category subset and different architectures. The qualitative pattern is nevertheless compatible: stronger models can exploit visual evidence more effectively, but increased language-model capacity alone does not remove the difficulty of grounding temporally and relationally complex information. The preliminary analysis further localizes this limitation by showing that part of the bottleneck arises before downstream reasoning, at the stage at which relevant information is extracted and organized from the visual stream.

Audio–Visual Performance

Adding raw audio to video improves item-level accuracy for every model. The descriptive

Model	Video Only	Audio + Video	Δ_{audio}
Gemma-4B	69.83%	77.17%	+7.33
Gemma-12B	73.08%	80.50%	+7.42
Qwen-3B	72.50%	78.92%	+6.42
Qwen-30B	73.25%	83.17%	+9.92

Table 4.10: Conditional gain produced by adding raw audio to the video-only condition.

gains, from 6.42 points for Qwen-3B to 9.92 for Qwen-30B, are all statistically robust: every confidence interval excludes zero and every Holm-corrected test remains significant. Raw audio is therefore not simply ignored once visual evidence is present, a pattern that differs from the 32-frame results in *Lost in Transcription*, where Qwen2.5-Omni-7B showed virtually identical accuracy with and without raw audio [4]. Between 44.75% and

Model	Comparison	Base Acc.	New Acc.	Δ	95% CI	Holm p
Gemma-4B	Video $\rightarrow$ Audio+Video	69.83%	77.17%	+7.33	[4.75, 10.00]	< .001
Gemma-12B	Video $\rightarrow$ Audio+Video	73.08%	80.50%	+7.42	[4.67, 10.25]	< .001
Qwen-3B	Video $\rightarrow$ Audio+Video	72.50%	78.92%	+6.42	[3.33, 9.42]	< .001
Qwen-30B	Video $\rightarrow$ Audio+Video	73.25%	83.17%	+9.92	[7.25, 12.58]	< .001

Table 4.11: Exact McNemar comparisons between the video-only condition and the audio-visual condition. Accuracy differences are in percentage points. Confidence intervals were obtained through cluster bootstrap resampling at the pool level, and reported p -values are Holm-corrected.

Model	Repair Rate	Damage Rate
Gemma-4B	44.75%	8.83%
Gemma-12B	48.92%	7.87%
Qwen-3B	46.97%	8.97%
Qwen-30B	49.84%	4.66%

Table 4.12: Repair and damage rates when raw audio is added to video input. Repair Rate measures the proportion of video-only errors corrected by audio. Damage Rate measures the proportion of video-only correct predictions that become incorrect after adding audio.

49.84% of video-only errors are corrected once raw audio is introduced, suggesting that the acoustic stream frequently carries information capable of resolving insufficient or ambiguous visual evidence, while damage rates stay below 9% for every model. Qwen-30B shows the strongest asymmetry: a Repair Rate of 49.84% against a Damage Rate of only 4.66%. Its audio-visual improvement does not come from generic prediction instability; it comes from the additional modality converting incorrect visual decisions into correct multimodal ones. The underlying McNemar contingency counts make the direction of this effect explicit: for Qwen-30B, adding raw audio repairs 160 video-only errors while introducing only 41 new ones. This kind of repair is consistent with several linguistic mechanisms. The acoustic modality may supply a **semantic representation that resolves underspecified visual evidence**. Spoken descriptions can lexicalize events that are difficult to infer from sparse frames. **Environmental sounds can disambiguate between**

competing event interpretations. In each case, the language backbone can use this evidence to judge which candidate statement is better supported. The reciprocal question, whether vision still adds anything once audio is already available, receives a more mixed answer. Vision significantly improves on audio-only performance for Gemma-12B, Qwen,

Model	Audio Only	Audio+Video	Δ	Holm p
Gemma-4B	74.58%	77.17%	+2.58	.153
Gemma-12B	75.92%	80.50%	+4.58	.002
Qwen-3B	73.50%	78.92%	+5.42	< .001
Qwen-30B	79.17%	83.17%	+4.00	< .001

Table 4.13: Exact McNemar comparisons between audio-only and audio–video conditions.

and Qwen-30B, but not for Gemma, whose +2.58 point gain does not survive Holm correction ($p = .153$). For this model, raw audio already captures a large share of the evidence that determines the final prediction, and adding vision produces no statistically robust item-level advantage, a pattern compatible with a degree of modality dominance in which the acoustic channel alone accounts for most of what the task requires. The other three models instead show clear, significant benefits from the joint availability of both channels: vision continues to contribute even once audio is already on the table.

Video–Transcription Performance

The effect of transcription becomes markedly dependent on model scale once visual information is already available. Gemma-4B and Qwen-3B gain 5.25 and 3.00 points

Model	Video Only	Video + Transcription	$\Delta_{\text{transcription}}$
Gemma-4B	69.83%	75.08%	+5.25
Gemma-12B	73.08%	84.75%	+11.67
Qwen-3B	72.50%	75.50%	+3.00
Qwen-30B	73.25%	84.08%	+10.83

Table 4.14: Conditional gain produced by adding transcription to the video-only condition.

Model	Comparison	Base Acc.	New Acc.	Δ	95% CI	Holm p
Gemma-4B	Video $\rightarrow$ Video+Transcription	69.83%	75.08%	+5.25	[2.25, 8.25]	< .001
Gemma-12B	Video $\rightarrow$ Video+Transcription	73.08%	84.75%	+11.67	[8.25, 15.17]	< .001
Qwen-3B	Video $\rightarrow$ Video+Transcription	72.50%	75.50%	+3.00	[0.25, 5.75]	.023
Qwen-30B	Video $\rightarrow$ Video+Transcription	73.25%	84.08%	+10.83	[8.17, 13.58]	< .001

Table 4.15: Exact McNemar comparisons between the video-only condition and the video-plus-transcription condition.

respectively, while Gemma-12B and Qwen-30B gain 11.67 and 10.83. Every one of these differences remains significant after Holm correction, including Qwen’s comparatively small gain ($p = .023$), so even the weakest effect in this comparison cannot be dismissed as noise. The contribution of linguistic auxiliary information clearly increases in the larger configurations. A transcription is already expressed in the representational domain where the language backbone is strongest, but exploiting it well is not automatic: the model still has to identify the portions semantically related to the question, suppress irrelevant discourse, resolve references, and integrate textual evidence with the visual representation. The larger models appear considerably better at these operations, so increased capacity may add stronger contextual filtering and semantic alignment between heterogeneous sources of evidence, not just richer linguistic knowledge. Unlike the raw-audio com-

Model	Transcription Only	Video+Transcription	Δ	Holm p
Gemma-4B	70.25%	75.08%	+4.83	.011
Gemma-12B	69.25%	84.75%	+15.50	< .001
Qwen-3B	66.50%	75.50%	+9.00	< .001
Qwen-30B	77.00%	84.08%	+7.08	< .001

Table 4.16: Exact McNemar comparisons between transcription-only and video-plus-transcription conditions.

parison of Section 4.2, vision significantly improves on transcription-only performance in every model. The effect is largest precisely where the video-plus-transcription gain was largest: Gemma-12B, where 15.50 additional points correspond to 238 items moving from incorrect to correct against only 52 moving the other way, for a Repair Rate of 64.09% and

Model	Repair Rate	Damage Rate
Gemma-4B	43.65%	11.34%
Gemma-12B	64.09%	7.64%
Qwen-3B	41.21%	11.49%
Qwen-30B	50.78%	3.75%

Table 4.17: Repair and damage rates when transcription is added to video input.

a Damage Rate of just 7.64%. Qwen-30B shows a similarly stable, asymmetric pattern: a Repair Rate of 50.78% and the lowest Damage Rate observed in any transcription-based comparison, 3.75%. This asymmetry indicates that transcription and visual content carry genuinely complementary information; a model relying purely on textualized audio and ignoring the video would show no such systematic, one-sided correction pattern. The behaviour is the opposite of what *Lost in Transcription* reported, where Qwen2.5-VL-7B decreased from 80.67% with 32 frames alone to 78.78% with video and transcription combined. Item-level McNemar analysis found that 1.89 point drop to be systematic, with correct predictions converted into errors more often than incorrect ones were repaired [4]. The current models instead find transcription beneficial, particularly at larger scale. Transcription-based multimodal integration, in this evaluation, depends substantially on model architecture, scale, and the language backbone’s ability to regulate auxiliary textual evidence.

Raw Audio and Transcription as Complementary Representations

The purely descriptive ranking between the two multimodal conditions suggests a split by model size. For Gemma-4B and Qwen, raw audio gives the strongest multimodal configuration, 77.17% against 75.08% for Gemma-4B and 78.92% against 75.50% for Qwen. For the larger configurations the ranking flips, 84.75% against 80.50% for Gemma-12B and 84.08% against 83.17% for Qwen-30B. Once tested directly, this small-versus-large story does not hold up cleanly. Gemma’s numerical advantage for raw audio and Qwen-30B’s numerical advantage for transcription both fail to reach significance ($p = .153$ and $p = .351$), so for these two models the item-level data cannot statistically distinguish be-

Model	Audio+Video	Video+Transcription	Δ_{VT-AV}	Holm p
Gemma-4B	77.17%	75.08%	-2.08	.153
Gemma-12B	80.50%	84.75%	+4.25	.001
Qwen-3B	78.92%	75.50%	-3.42	.009
Qwen-30B	83.17%	84.08%	+0.92	.351

Table 4.18: Direct comparison between native audio–video input and video enriched with transcription. Positive differences favour transcription, negative differences favour raw audio.

tween the two multimodal representations. A significant, representation-dependent preference emerges only for the remaining two models, and it does not line up neatly with scale: Gemma-12B performs significantly better with transcription-derived linguistic information ($\Delta = +4.25$ points, $p = .001$), while Qwen-3B shows the opposite pattern, with native raw audio outperforming transcription by 3.42 points ($p = .009$). The data therefore do not support a general claim that smaller models systematically prefer raw audio and larger models systematically prefer transcription. The advantage of one representation over the other is model-dependent, and likely reflects the specific interaction between a model’s modality-specific encoders and the linguistic reasoning capacity of its backbone, more than any property of scale on its own. This also clarifies why transcription should not be treated as a simple stand-in for audio. Raw audio preserves acoustic and temporal information at the cost of requiring cross-modal representation alignment. Transcription discards part of that information in exchange for direct access to lexical and propositional structure. Which of the two wins out for a given model depends on the balance it strikes between perceptual richness and linguistic accessibility, and the present analysis shows that balance to be architecture-specific, not simply a matter of parameter count.

Aggregate Accuracy and Prediction Consistency

Item-level accuracy does not reveal whether a model behaves consistently across the several minimally contrastive realizations tied to the same underlying question. Aggregate Accuracy, which counts a pool as correct only when every one of its items is answered cor-

rectly, gives a stricter estimate of robustness. The gap between item-level and aggregate

Model	Audio Only	Transc. Only	Video Only	Audio + Video	Video + Transc.
Gemma-4B	41.67%	37.67%	34.00%	49.00%	42.67%
Gemma-12B	47.00%	36.00%	42.00%	55.33%	64.00%
Qwen-3B	40.33%	31.33%	38.33%	50.00%	44.33%
Qwen-30B	55.67%	52.00%	41.33%	60.67%	63.33%

Table 4.19: Aggregate Accuracy across experimental conditions. A pool is counted as correct only when the model correctly resolves all minimally contrastive items associated with it.

performance is substantial for every model, confirming that isolated correct predictions do not necessarily correspond to stable semantic discrimination. Multimodal conditions improve consistency as well as raw correctness: Gemma-12B rises from 42.00% Aggregate Accuracy in video-only to 64.00% with video and transcription, Qwen-30B from 41.33% to 63.33%. This connects directly to MAIA’s own evaluation philosophy, which treats visually grounded competence as unproven by isolated correct decisions if semantically equivalent reformulations produce inconsistent predictions [24]. The original 32-frame VSV results reached 0.44 pool accuracy for Qwen2.5-VL-7B and 0.54 for Qwen2.5-VL-72B [24], while the present multimodal conditions push above 0.60 for the larger models. Direct comparison remains limited, though, by the different model families and category subsets involved. Repeating the paired significance tests at pool level largely confirms these item-level results while also exposing an important exception in prediction consistency. Adding raw audio to video significantly improves Aggregate Accuracy for every model. Transcription produces a significant pool-level gain for Gemma, Gemma-12B, and Qwen-30B, but not for Qwen. Its item-level improvement from 72.50% to 75.50% (Section 4.2, $p = .023$) does not carry over to its Aggregate Accuracy gain from 38.33% to 44.33% ($p = .071$). For this model, transcription improves local, item-level decisions without correcting them systematically enough to raise the number of fully consistent pools to a statistically robust degree: the effect is detectable in isolated predictions but too weak to survive the stricter, semantic-consistency criterion MAIA is specifically designed

Model	Comparison	Base AA	New AA	Δ	Holm p
Gemma-4B	Video $\rightarrow$ Audio+Video	34.00%	49.00%	+15.00	< .001
Gemma-4B	Video $\rightarrow$ Video+Transcription	34.00%	42.67%	+8.67	.019
Gemma-12B	Video $\rightarrow$ Audio+Video	42.00%	55.33%	+13.33	< .001
Gemma-12B	Video $\rightarrow$ Video+Transcription	42.00%	64.00%	+22.00	< .001
Qwen-3B	Video $\rightarrow$ Audio+Video	38.33%	50.00%	+11.67	< .001
Qwen-3B	Video $\rightarrow$ Video+Transcription	38.33%	44.33%	+6.00	.071
Qwen-30B	Video $\rightarrow$ Audio+Video	41.33%	60.67%	+19.33	< .001
Qwen-30B	Video $\rightarrow$ Video+Transcription	41.33%	63.33%	+22.00	< .001

Table 4.20: Exact McNemar comparisons at pool level using Aggregate Accuracy.

to enforce. Gemma-12B and Qwen-30B, by contrast, both gain 22 percentage points in Aggregate Accuracy from transcription, remaining highly significant after correction. For these two models, the additional linguistic evidence improves not just how many decisions are correct, but how stable those decisions are across minimally different formulations of the same content.

Audio–Video Interaction Effect

Positive multimodal gains do not by themselves imply that audio and video combine independently. A logistic interaction analysis over the 2×2 factorial design of no-input, audio-only, video-only, and audio-visual conditions returns a negative interaction coefficient for every model. All four coefficients are negative and statistically significant: Qwen-3B shows the strongest interaction at $\beta_{A \times V} = -0.636$, Qwen-30B the weakest at -0.298 . A negative coefficient does not mean combining audio and video reduces absolute performance; audio-visual accuracy is higher than video-only accuracy for every model, as already established by the significant McNemar comparisons in Section 4.2. The coefficient instead indicates that the joint gain is smaller than an additive model of independent modality contributions would predict, pointing to partial informational redundancy or competition between the two channels. Some of what audio contributes may already be recoverable from the video, and the model may fail to preserve the full benefit

Model	$\beta_{A \times V}$	p -value	Odds Ratio
Gemma-4B	−0.399	< .001	0.671
Gemma-12B	−0.355	.001	0.701
Qwen-3B	−0.636	< .001	0.529
Qwen-30B	−0.298	.011	0.743

Table 4.21: Audio–video interaction coefficients estimated from the factorial comparison between no-input, audio-only, video-only, and audio–video conditions. Negative coefficients indicate a sub-additive interaction.

of each source once both compete within a shared representational or attentional space. Read together, these two results, a significant audio-visual improvement over video alone and a significant sub-additive interaction, describe two distinct properties of multimodal behaviour rather than contradicting one another. The pairwise tests establish that audio provides **useful additional evidence whenever video is available**. The interaction coefficient establishes that **this additional evidence is not combined with the visual channel as an independent, additive source**. Audio and video are complementary, in that both improve performance, but part of their informational content overlaps or is not fully preserved during integration: exactly the distinction between *multimodal benefit* and *multimodal integration* that this thesis has argued for throughout. The result is compatible with the modality-balancing problem already identified in *Lost in Transcription*, where Qwen2.5-Omni dynamically reduced attention to audio as visual information increased, without the attention still allocated to audio necessarily translating into corresponding gains in reasoning performance [4]. The behavioural results reported here extend this observation across different architectures: additional modalities are used, but their informational contribution stays partly redundant, and their combination falls short of a fully additive effect.

Model Scale and Linguistic Integration

The influence of model scale varies depending on the input modality. Its impact is least noticeable in video-only scenarios, where the difference between Qwen-3B and Qwen-30B

is minimal at 0.75 points. In contrast, the effect is more pronounced in transcription-based contexts, with Qwen-30B surpassing Qwen-3B by 10.50 points when using transcription alone and by 8.58 points when both video and transcription are combined. Similarly, the Gemma-4B family shows that Gemma-12B has a 3.25-point advantage over Gemma-4B in video-only inputs, which increases to 9.67 points with both video and transcription. Scaling enhances a model’s ability to use linguistically structured context effectively. A larger language backbone enables finer lexical distinctions, maintains discourse dependencies better, resolves referential relations more efficiently, and evaluates semantic compatibility between questions, transcriptions, and potential answers—all critical after acoustic information is translated into text. However, visual input presents a different challenge; the model must first derive an informative semantic representation from perceptual evidence before applying linguistic reasoning. Increasing language model capacity doesn’t ensure that relevant entities, events, or spatial and temporal relations are correctly extracted. This asymmetry is supported by preliminary analysis in Section 4, which found that scaling improved scores for entity, event, and inference recovery but did not change the underlying entity-event-inference hierarchy. It is further supported by the video-only bottleneck described in Section 5, which traced the limitation back to what the visual representation contains, independent of post-linguistic reasoning. The original MAIA results support this finding: Qwen2.5-VL-72B significantly outperformed Qwen2.5-VL-7B in a 32-frame pool-based evaluation (0.54 versus 0.44), yet it still faced challenges with temporally demanding categories despite its scaling advantage. This suggests that increased capacity does not consistently resolve perceptual grounding issues. Overall, these findings highlight that scaling is particularly effective when evidence is already in linguistic form, while improvements with isolated visual inputs remain limited. This asymmetry aligns with a multimodal performance framework based on two distinct capacities: **extracting task-relevant evidence from a given modality and reasoning over that evidence once it is in a semantically accessible form**. The preliminary analysis in Section 4.1 primarily addresses the first capacity, while the experimental results discussed here mainly relate to the second. The interaction between both capacities determines how well a model integrates audio, video, and language in Visual Statement Verification.

Attention Allocation and Modality Integration

The behavioral analyses above establish whether adding an auxiliary modality changes VSV performance. Attention analysis addresses a different question: *how the model distributes its attention during generation when the same visual evidence is combined either with raw audio or with an audio-derived transcription*. The two levels should not be conflated. Attention weights provide a diagnostic view of information access and allocation, but they are not causal attributions and **do not measure the effective contribution of a modality to the final decision**, their interpretation is therefore most informative when considered together with Conditional Modality Gain and the corresponding Repair and Damage rates. Figure 4.1 reports the mean attention fractions over the complete autoregressive attention distribution. The values distinguish attention assigned to video, raw audio, transcription, question tokens, the remaining textual context, and tokens generated earlier in the answer. Raw attention fractions are used rather than source-normalized shares in order to preserve a common denominator across models. In particular, the source-normalized estimates obtained for Qwen-30B are not directly comparable with those of the other systems, whereas the raw attention components form a coherent decomposition of the complete attention distribution. A zero value in the question column should likewise be interpreted as a property of the token-span segmentation used for that model, not as evidence that the model ignores the question. The first relevant result is that **replacing raw audio with its transcription does not induce a uniform redistribution of attention**: the same transformation produces qualitatively different patterns across models. For Gemma, transcription receives 9.79% of the total attention budget, compared with 5.99% for raw audio. This increase is accompanied by a modest reduction in visual attention, from 61.22% to 59.02%, and by a reduction in attention to previously generated tokens. The transcription therefore becomes more salient than the native acoustic representation within the same visual context, although this does not establish that transcription contributes more useful evidence to the decision; it shows instead that converting acoustic information into linguistic form changes how strongly that auxiliary source competes with the visual stream for attention. Gemma-12B exhibits the opposite relation. Raw au-

dio receives 7.34% of total attention, while transcription receives 4.60%. Visual allocation remains nearly stable across the two conditions, increasing only from 42.97% to 43.80%, and generated-history attention is effectively unchanged at approximately 13.7%. The substitution of raw audio with text therefore alters the salience of the auxiliary modality without producing a broader reorganization of the model’s autoregressive attention regime. Taken together with the Gemma-4B result, this already excludes a general interpretation in which transcription is systematically more accessible than native audio, or vice versa. The smaller Qwen-3B model provides an especially useful control for separating attentional allocation from behavioral contribution. Raw audio receives 5.19% of total attention in the A+V condition, while transcription receives 5.24% in V+T. Video attention also remains nearly unchanged, from 6.48% to 6.28%, as do the remaining textual context and generated history. The model therefore allocates almost the same amount of attention to the auxiliary source independently of whether that source is presented acoustically or textually, so any substantial difference between the corresponding Conditional Modality Gain, Repair Rate, or Damage Rate cannot be explained by a simple difference in the amount of attention devoted to the added modality. It would instead indicate a difference in how the attended representation is transformed into evidence useful for the VSV decision. Qwen-30B shows the most pronounced condition-dependent reorganization. In the A+V condition, 76.34% of total attention is assigned to video and 21.69% to raw audio. The two perceptual sources therefore account for approximately 98.0% of the measured attention distribution, while generated history receives only 1.60%. The model is thus clearly capable of allocating substantial attention to native acoustic input when audio and video are jointly available, although high audio attention should not be equated with effective audiovisual reasoning: its behavioral value must still be determined from the paired change in prediction relative to Video Only. The V+T condition produces a markedly different profile for the same model. Visual attention decreases to 61.53%, transcription receives only 0.62%, and attention to previously generated tokens rises from 1.60% to 36.51%. The substitution of native audio with transcription therefore coincides with a transition from an attention pattern dominated by external perceptual sources to one with substantially greater allocation to the model’s own generated history. This pattern is compatible

with stronger autoregressive self-conditioning in the transcription condition, although attention weights alone cannot establish the causal role of this shift. The important result is that raw audio and transcription do not behave as interchangeable carriers of the same semantic content: changing the representational form of the auxiliary signal can reorganize the complete attentional context in which the answer is generated. Across the four models, the relative ordering between raw audio and transcription is therefore heterogeneous. Gemma-4B allocates more attention to transcription than to audio, Gemma-12B allocates more to audio than to transcription, Qwen-3B assigns nearly identical attention to the two representations, and Qwen-30B strongly favors native audio. The common thread across models is not a universal preference for one modality, but the dependence of multimodal allocation on the representation through which auxiliary information is supplied. This provides an internal counterpart to the behavioral comparison between Audio+Video and Video+Transcription: the conditional experiments establish whether the additional source changes decisions, while the attention analysis shows that similar behavioral outcomes can arise from different internal allocation patterns and that similar attention levels can support different behavioral effects. This distinction matters when interpreting modality gain. A large attention share cannot by itself be taken as evidence that the corresponding modality provides discriminative information. Conversely, a small share does not imply that the modality is irrelevant, since a limited amount of highly diagnostic evidence may still alter the decision. Effective multimodal integration is therefore better characterized by the relation between three levels of evidence: attentional allocation, behavioral change, and error transitions. Conditional Modality Gain captures the net change in accuracy, while Repair and Damage rates distinguish useful correction from cross-modal interference. Attention analysis complements these measures by showing whether the additional modality is accessed during generation, but it does not replace them. The same dissociation was already visible in *Lost in Translation*. In that study, Qwen2.5-Omni-7B allocated a non-negligible part of its attention budget to raw audio even when visual evidence was available, yet the 32-frame configuration changed from 82.70% without audio to 82.68% with audio. Raw audio became substantially beneficial only when visual evidence was removed, increasing black-frame accuracy from 67.40% to 73.19% [4]. The present analysis extends this ob-

servation: what matters is not simply whether a model attends to audio, but whether that attended acoustic information is transformed into discriminative evidence that changes the VSV decision in a useful direction. Absolute attention percentages should not be compared directly across the two studies because the extraction and aggregation procedures differ; the relevant continuity is the qualitative separation between attentional access and behavioral utility. This interpretation is also consistent with the original design of MAIA. The benchmark was constructed from visual evidence, and audio-based questions were explicitly excluded during annotation [24]. Strong allocation to raw audio must therefore be interpreted as exploitation of an auxiliary source that can contain redundant or correlated information about visually defined events, not as evidence that the benchmark itself requires audio. The attention results consequently reinforce the need for the controlled modality ablations used here: a model may devote substantial computational attention to an auxiliary channel even when that channel adds little task-relevant information, while another model may use a smaller attentional share more selectively. Finally, the attention analysis complements the preliminary visual analysis by operating at a different stage of processing. The preliminary analysis measures whether entities, events, and relations required by a question are recoverable from visual evidence before the question is exposed. The present analysis instead observes allocation during question-conditioned generation, when the task formulation can direct processing toward specific portions of the multimodal context. A high visual attention share therefore does not imply that the visual representation already contains the semantic structure required by the question. Conversely, an incomplete question-independent representation does not prevent the model from later retrieving useful evidence once the linguistic task constrains what should be inspected.



Figure 4.1: Attention allocation across the two multimodal conditions for the four evaluated models. The figure reports, from top to bottom, Qwen2.5-Omni-3B, Qwen-30B, Gemma-3n-E4B, and Gemma-4-12B.

5. Conclusions

Throughout this work, the experimental approach has progressively sharpened a single question: not simply whether a model gives the right answer, but where that answer comes from. The focus has been on what information the model actually has access to, and whether the different modalities are genuinely integrated during reasoning, rather than merely present. The MAIA framework opens this line of inquiry. Its competence-oriented evaluation shows that high performance on individual items does not translate into robust multimodal understanding. When correctness is required to hold across minimally varied formulations and across both discriminative and generative tasks, the models become noticeably more fragile. The benchmark exposes a gap between producing a plausible-looking answer and maintaining a video representation that is grounded, consistent, and stable across different linguistic phrasings of the same underlying content. The difficulty models show with temporally distributed, relational, and implicit phenomena suggests that the core limitation goes beyond object recognition, extending to the construction of a systematic representation of events and how they relate to one another. *Lost in Translation* extends this line of inquiry from visual grounding to cross-modal interaction. Adding audio shows that the presence of an additional information channel does not by itself guarantee that a model exploits it. Audio can help when visual evidence is weak, but its effect shrinks once strong visual evidence is already available. Converting the same acoustic content into a transcription changes its effect on the model’s behavior, so raw audio and its textual counterpart cannot be treated as interchangeable representations of the same information. This points to a basic property of current multimodal systems: how much a modality contributes depends on what other information is already available and on the representational form in which that additional evidence is supplied. MAIA-AV takes this observation as its central object of study. The results show that multimodal reasoning is not simply the sum of independent unimodal capabilities. Audio, video, and transcription each supply evidence that partially overlaps but is not functionally identical. An additional modality can correct an error, leave a decision unchanged, or overturn a decision that was

previously correct. Improved overall accuracy is therefore only one possible outcome of adding a modality; just as informative are the cases where a new modality changes a decision without any net gain, which show that integration involves competition between information sources as much as complementarity. The unimodal conditions clarify this further. The strong contribution of audio on its own shows that speech, environmental sound, and their timing can carry information directly relevant to visually defined events. Transcription behaves differently, confirming that linguistic accessibility and perceptual depth are not the same property: once information is expressed in language, larger models draw on semantic associations more effectively, while increasing model size does little to resolve the difficulty of extracting information from visual input in the first place. This points to a bottleneck that sits before linguistic reasoning even begins, at the point where dynamic perceptual evidence must first be converted into the structured semantic representation the task requires. The preliminary analysis supports the same interpretation. Models recover entities more reliably than they build complete representations of events and the relations between them. Evidence retrieved before the question is known correlates with downstream success but does not determine it: the question itself can direct the model’s attention toward evidence it would otherwise miss. Correct task execution therefore depends on several stages, the initial availability of perceptual evidence, question-guided access to it, and its transformation into a representation suitable for comparison or inference, and a correct answer can still emerge even when one of these stages is only partially completed. This is why behavioral accuracy on its own cannot be taken as proof of genuine multimodal comprehension. The no-input experiments confirm this from the opposite direction. Some decisions remain possible even with no perceptual evidence at all, showing that linguistic plausibility, world knowledge, and structural properties of the answer options can partly resolve the task by themselves. This does not make the perceptual component of the evaluation less important; it shows why it is necessary to separate what a model derives from the stimulus from what it already knows going in. A model should not be credited with genuine grounding simply because it reaches the correct answer, if that same answer is reachable through non-perceptual shortcuts. The attention analysis adds a further layer to this picture. The models tested here do not distribute attention across

audio, video, transcription, textual context, and their own generated history in the same way. The same modality can receive very different amounts of attention depending on the architecture and on the form in which the information is presented, and similar levels of attention can still lead to different behavioral outcomes. Attention therefore reflects access to information, but not whether that information is useful, discriminative, or causally responsible for the final decision. Evaluating multimodal integration means relating internal attention allocation to observed behavioral change, rather than treating either one as sufficient evidence on its own. Together, these results trace a progression from *multimodal availability* to *multimodal integration*. MAIA showed that visual access and correct prediction alone do not guarantee robust grounding. *Lost in Translation* showed that adding a modality does not guarantee that it contributes useful information. MAIA-AV shows that even when several modalities are each individually informative, using them together remains selective, asymmetric, and strongly shaped by architecture. The central limitation of current systems is accordingly not an inability to perceive each modality on its own, but a difficulty in building, selecting, and combining the information each one provides in a way that actually matches what the reasoning task requires. Seen this way, genuine multimodal competence is not defined by how many input modalities a model can accept. It is defined by the ability to identify relevant evidence, keep complementary information from different representational spaces available, discard redundant or misleading signals, and combine what remains into a single coherent semantic representation. The results from MAIA-AV show that current models already have parts of this competence, even if putting them together is still far from complete. Multimodality, in this sense, is not a property of the input interface. It is a reasoning ability that has to be evaluated in its own right.

6. Future Perspectives

The four given directions can be interpreted as one continuous line segment motion rather than four discrete motions. The goal is to gradually develop a more accurate diagnosis of multimodal integration. Starting from differences in behaviors between reasoning categories, then moving to the structure of what each model retrieves from the video, then into inner dynamics of attention, and eventually concludes in causal verification.

Expansion to the Twelve Reasoning Categories The experiments in this thesis treat spatial, temporal, and causal reasoning as the primary axes of analysis, without asking whether modality contribution itself varies across the full set of reasoning phenomena covered by *MAIA*. A natural extension is to apply the same behavioural and modality-contribution metrics separately within each of the twelve categories.

Joint Analysis of Semantic Availability, Attention, and Behaviour This direction develops directly out of the attention analysis introduced in this thesis. Rather than treating semantic availability, attention allocation, and behavioural contribution as three separate questions, a natural next step is to study their joint relationship:

$$\text{Semantic Availability} \longleftrightarrow \text{Attention Allocation} \longleftrightarrow \text{Behavioural Contribution}. \quad (6.1)$$

The question is no longer only *whether the model found a given piece of information, how much it attends to a given modality, or how much accuracy improves when that modality is added*, but **how these three quantities relate to one another at the level of individual items**. This connects directly back to the question raised by *Lost in Transcription* [4]: whether a modality that is available is also genuinely exploited, and whether its internal use corresponds to an observable contribution to the final decision, rather than being available, attended, and yet behaviourally inert.

List of Abbreviations

VLM	Visual Language Model
VQA	Visual Question Answering
AGQA	Action Genome Question Answering
MVBench	Multi-modal Video Understanding Benchmark
VILMA	Video Language Model Assessment
MAIA	Multimodal AI Assessment
VSV	Visual Statement Verification
OEVQA	Open-Ended Visual Question Answering
MLM	Multimodal Language Model
AVQA	Audio-Visual Question Answering
Dave	Diagnostic benchmark for Audio Visual Evaluation
MMAU-Pro	Massive Multi-Task Audio Understanding and Reasoning benchmark, Pro version
SONIC-01	SOcial Natural Interaction Corpus (Omnimodal, v1)
GPT	Generative Pre-trained Transformer
MIRAGE	Assessing Hallucination in Multimodal Reasoning Chains of Multimodal Large Language Models

Bibliography

- [1] Stanislaw Antol et al. “VQA: Visual Question Answering”. In: *Proceedings of the IEEE International Conference on Computer Vision*. 2015, pp. 2425–2433.
- [2] Marco Baroni. “Grounding Distributional Semantics in the Visual World”. In: *Language and Linguistics Compass* 10 (2015). DOI: 10.1111/lnc3.12170.
- [3] Zesen Cheng et al. “VideoLLaMA 2: Advancing Spatial-Temporal Modeling and Audio Understanding in Video-LLMs”. In: *arXiv preprint arXiv:2406.07476* (2024). arXiv: 2406.07476. URL: <https://arxiv.org/abs/2406.07476>.
- [4] Michela Deodati et al. “Lost in Transcription: Investigating the Role of Audio Information for Video-Based Visual Statement Verification”. In: *Proceedings of the Twelfth Italian Conference on Computational Linguistics (CLiC-it 2026)*. Forthcoming. 2026.
- [5] Bowen Dong et al. “MIRAGE: Assessing Hallucination in Multimodal Reasoning Chains of MLLMs”. In: *Advances in Neural Information Processing Systems*. 2025. URL: <https://bit.ly/25mirage>.
- [6] Fondazione Bruno Kessler. *Fondazione Bruno Kessler Official Website*. <https://www.fbk.eu/>. Accesso effettuato il 29 agosto 2026. 2026.
- [7] Gemma Team. *Gemma 4 Technical Report*. 2026. DOI: 10.48550/arXiv.2607.02770. arXiv: 2607.02770 [cs.CL]. URL: <https://arxiv.org/abs/2607.02770>.
- [8] Rohit Girdhar et al. “ImageBind: One Embedding Space to Bind Them All”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2023, pp. 15180–15190.
- [9] Madeleine Grunde-McLaughlin, Ranjay Krishna, and Maneesh Agrawala. “AGQA: A Benchmark for Compositional Spatio-Temporal Reasoning”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2021, pp. 11287–11297.

- [10] Drew A. Hudson and Christopher D. Manning. “GQA: A New Dataset for Real-World Visual Reasoning and Compositional Question Answering”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2019, pp. 6700–6709.
- [11] Ilker Kesen et al. “VILMA: A Zero-Shot Benchmark for Linguistic and Temporal Grounding in Video-Language Models”. In: *International Conference on Learning Representations*. 2024. arXiv: 2311.07022. URL: <https://arxiv.org/abs/2311.07022>.
- [12] Sonal Kumar et al. “MMAU-Pro: A Challenging and Comprehensive Benchmark for Holistic Evaluation of Audio General Intelligence”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 40. 2026, pp. 22688–22697. URL: <https://sonalkum.github.io/mmau-pro>.
- [13] Guangyao Li et al. “Learning to Answer Questions in Dynamic Audio-Visual Scenarios”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2022, pp. 19108–19118.
- [14] K. Li et al. “MVBench: A Comprehensive Multi-Modal Video Understanding Benchmark”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2024. DOI: 10.48550/arXiv.2311.17005. arXiv: 2311.17005.
- [15] Pouya Pezeshkpour and Estevam Hruschka. “Large Language Models Sensitivity to The Order of Options in Multiple-Choice Questions”. In: *Findings of the Association for Computational Linguistics: NAACL 2024*. Ed. by Kevin Duh, Helena Gomez, and Steven Bethard. Mexico City, Mexico: Association for Computational Linguistics, June 2024, pp. 2006–2017. DOI: 10.18653/v1/2024.findings-naacl.130. URL: <https://aclanthology.org/2024.findings-naacl.130/>.
- [16] Qwen Team. “Qwen2.5-Omni Technical Report”. In: *arXiv preprint arXiv:2503.20215* (2025). arXiv: 2503.20215. URL: <https://arxiv.org/abs/2503.20215>.

- [17] Qwen Team. “Qwen3.5-Omni Technical Report”. In: *arXiv preprint arXiv:2604.15804* (2026). arXiv: 2604.15804. URL: <https://arxiv.org/abs/2604.15804>.
- [18] Gorjan Radevski et al. “Dave: Diagnostic benchmark for audio visual evaluation”. In: *Advances in Neural Information Processing Systems* 38 (2026).
- [19] Ahmed Y. Radwan et al. “SONIC-O1: A Real-World Benchmark for Evaluating Multimodal Large Language Models on Audio-Video Understanding”. In: *arXiv preprint arXiv:2601.21666* (2026). arXiv: 2601.21666. URL: <https://arxiv.org/abs/2601.21666>.
- [20] Omar Sanseviero and Ian Ballantyne. *Introducing Gemma 3n: The Developer Guide*. Google Developers Blog. June 2025. URL: <https://developers.googleblog.com/en/introducing-gemma-3n-developer-guide/>.
- [21] Ravi Shekhar et al. “FOIL it! Find One Mismatch between Image and Language Caption”. In: *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2017, pp. 255–265. DOI: 10.18653/v1/P17-1024. URL: <https://aclanthology.org/P17-1024/>.
- [22] Guangzhi Sun et al. “Video-SALMONN: Speech-Enhanced Audio-Visual Large Language Models”. In: *arXiv preprint arXiv:2406.15704* (2024). arXiv: 2406.15704. URL: <https://arxiv.org/abs/2406.15704>.
- [23] Davide Testa et al. “All-in-One: Understanding and Generation in Multimodal Reasoning with the MAIA Benchmark”. In: *Findings of the Association for Computational Linguistics: EMNLP 2025*. 2025, pp. 20030–20050. DOI: 10.18653/v1/2025.findings-emnlp.1091. URL: <https://aclanthology.org/2025.findings-emnlp.1091/>.
- [24] Davide Testa et al. “MAIA: A Benchmark for Multimodal AI Assessment”. In: *Proceedings of the Eleventh Italian Conference on Computational Linguistics (CLiC-it 2025)*. 2025, pp. 1121–1134. URL: <https://aclanthology.org/2025.clicit-1.106/>.

- [25] Tristan Thrush et al. “Winoground: Probing Vision and Language Models for Visio-Linguistic Compositionality”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2022, pp. 5238–5248.
- [26] Jin Xu et al. “Qwen2.5-Omni Technical Report”. In: *arXiv preprint arXiv:2503.20215* (2025). DOI: 10.48550/arXiv.2503.20215. URL: <https://arxiv.org/abs/2503.20215>.
- [27] Jin Xu et al. “Qwen3-Omni Technical Report”. In: *arXiv preprint arXiv:2509.17765* (2025). DOI: 10.48550/arXiv.2509.17765. URL: <https://arxiv.org/abs/2509.17765>.
- [28] Chujie Zheng et al. *Large Language Models Are Not Robust Multiple Choice Selectors*. 2024. arXiv: 2309.03882 [cs.CL]. URL: <https://arxiv.org/abs/2309.03882>.